LLMs Unplugged

Understand AI by building it yourself

Speaker notes

Open with the promise: understand AI by building one yourself, by hand. No computers, no maths background needed.

Acknowledgement of Country

Speaker notes

Acknowledge the Traditional Custodians of the land you're meeting on, and pay respects to Elders past and present. Keep it brief and sincere; say it in your own words.

activity

everyone stand up

Speaker notes

Quick energiser. Read each line and let people sit; it speeds up, and by "the last 5 minutes" almost everyone's down. Lands the point that this stuff is already everywhere. ~2 min, don't linger.

sit down if you

have **never** used ChatGPT

sit down if you

haven’t used it in the last **month**

sit down if you

haven’t used it in the last **week**

sit down if you

haven’t used it in the last **day**

sit down if you

haven’t used it in the last **hour**

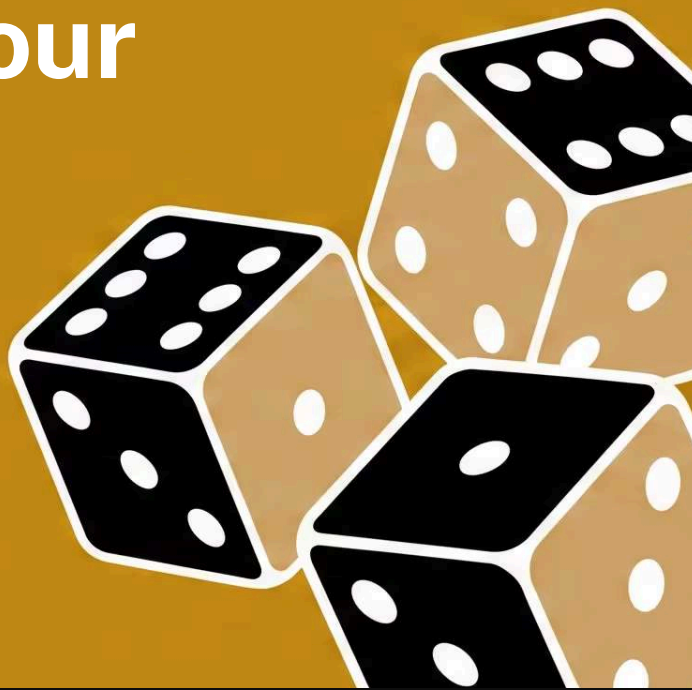

sit down if you

haven’t used it in the last **5 minutes**

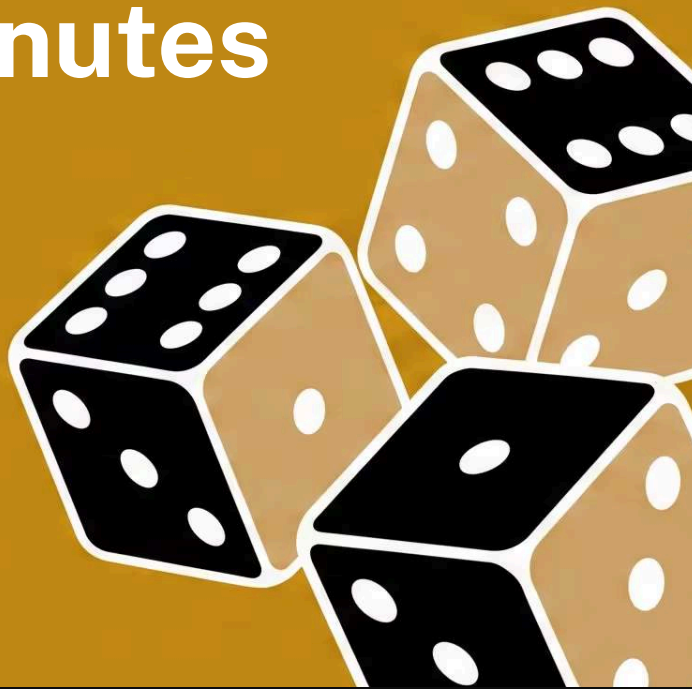

What is this about?

you'll **build** your own language model — from scratch — with just a kids book, pen & paper, and some dice rolling

you'll **learn** how language models work by spotting patterns in text to generate new text

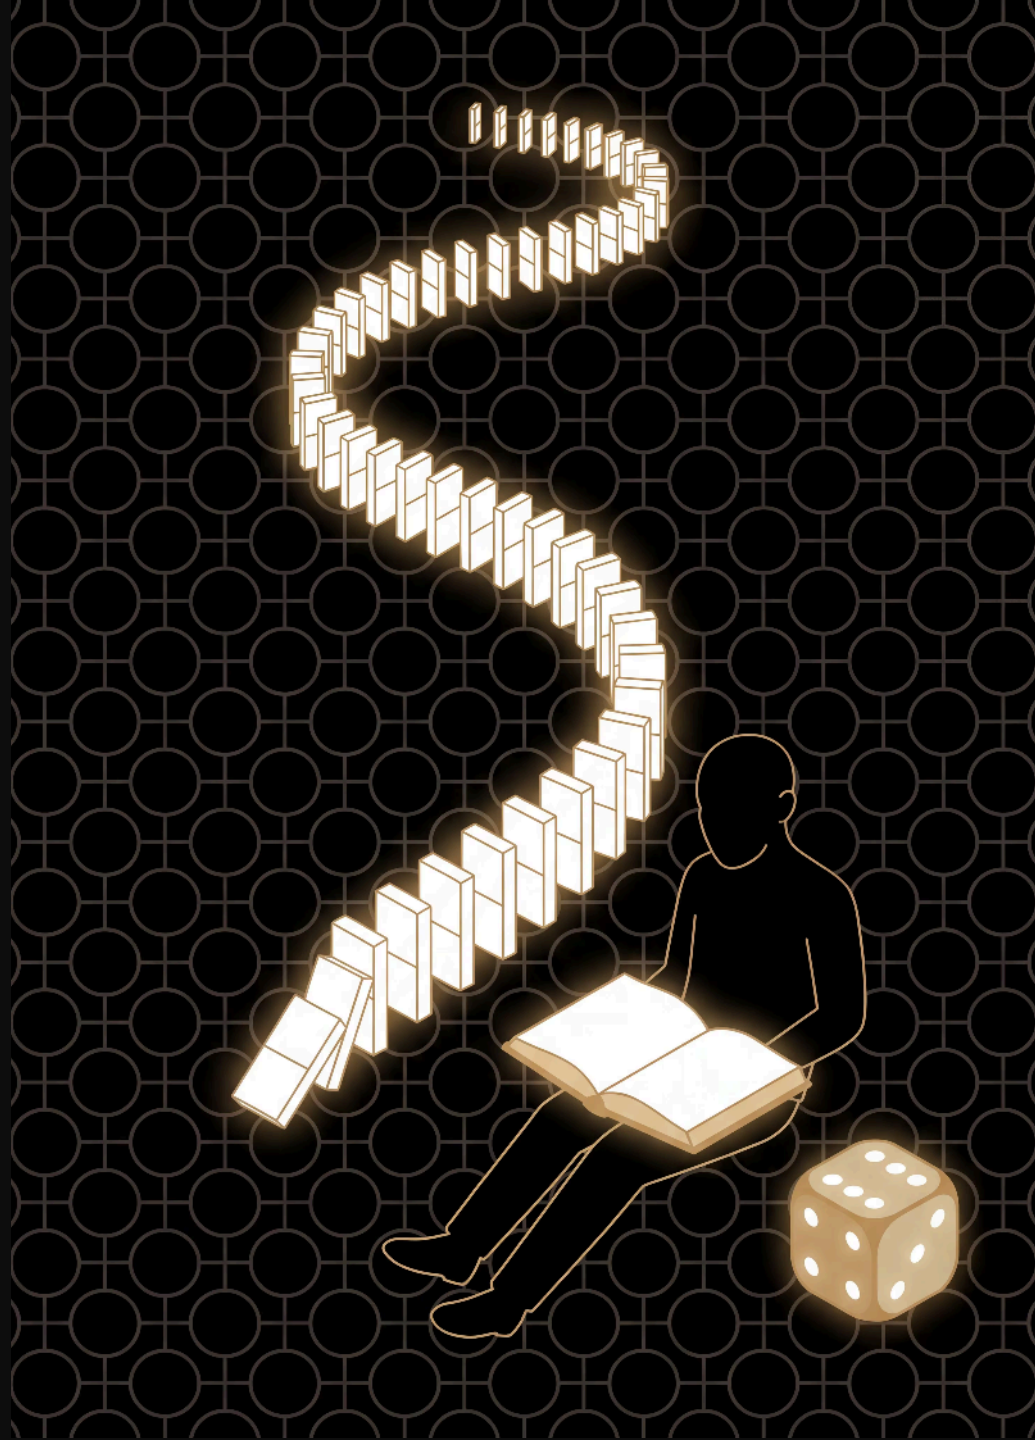

Speaker notes

Frame the next stretch: build a model from a kids' book, pen, paper and dice, by spotting patterns. ~1 min, then move.

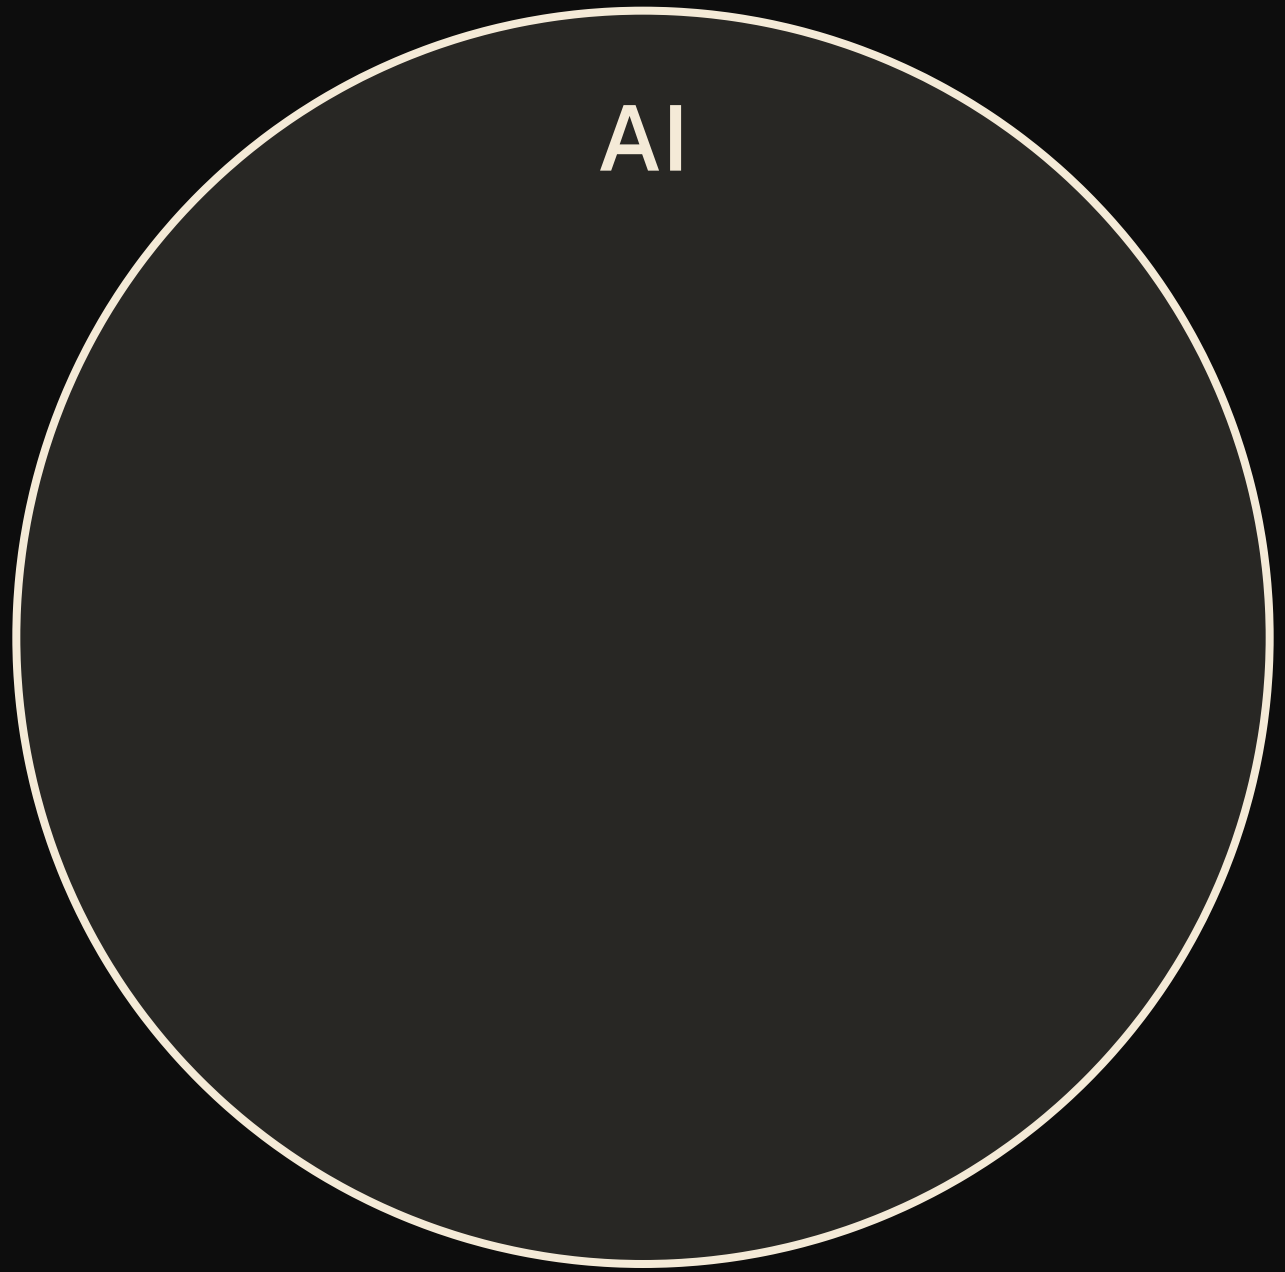
AI

Speaker notes

No text on this one — talk over the build. AI: the umbrella term, any system doing things we'd call "smart" (chess engines, spam filters, route planners). Click — ML: the subset that learns patterns from data instead of following hand-written rules; the model you'll build today lives here. Click — LLMs: machine learning models trained on huge amounts of text to predict the next token — exactly what you're about to do by hand, just scaled up. Click — and the chatbots in the news (Claude, ChatGPT, Gemini, DeepSeek) are LLMs wrapped in an app. Same species as the thing you'll build with pen and dice.

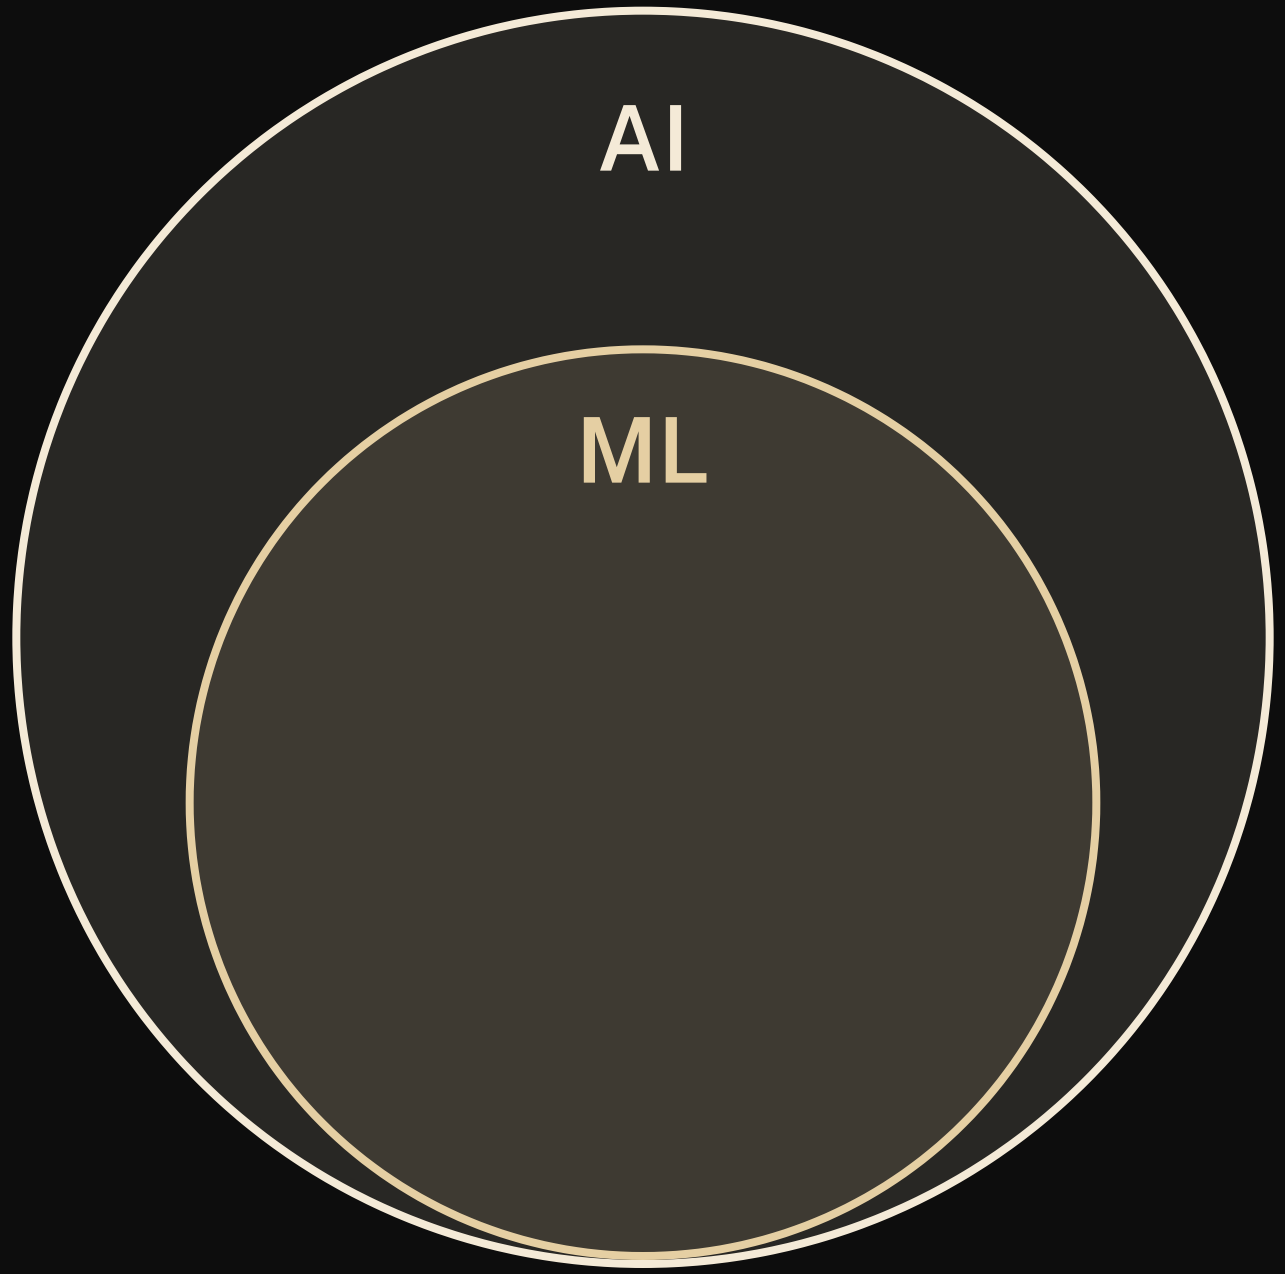
AI
ML

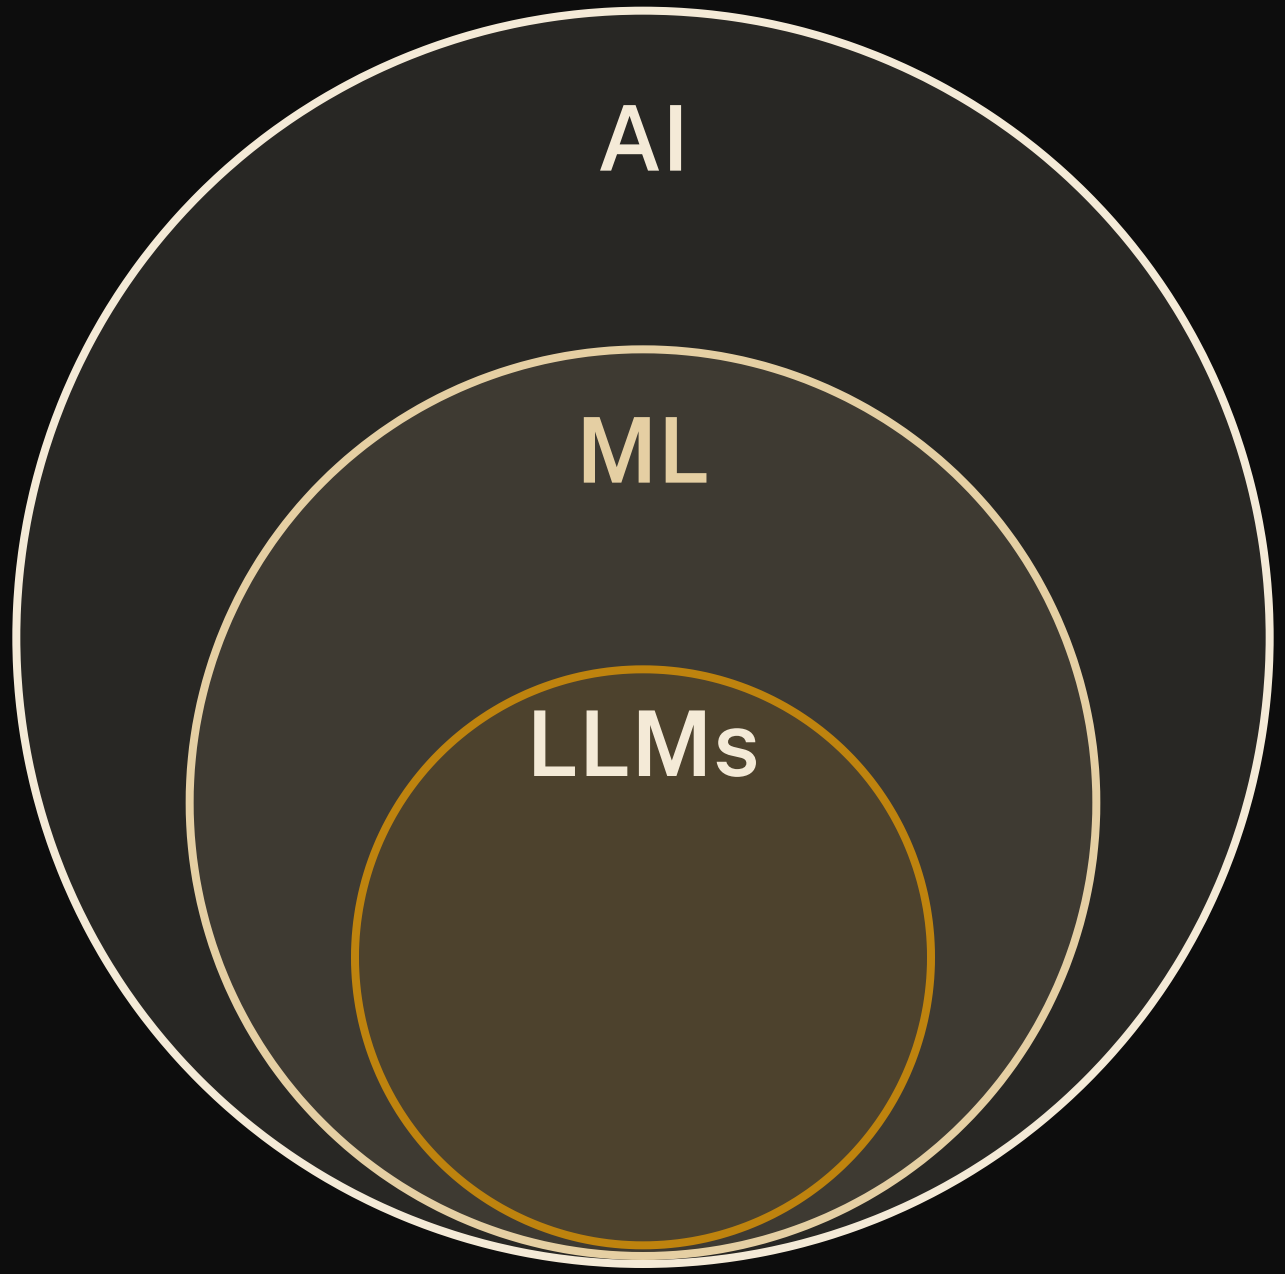
AI
ML
LLMs

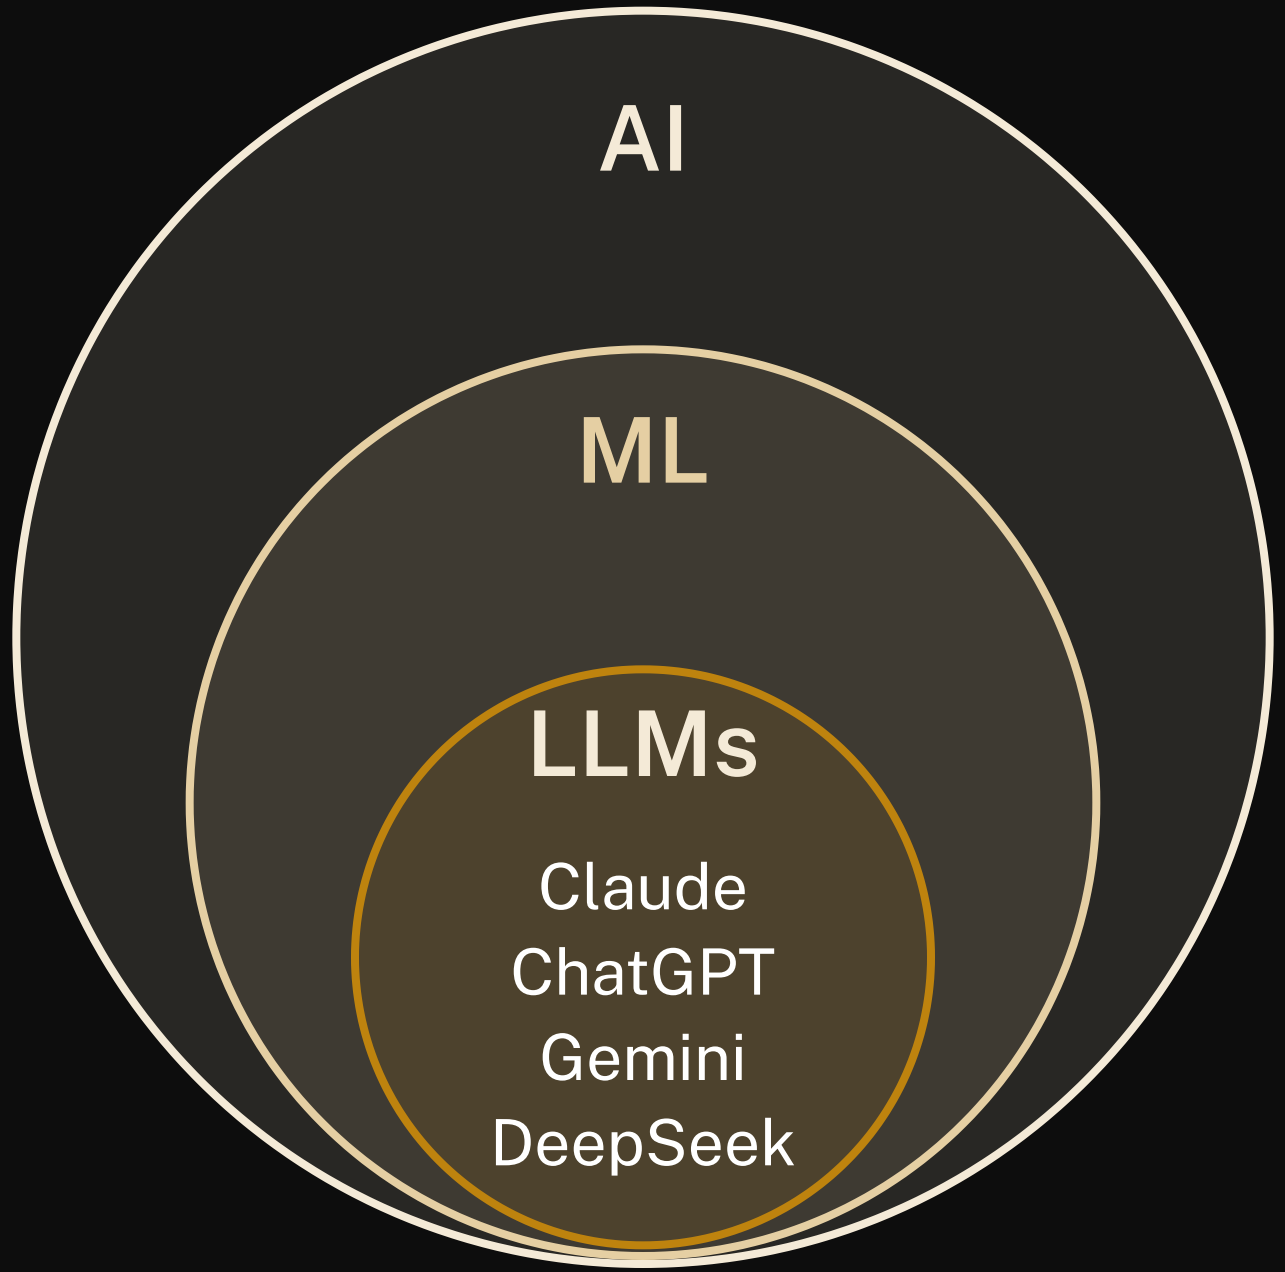
AI
ML
LLMs
Claude
ChatGPT
Gemini
DeepSeek

llmsunplugged.org

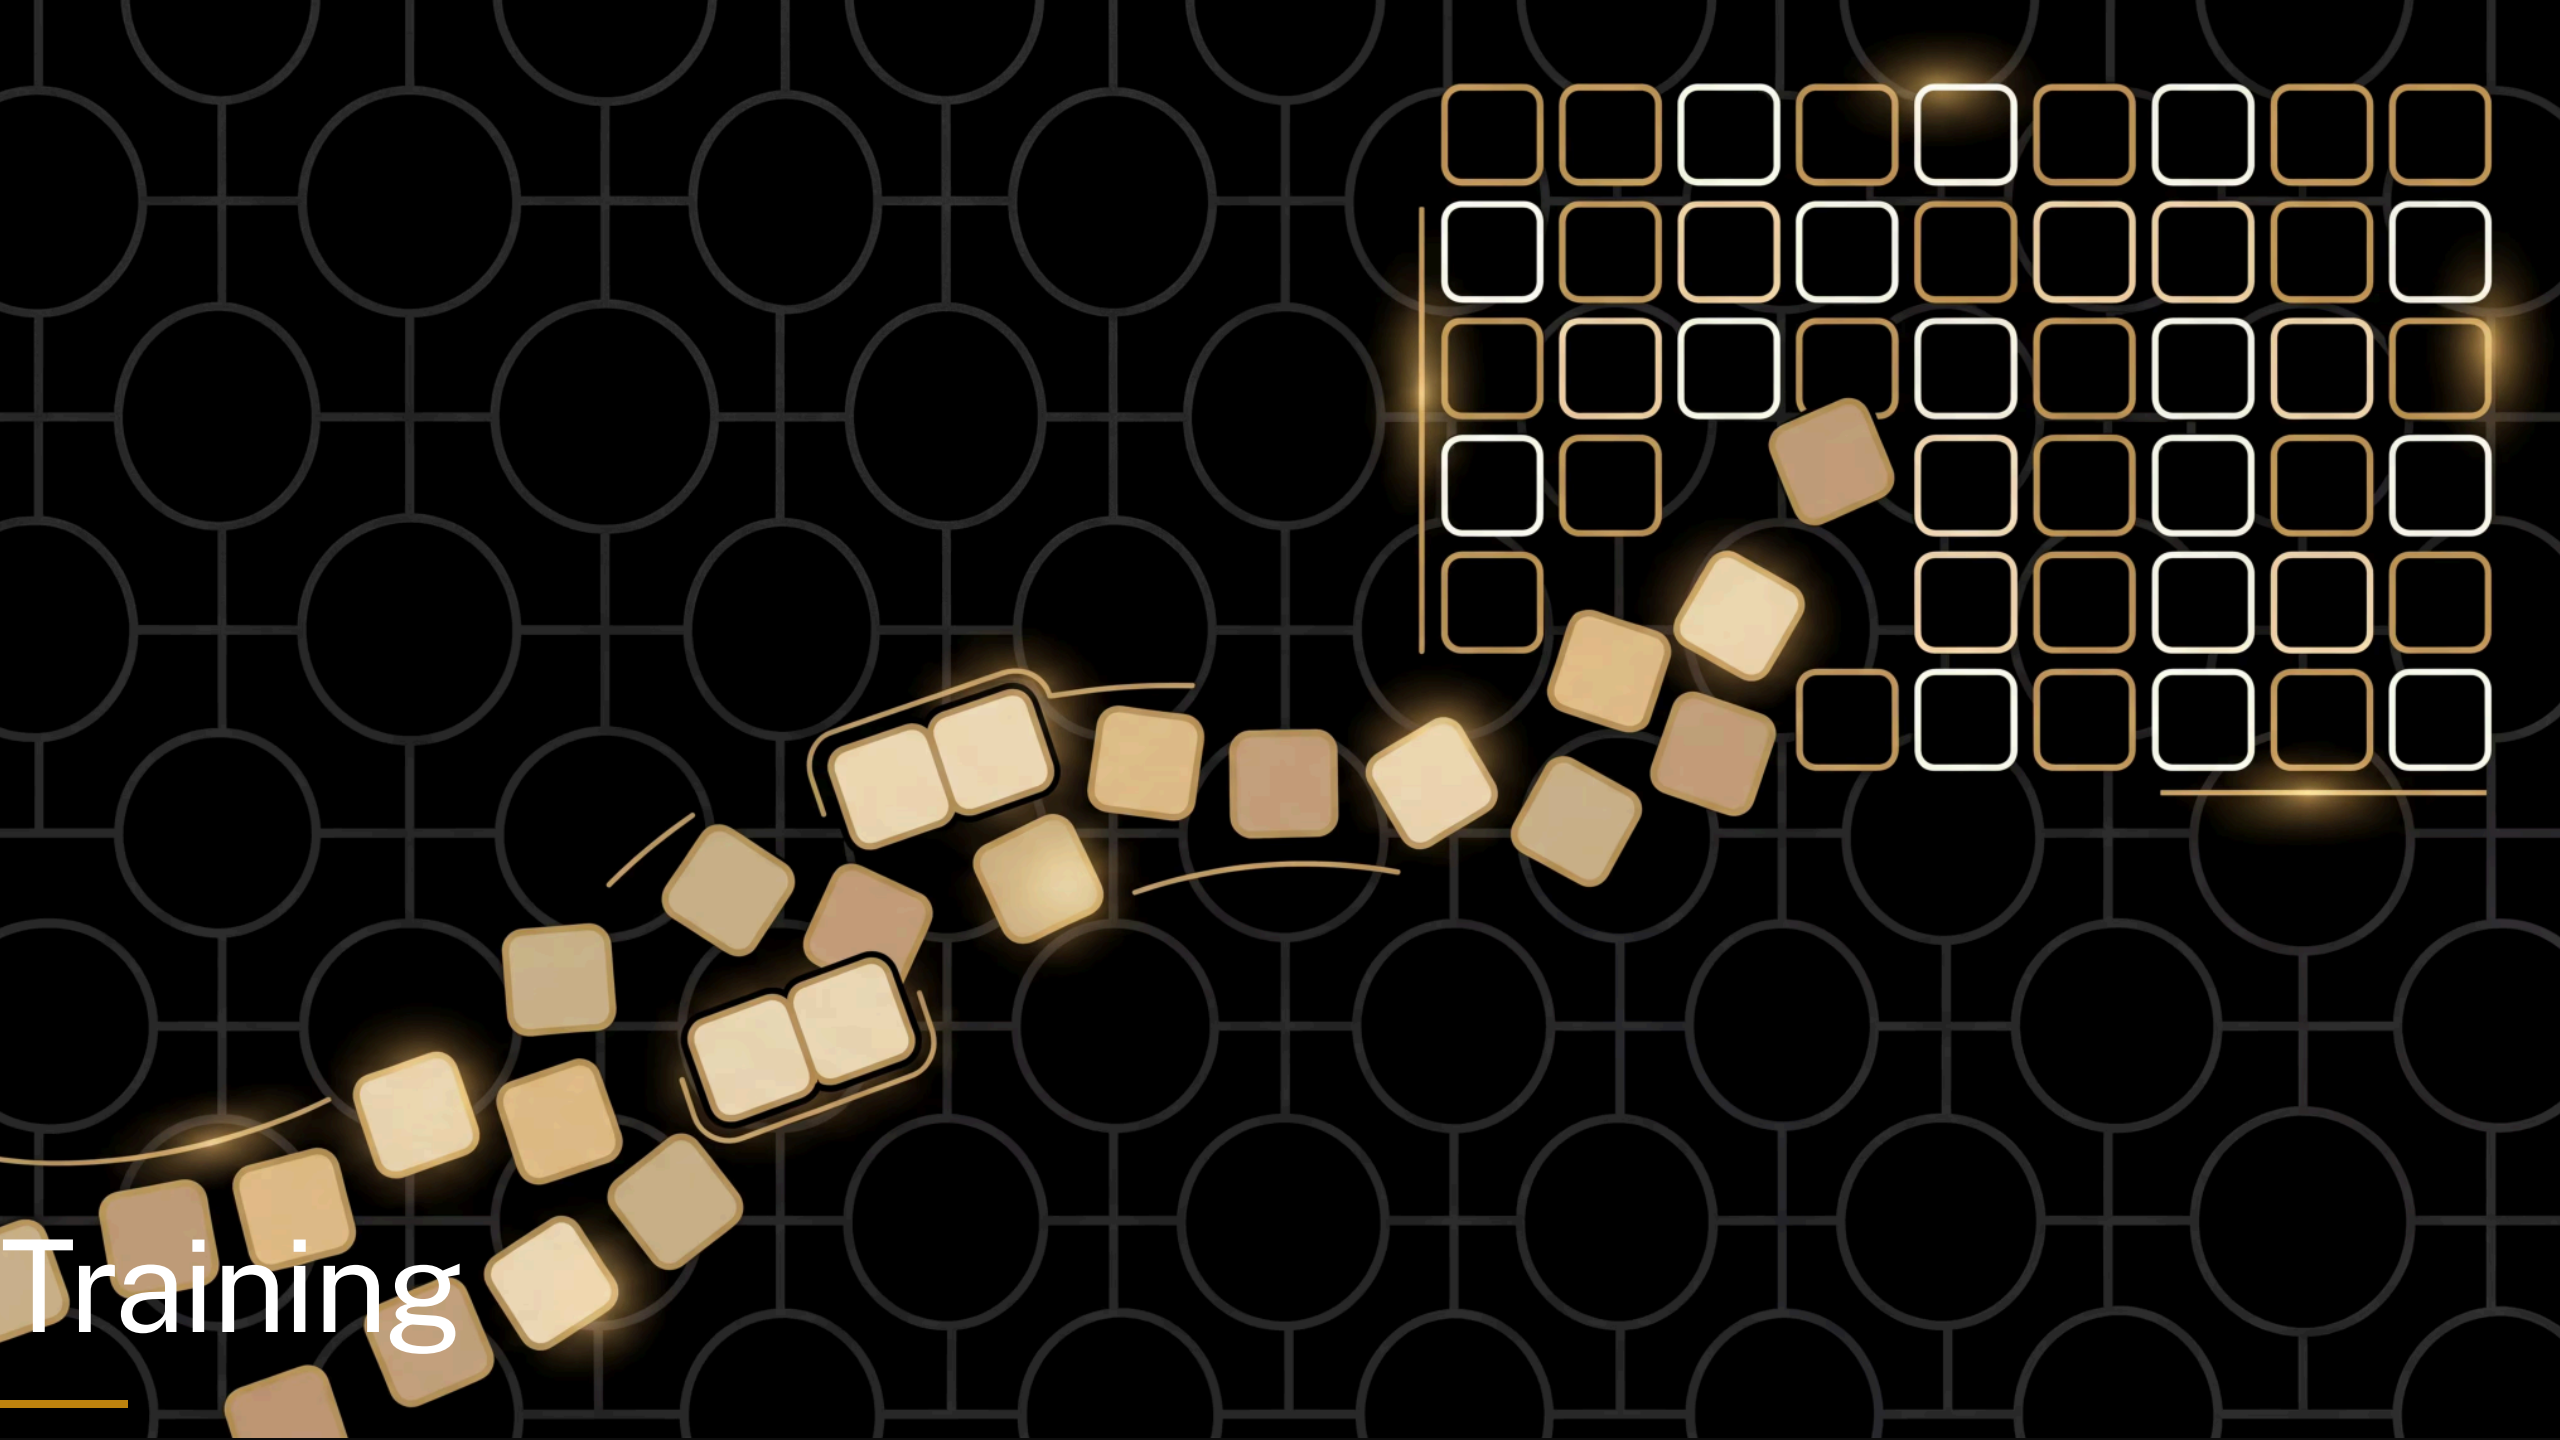
Training

The recipe

walk through your text and tally up which tokens follow which in a grid

Speaker notes

1. **tidy up** your text: lowercase everything, treat words and single punctuation marks (`.` `,` `!` `?` `;` `:`) as separate tokens
2. **set up** your grid: first token goes in both the row & column headers
3. **fill in** the grid: for each consecutive pair of tokens, add a tally mark in that cell (add new rows/columns as needed)

Example

“Run, Spot, run. See Spot run.”

after tidying up:

`run` `,` `spot` `,` `run` `.` `see` `spot` `run` `.`

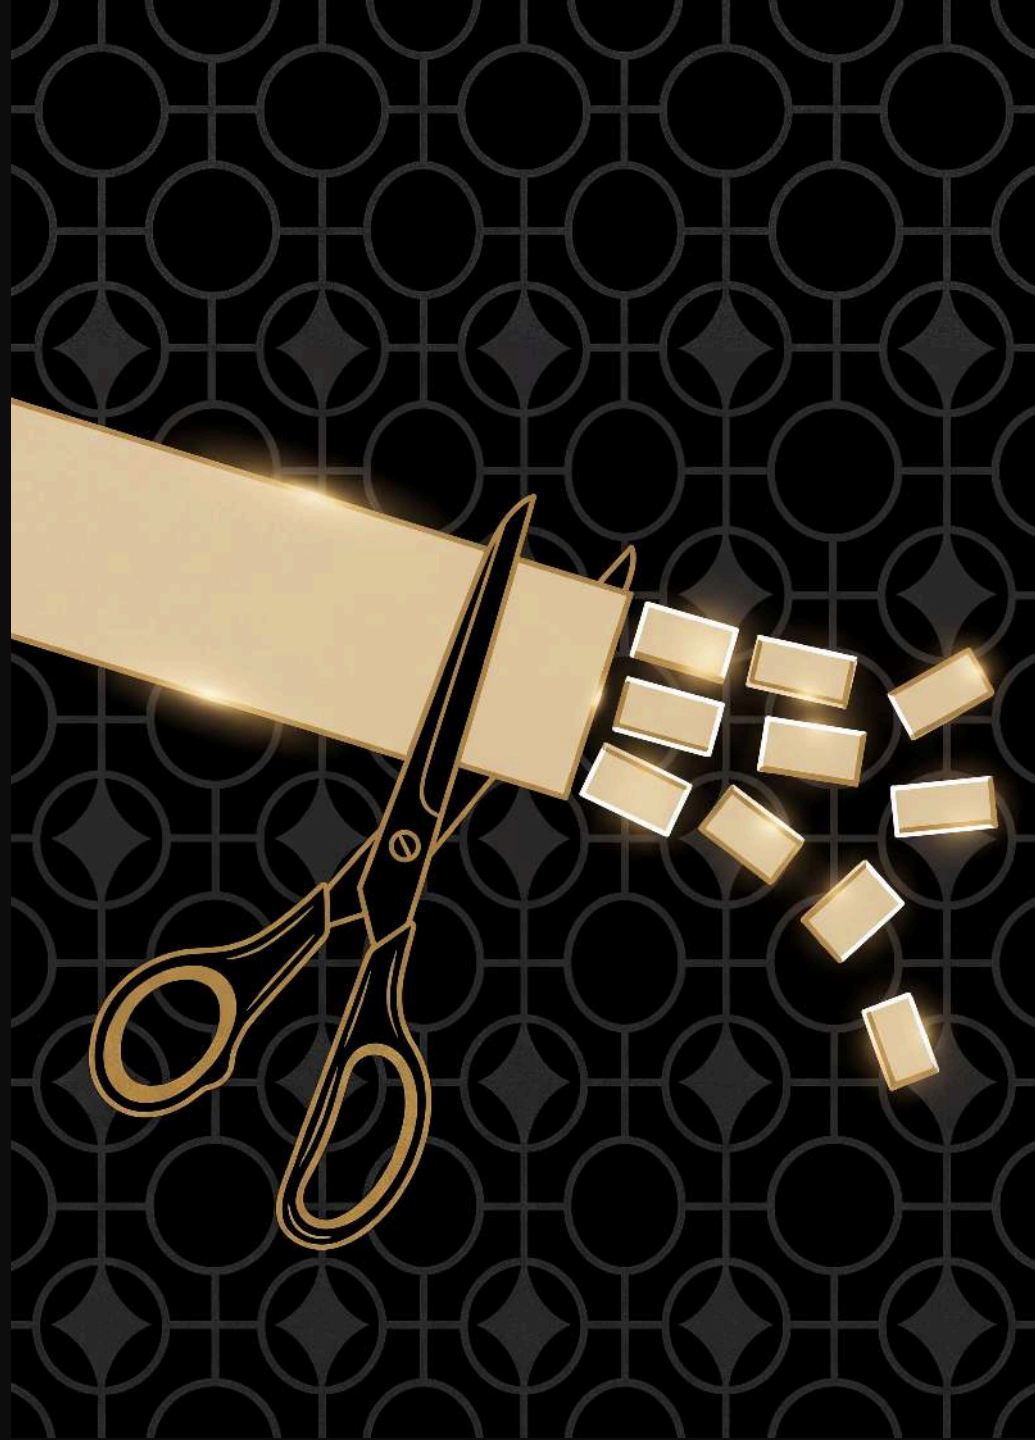

The empty grid

`run` `,` `spot` `,` `run` `.` `see` `spot` `run` `.`

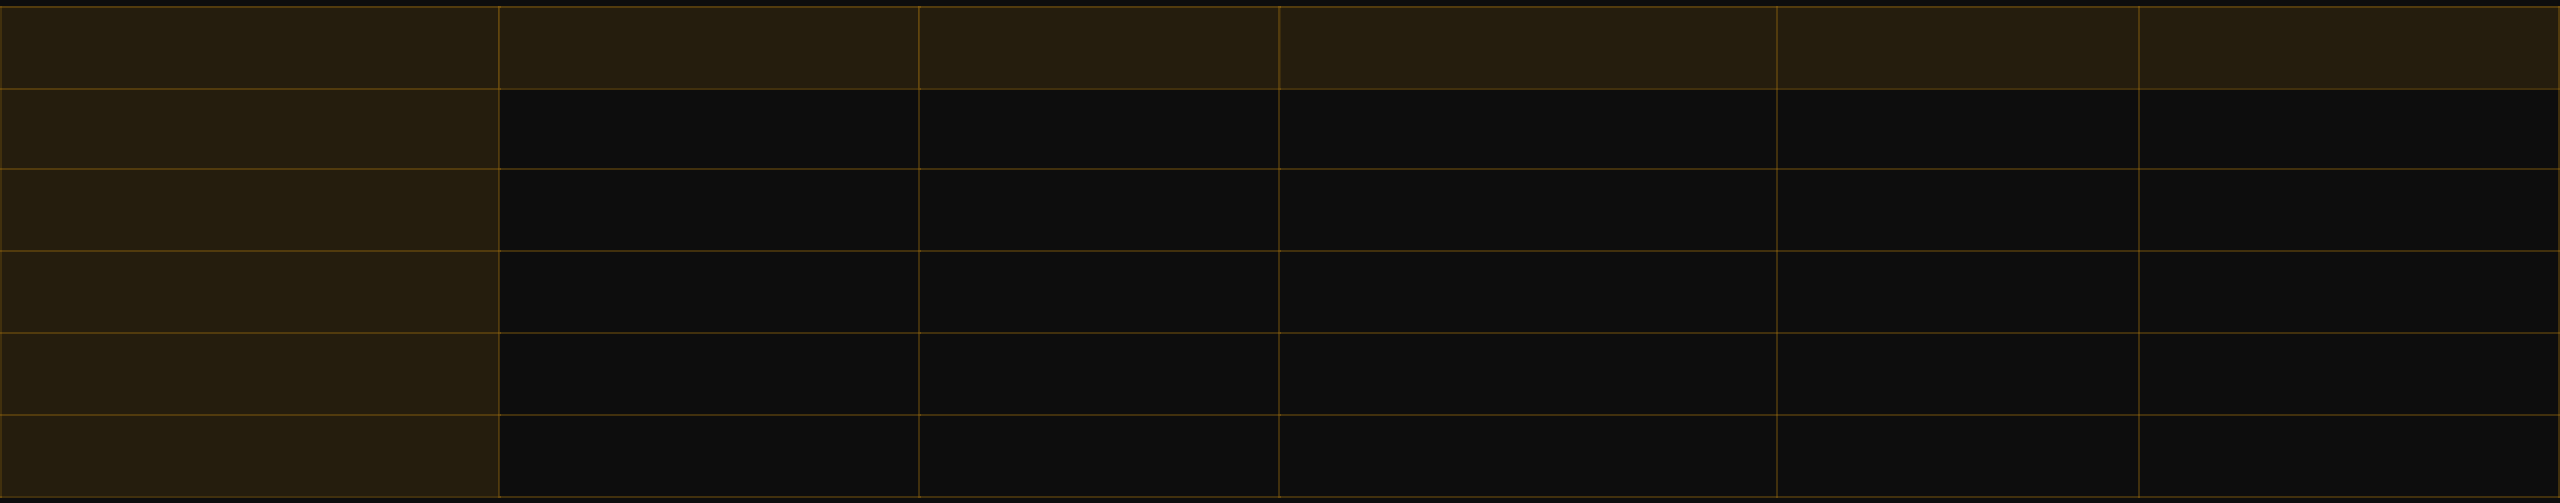

Training: run → ,

run , spot , run . see spot run .

	run	,			
run					
,					

Training: , → spot

run , spot , run . see spot run .

	run	,	spot		
run					
,					
spot					

Training: spot → ,

run , spot , run . see spot run .

	run	,	spot		
run		|			
,			|		
spot		|			

Speaker notes

`,` is already in our vocabulary — no new column needed!

Training: , → run

run , spot , run . see spot run .

	run	,	spot		
run		|			
,			|		
spot		|			

Training: run → .

run , spot , run . see spot run .

	run	,	spot	.	
run		|		|	
,	|		|		
spot		|			
.					

Speaker notes

. is a new token — punctuation gets its own row and column too

Complete model

`run` `,` `spot` `,` `run` `.` `see` `spot` `run` `.`

	run	,	spot	.	see
run		|		||	
,	|		|		
spot	|	|			
.					|
see			|		

Speaker notes

with the rest of the text tallied the same way, the grid is complete. `run` → `.` came up twice, so its cell has two marks — those counts are what make some next-words more likely than others.

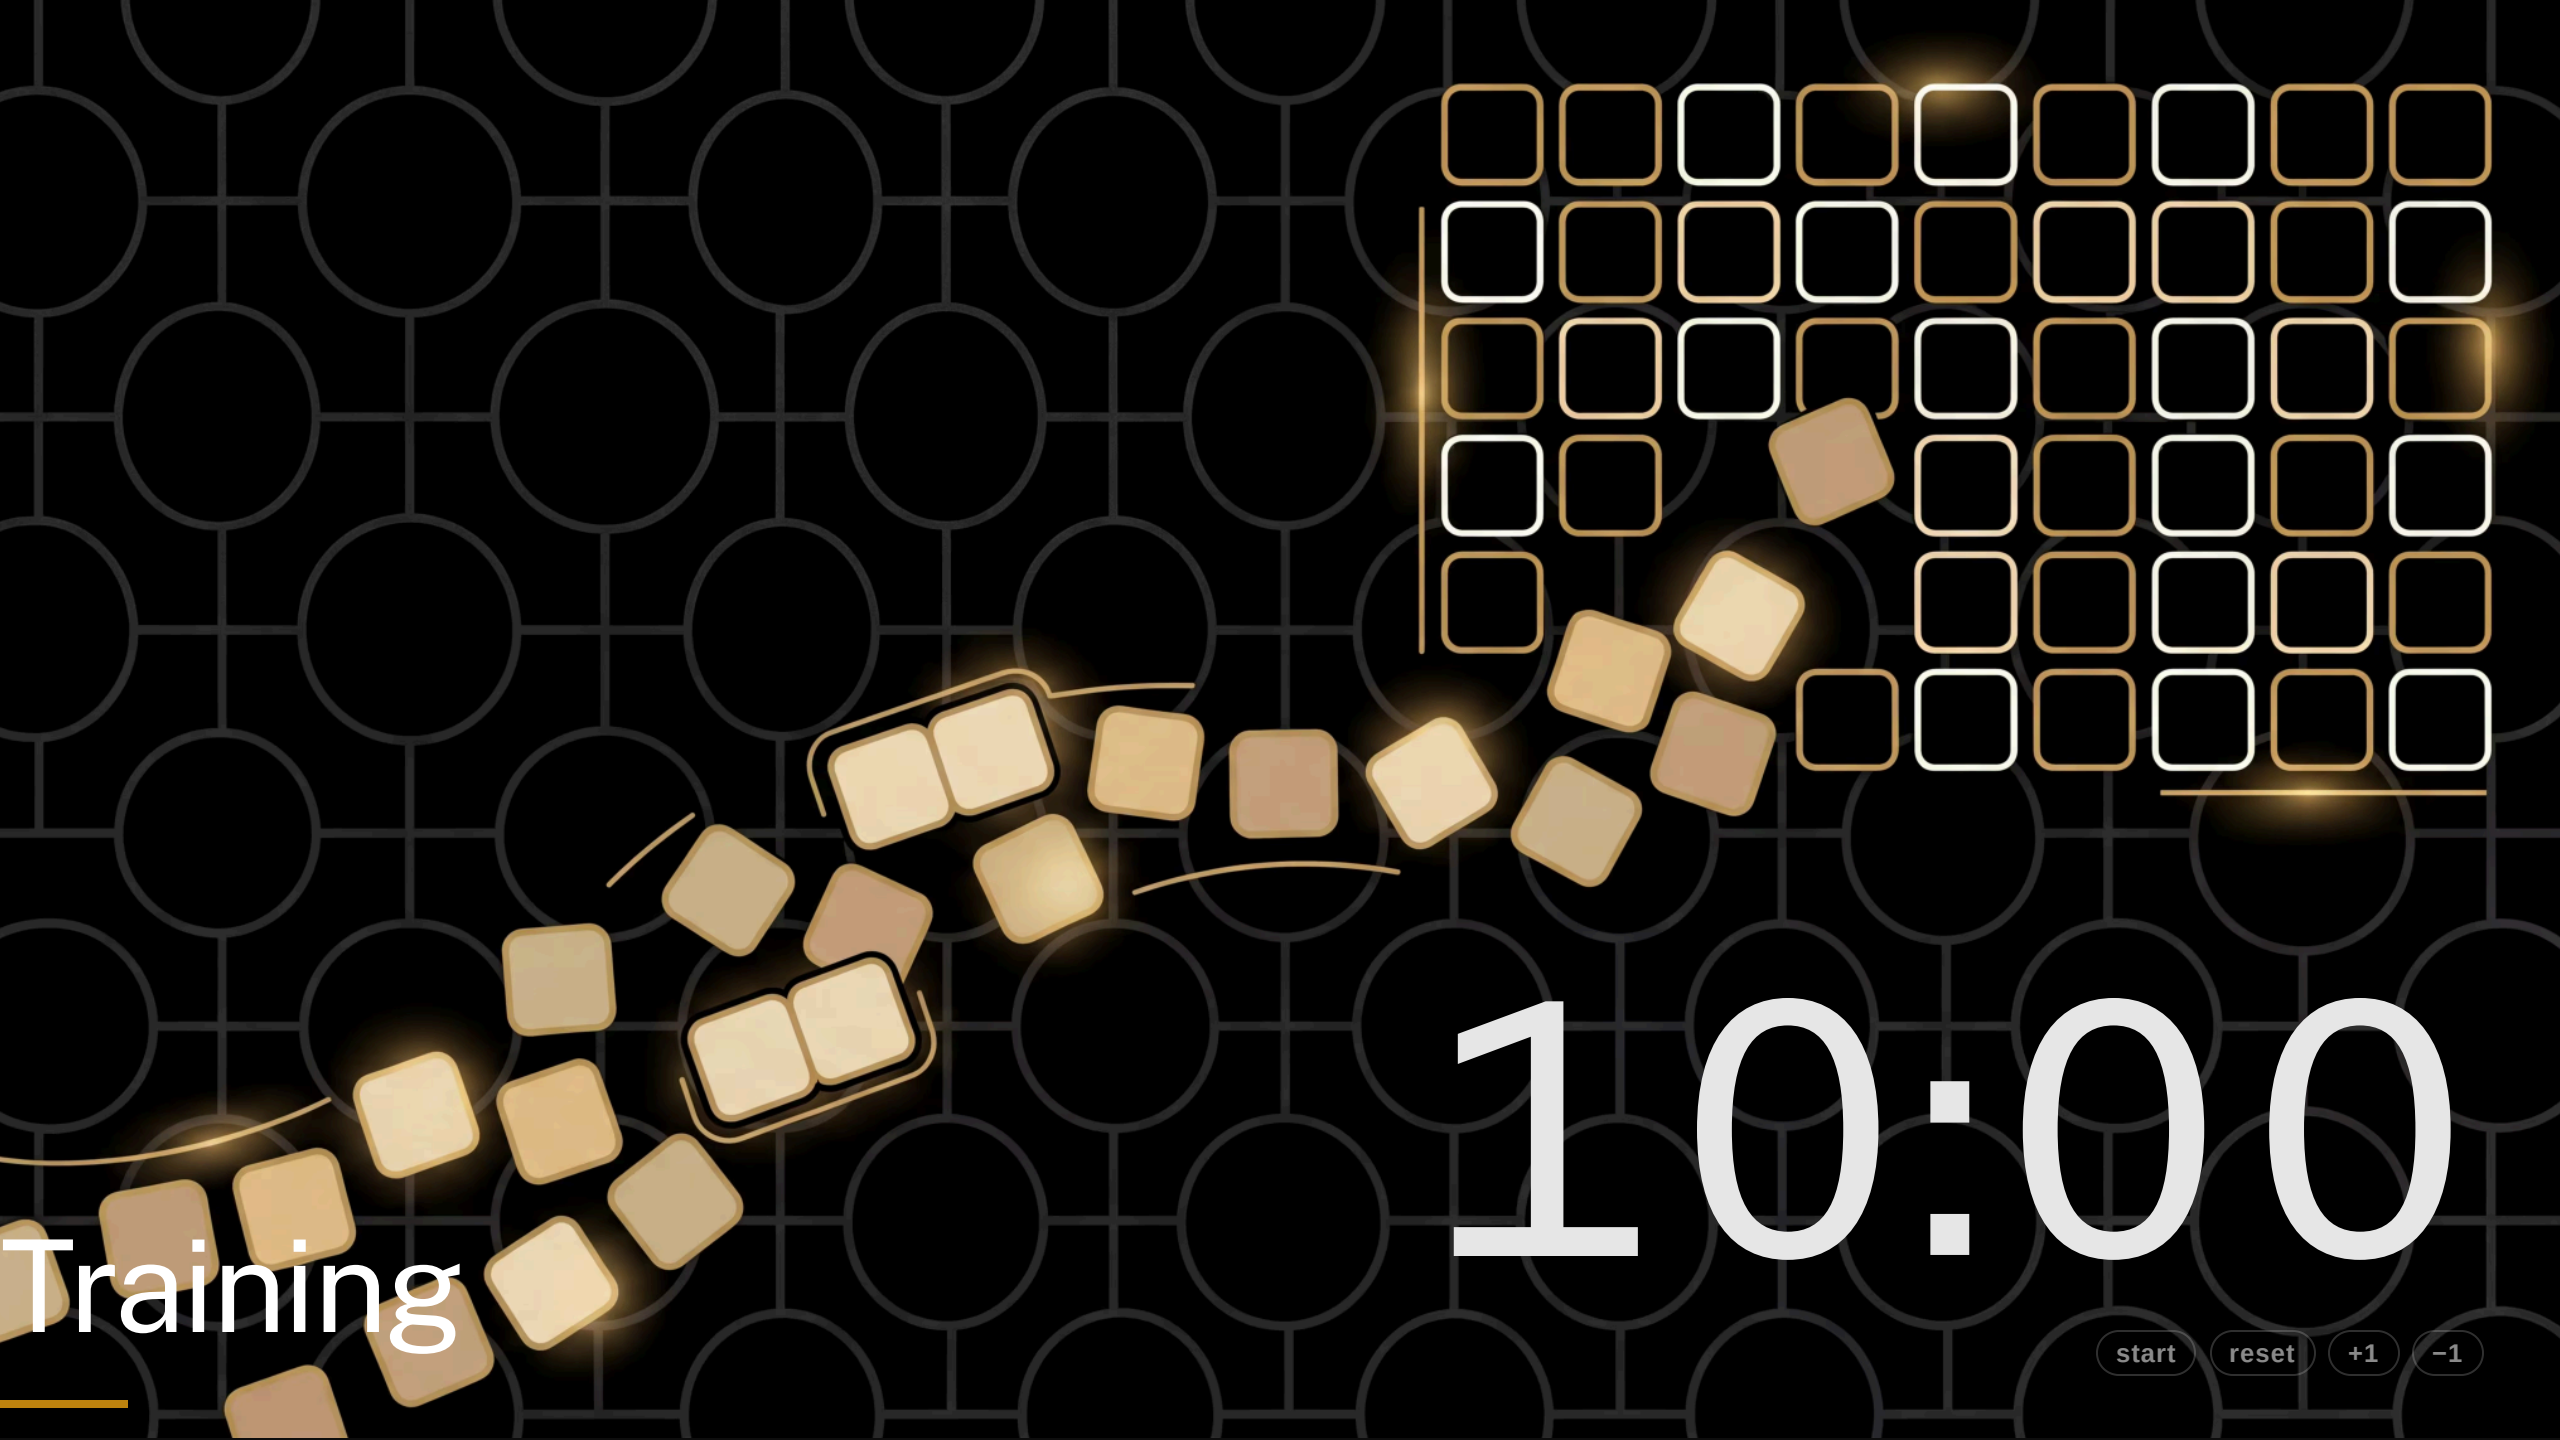
10:00
Training
start
reset
+1
−1

Speaker notes

Hand out grids; groups tally their own text. 10 min, circulate --- the usual snag is that a token only gets a new row and column the first time it appears.

The *language* of language models

- model
- token
- weights

Speaker notes

Everyday sense first, then ours: a model is a fashion model or a model aeroplane — here it's the thing that captures the patterns, your grid. A token is a bus token or a token gesture — here it's one word or punctuation mark. Weights are what you lift at the gym, or what you weigh up before deciding — here they're just the tally marks, the numbers that make some next words likelier than others. Point at a cell: when you hear a model has billions of weights, that's what's being counted.

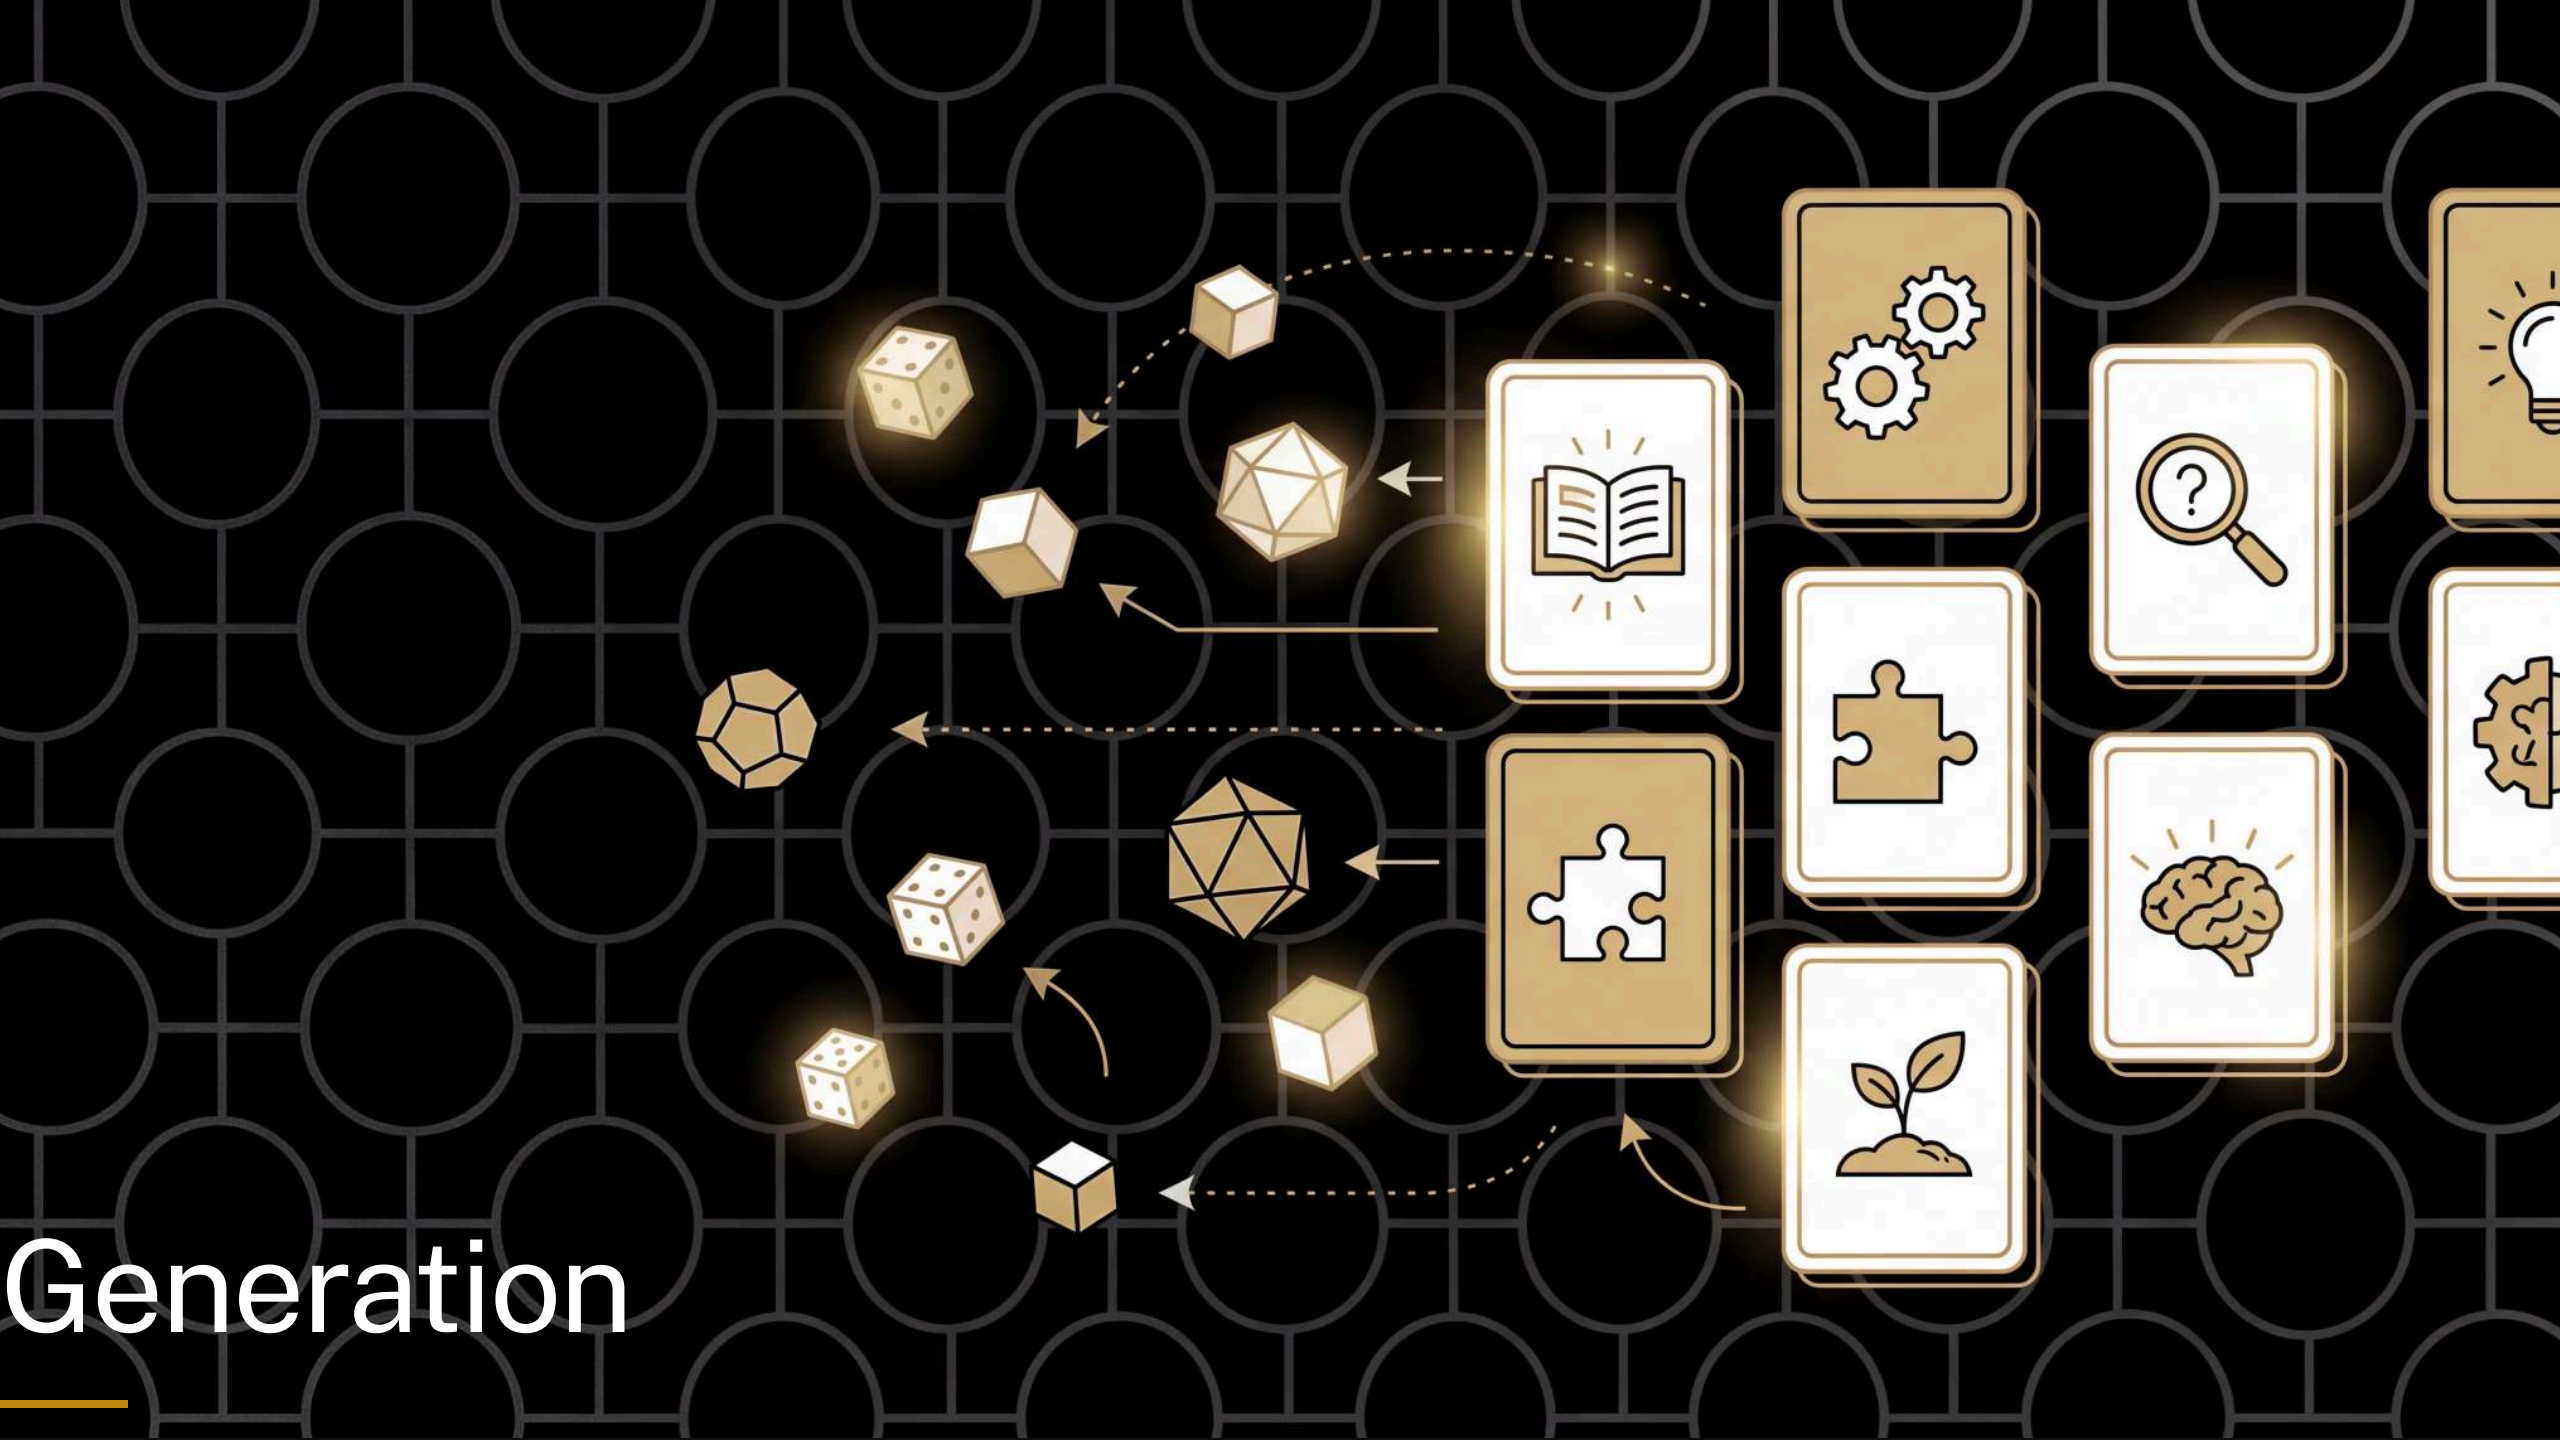

Generation

The recipe

use your grid to generate new text, rolling dice to choose each next word

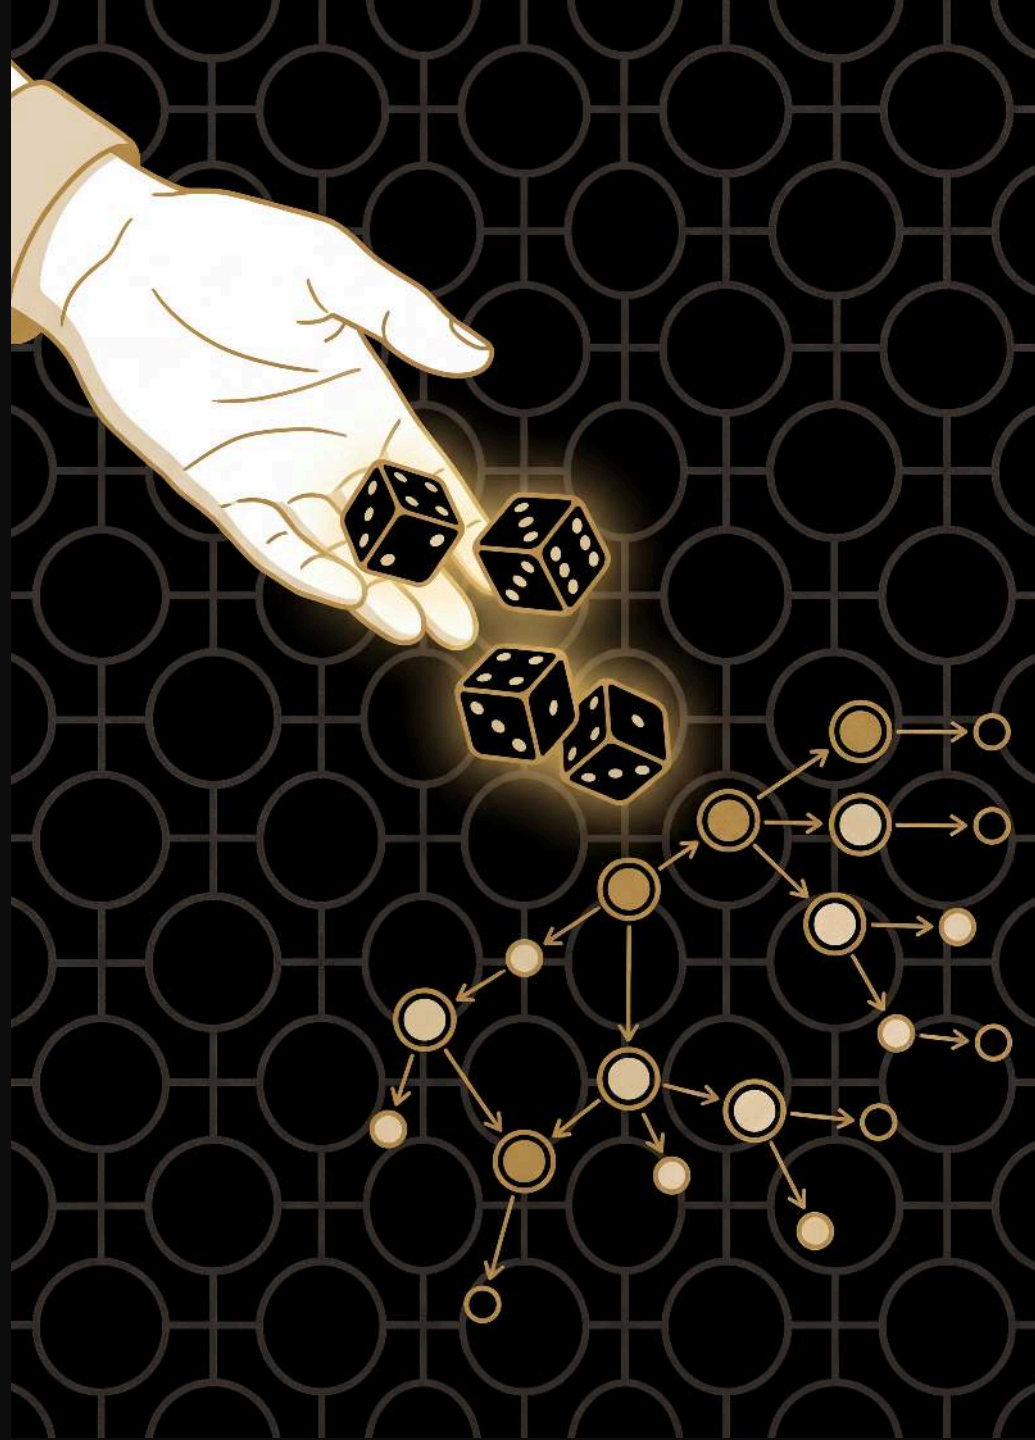

Speaker notes

1. **choose** a starting word from your grid (any word---not just the first one)
2. **look** at that word's row to find all possible next words and their counts
3. **roll dice** weighted by the counts to pick the next word
4. **write down** the chosen word and make it your new starting word
5. **repeat** from step 2

Generation: start with `see`

`see` `spot`

one option — no roll needed

	`run`	`,`	`spot`	`.`	`see`
`run`		|		||	
`,`	|		|		
`spot`	|	|			
`.`					|
`see`			|		

Speaker notes

any word will do as a seed---we're starting from `see`, which is the last row, to show you don't have to begin at the top. Only one option here, so no roll.

Generation: from `spot`

`see` `spot`

2 options — roll the die!

	`run`	`,`	`spot`	`.`	`see`
`run`		|		||	
`,`	|		|		
`spot`	|	|			
`.`					|
`see`			|		

Speaker notes

two options in this row (`run` and `,`) with equal tallies---highlight both, then roll

How the die chooses: `spot`

`spot` → ? roll a d10

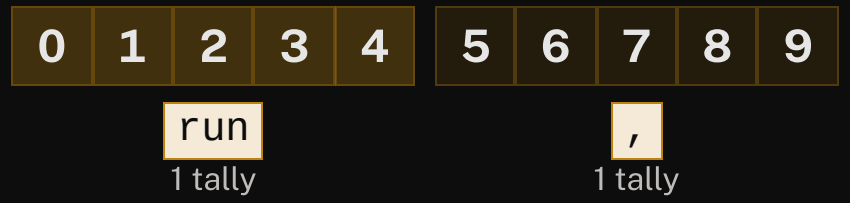

equal tallies → equal chances

Speaker notes

equal tallies, so each option gets an equal share of the ten d10 faces---`run` gets 0-4 and `,` gets 5-9, reading the row's cells left to right. Pause here before rolling.

How the die chooses: `spot`

`spot` → ? roll a d10

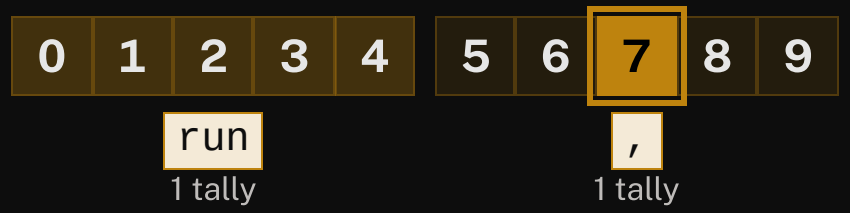

equal tallies → equal chances

rolled 7 → `,`

Speaker notes

now roll --- a 7 lands in the `,` band

Generation: spot → ,

see spot ,

rolled 7 → ,

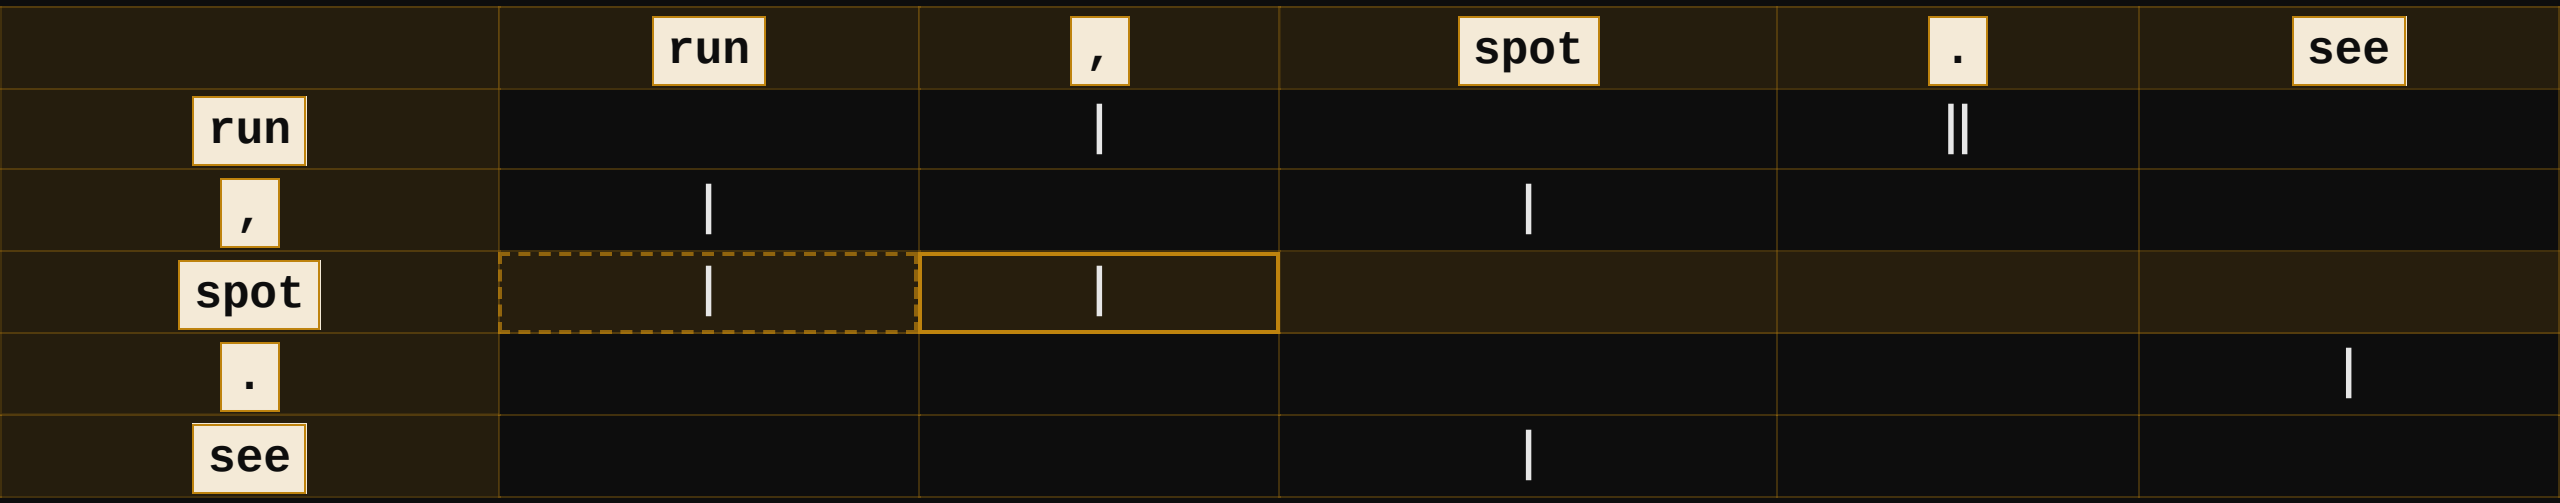

Speaker notes

the 7 landed in the `,` band, so `,` is our next word

Generation: from ,

see spot ,

2 options — roll the die!

	run	,	spot	.	see
run		|		||	
,	|		|		
spot	|	|			
.					|
see			|		

Speaker notes

another two-option row (`run` and `spot`), still equal --- both highlighted, then roll

Generation: , → run

see spot , run

rolled 2 → run

	run	,	spot	.	see
run		|		||	
,	|		|		
spot	|	|			
.					|
see			|		

Speaker notes

rolled a 2, which lands in the `run` band

Generation: from `run`

`see` `spot` `,` **`run`**

2 options — roll the die!

	`run`	`,`	`spot`	`.`	`see`
`run`		|		||	
`,`	|		|		
`spot`	|	|			
`.`					|
`see`			|		

Speaker notes

two options again --- but `.` has two tallies to `,`'s one, so this isn't a 50/50 roll

How the die chooses: `run`

`run` → ? roll a d10

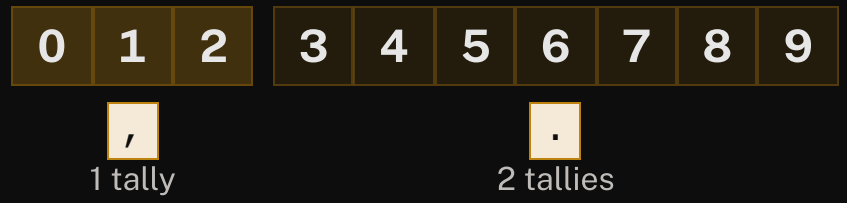

more tallies → more faces → more likely

Speaker notes

`.` has twice the tallies, so it gets more of the faces (3-9) than `,` (0-2). With a ten-sided die that's the closest split rather than exactly 2:1 --- the point is more tallies → more faces → more likely. Pause here before rolling.

How the die chooses: `run`

`run` → ? roll a d10

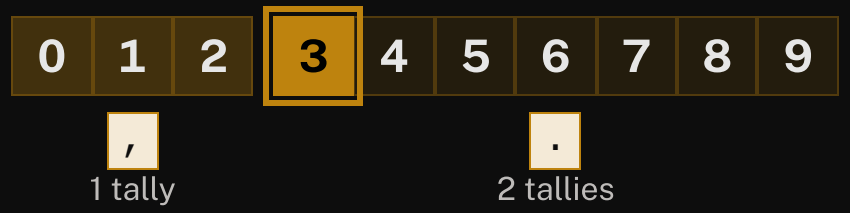

more tallies → more faces → more likely

rolled **3** → `.`

Speaker notes

now roll --- a 3 sits in the `.` band

Generation: run → .

see spot , run .

rolled 3 → .

	run	,	spot	.	see
run		|		||	
,	|		|		
spot	|	|			
.					|
see			|		

Speaker notes

the 3 sits in the `.` band, so we pick `.`

Generation: from `.`

`see` `spot` `,` `run` `.` `see`

one option — no roll needed

	`run`	`,`	`spot`	`.`	`see`
`run`		|		||	
`,`	|		|		
`spot`	|	|			
`.`					|
`see`			|		

Speaker notes

only one option (`see`) --- no roll. Note we keep rolling straight past the full stop; the model has no idea a sentence just ended.

Generation: back to `see`

`see` `spot` `,` `run` `.` `see`

one option — no roll needed

	`run`	`,`	`spot`	`.`	`see`
`run`		|		||	
`,`	|		|		
`spot`	|	|			
`.`					|
`see`			|		

Speaker notes

and so on---we've looped back to where we started. Keep going for as long as you like; you decide when to stop.

Generated text

“see spot, run. see”

a new sentence — not in the training data!

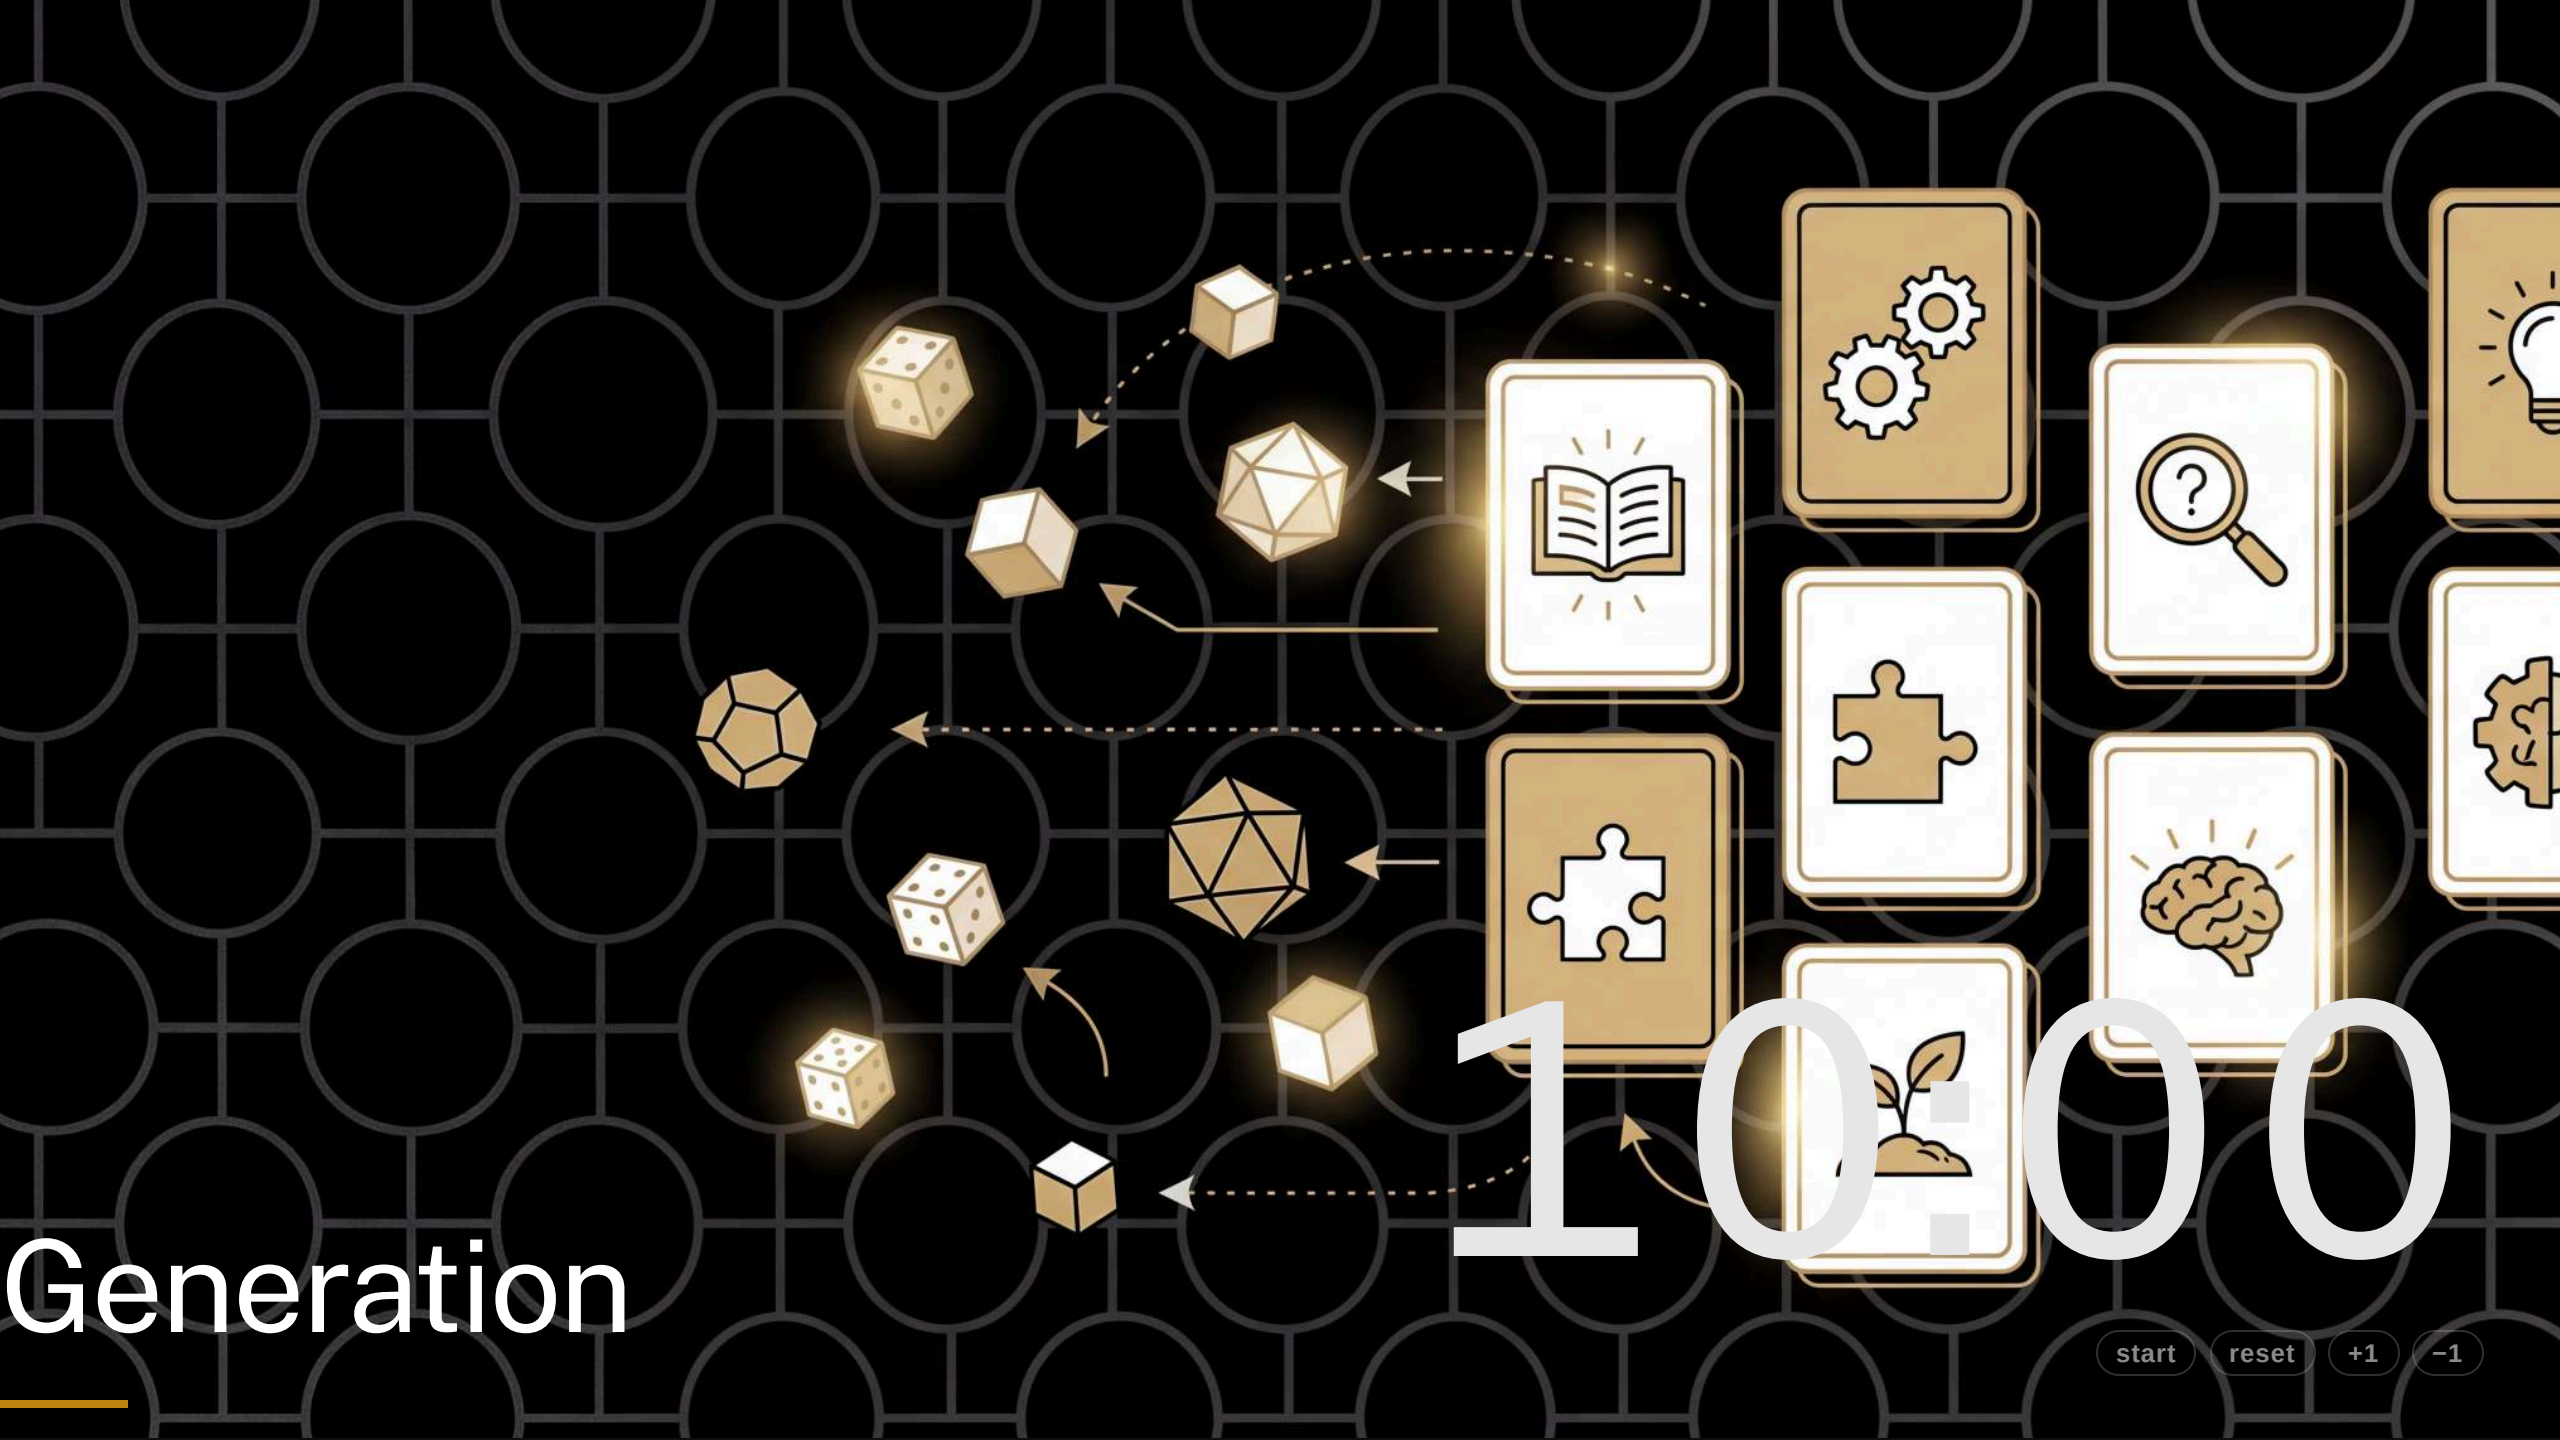

Generation

Speaker notes

Groups generate from their grid, rolling dice on each choice. 10 min, circulate. Where a row has only one option, no roll needed.

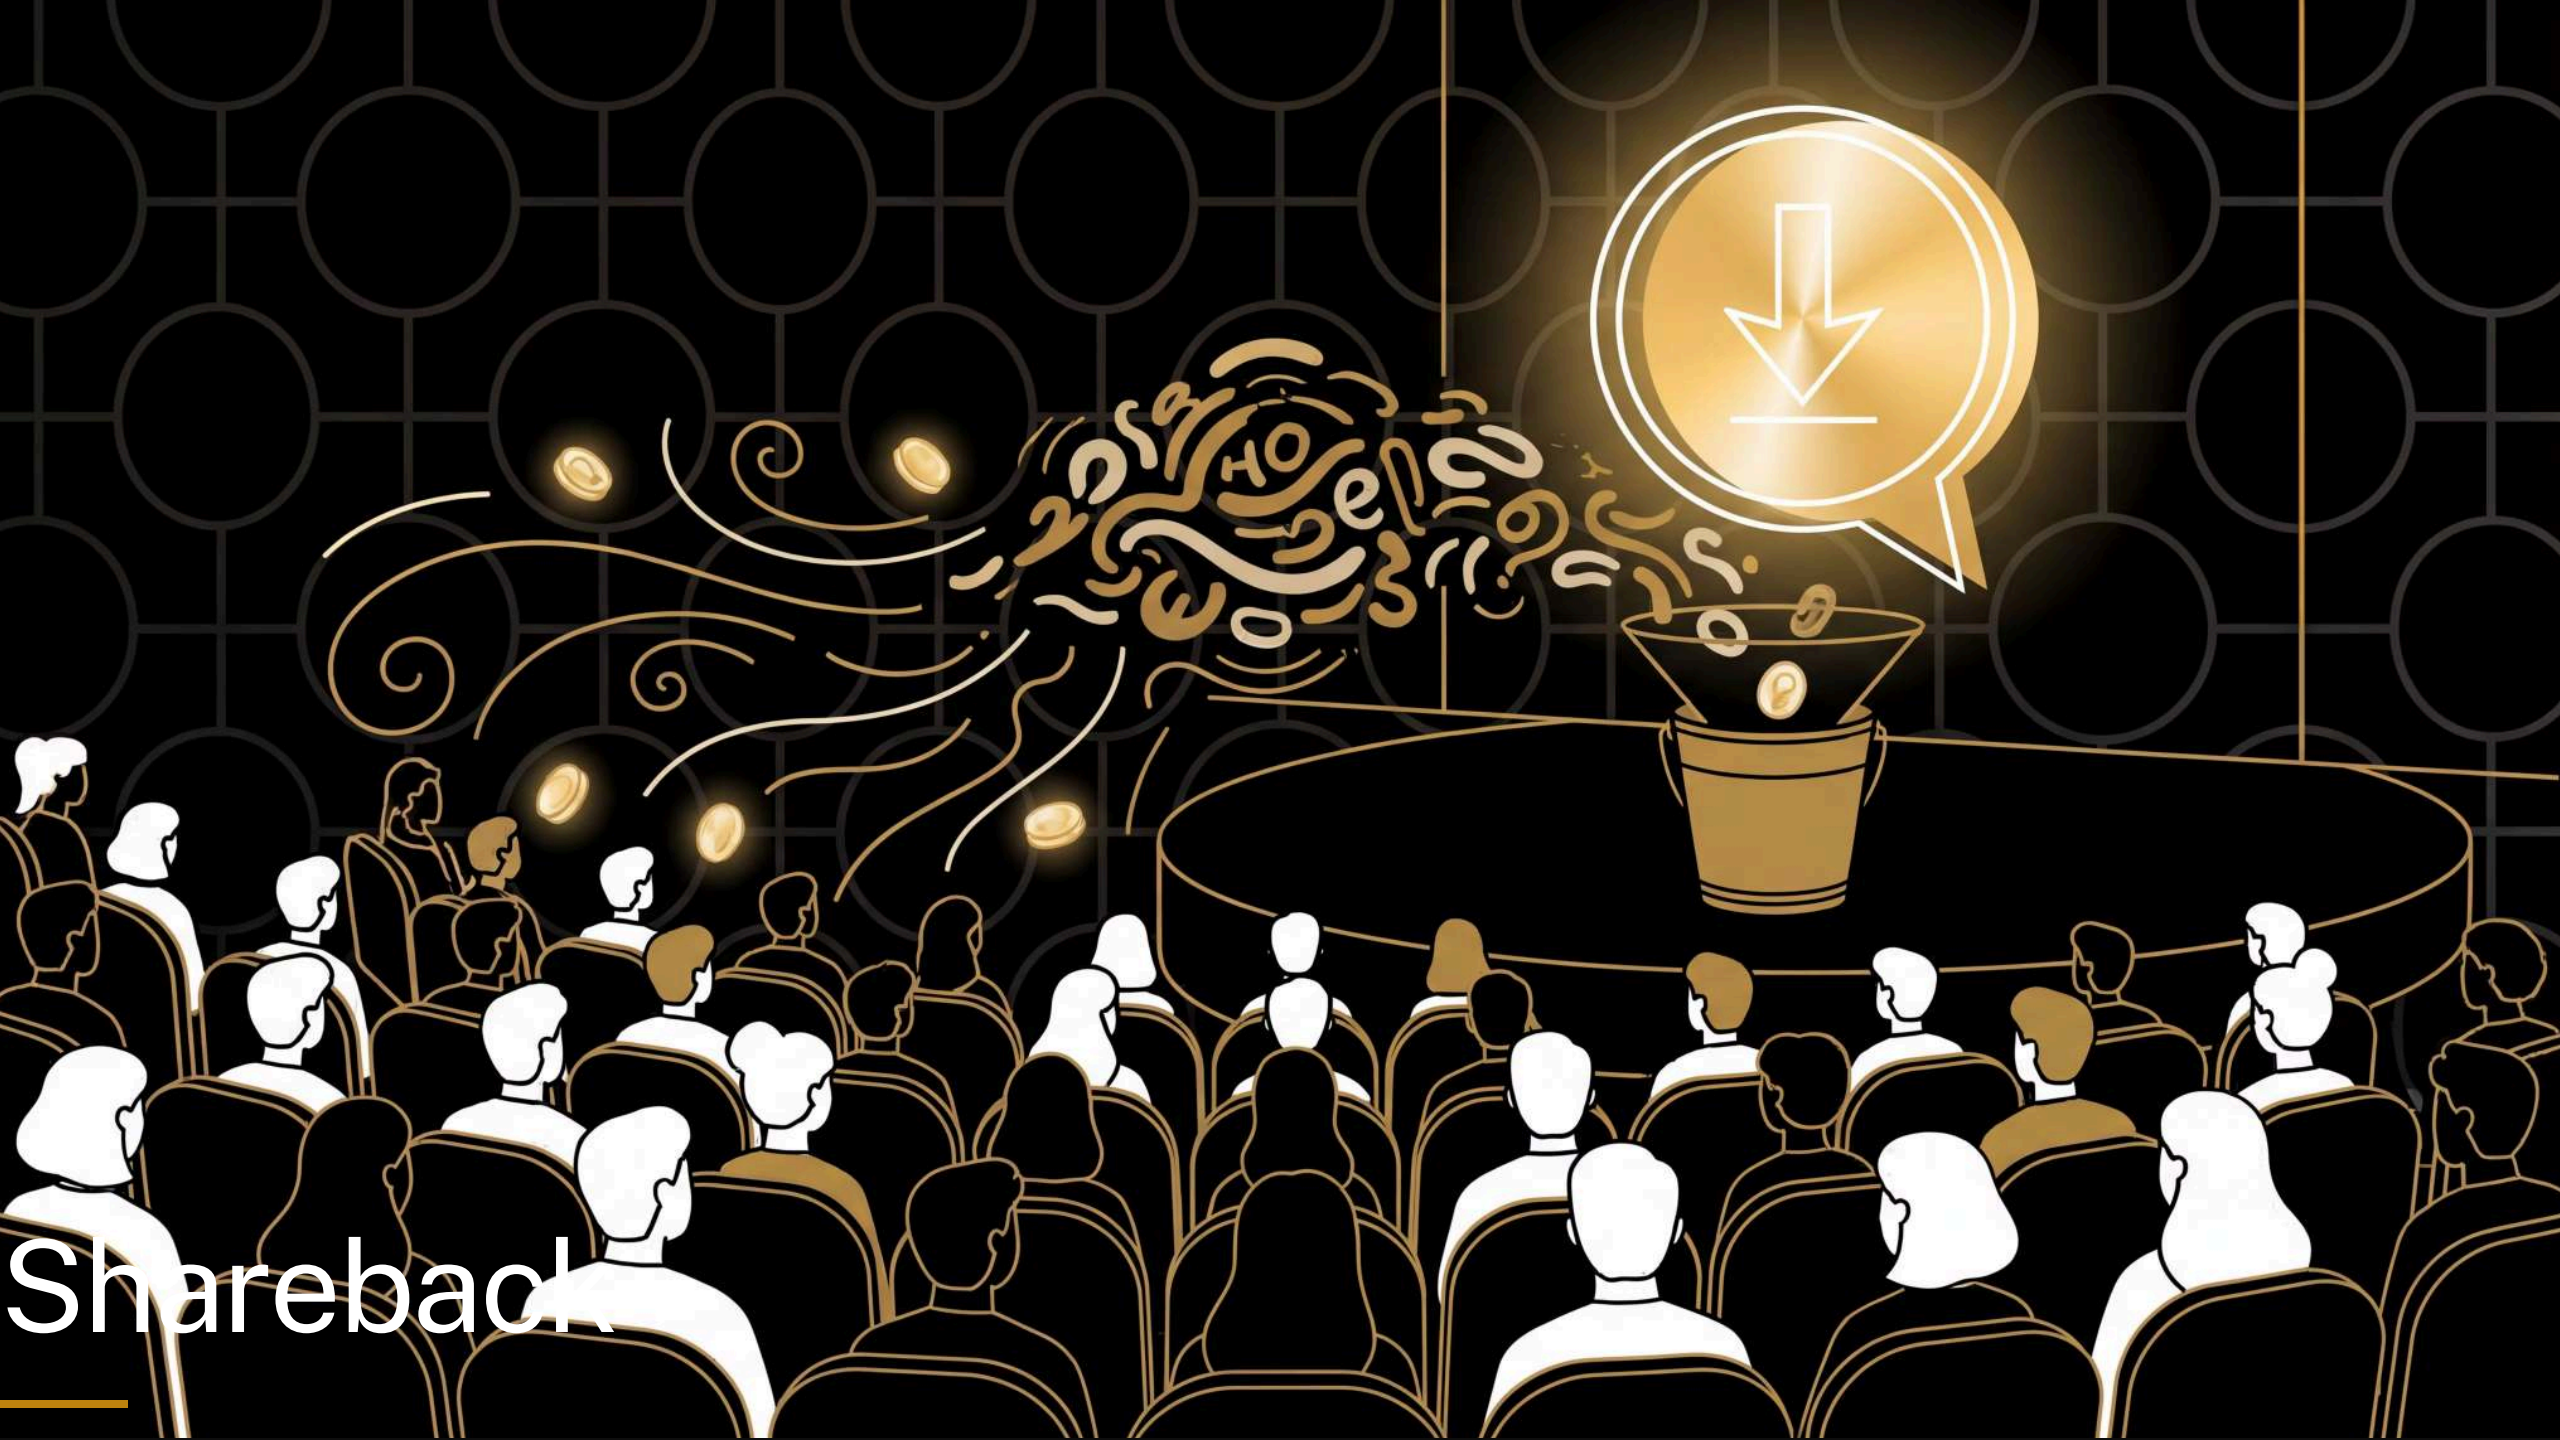

Shareback

Speaker notes

Take a couple of groups' generated sentences; celebrate the weird ones. ~5 min.

The *language* of language models

- prompt
- completion
- hallucination

Speaker notes

Everyday sense first, then ours: to prompt is to nudge an actor who's forgotten a line---here the prompt is the starting text you give the model. A completion is finishing something off---here it's the text the model generates back, and it doesn't know when to stop. A hallucination is seeing something that isn't there, a malfunction---but the sentence you just generated was never in the text either, and it came out of the same tallies and the same dice as everything else. There's no separate mode the model slips into; we only call it hallucination when we didn't want the answer.

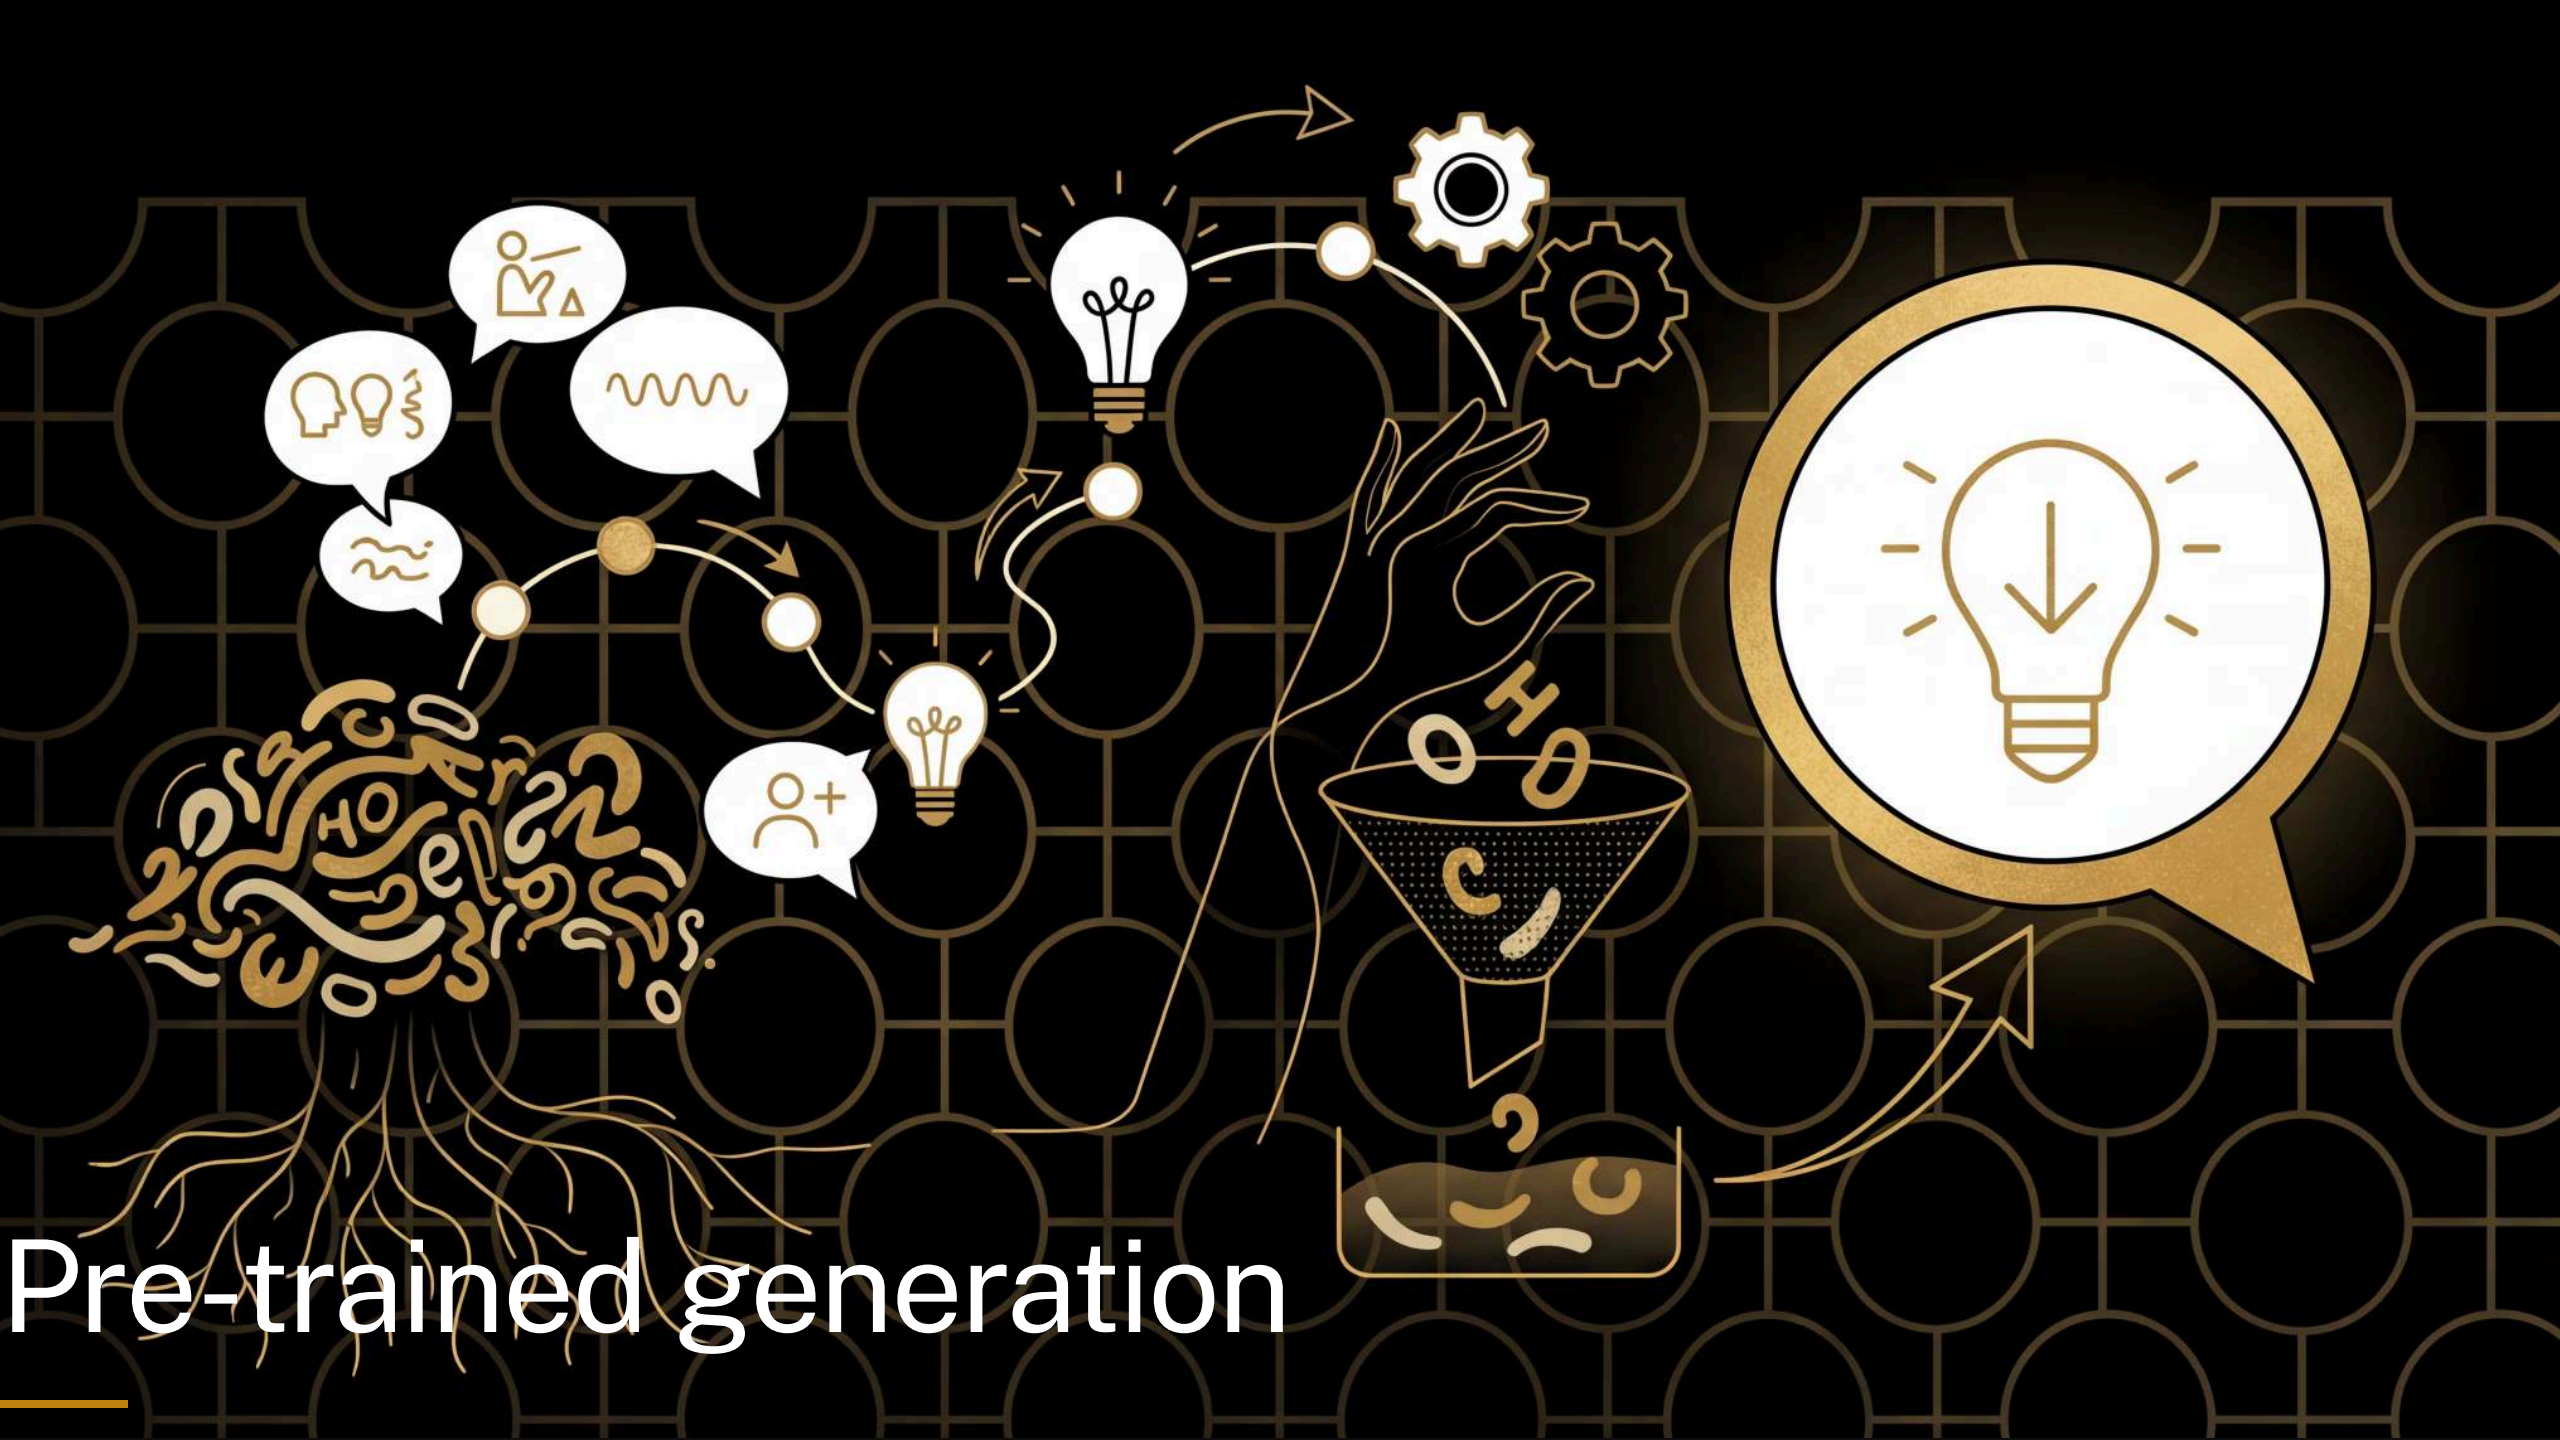

Pre-trained generation

The recipe

use the booklet to generate new text, rolling one or more d10s to choose each next word

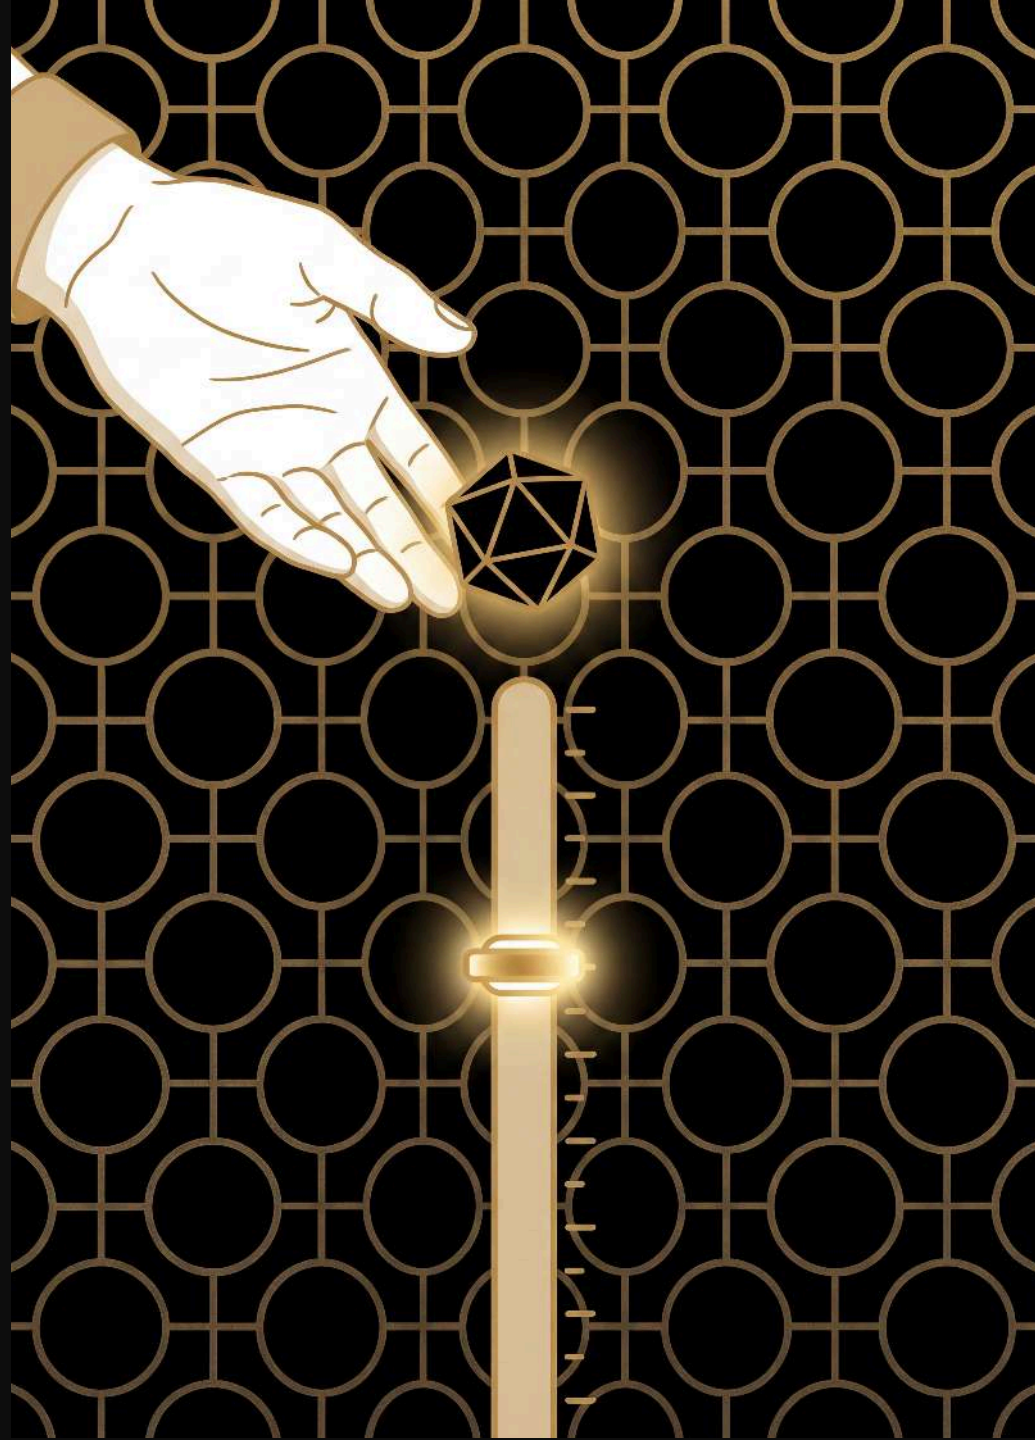

Speaker notes

1. **choose** a starting word --- any bold word in the booklet
2. **look up** that word's entry to see possible next words and their thresholds
3. **roll** your d10(s) and scan the thresholds to find the next word
4. **write down** the chosen word and make it your new starting word
5. **repeat** from step 2

Pre-trained: start with `the`

`the` `cat`

the ♦♦ **49**|cat **79**|dog **99**|hat **rolled 27**

cat ♦ **5**|sat **9**|ran

sat .

. the

ran .

dog ♦ **6**|sat **9**|ran

hat ♦ **4**|sat **9**|ran

Speaker notes

a bigger model — `the` is so common it needs **two** dice (roll two d10s for 00 to 99)

Pre-trained: from `cat`

`the` `cat` `sat`

the ♦♦ **49**|cat **79**|dog **99**|hat

cat ♦ **5**|sat **9**|ran **rolled 4**

sat .

. the

ran .

dog ♦ **6**|sat **9**|ran

hat ♦ **4**|sat **9**|ran

Speaker notes

`cat` is rarer than `the` — just **one** die, roll a d10

Pre-trained: from `sat`

`the` `cat` `sat` `.`

the ♦♦ 49|cat 79|dog 99|hat

cat ♦ 5|sat 9|ran

sat .

. the

ran .

dog ♦ 6|sat 9|ran

hat ♦ 4|sat 9|ran

Speaker notes

only one option — no dice roll needed

Pre-trained: from .

the cat sat . the

the ♦♦ 49|cat 79|dog 99|hat

cat ♦ 5|sat 9|ran

sat .

. the

ran .

dog ♦ 6|sat 9|ran

hat ♦ 4|sat 9|ran

Speaker notes

only one option — no dice roll needed

Pre-trained: from `the` again

`the` `cat` `sat` `.` `the` `dog`

the ♦♦ 49|cat 79|dog 99|hat rolled 63

cat ♦ 5|sat 9|ran

sat .

. the

ran .

dog ♦ 6|sat 9|ran

hat ♦ 4|sat 9|ran

Speaker notes

two dice again — 63 is past `cat`'s 49, so we land on `dog`

Pre-trained: from dog

`the` `cat` `sat` `.` `the` `dog`

the ♦♦ **49**|cat **79**|dog **99**|hat

cat ♦ **5**|sat **9**|ran

sat .

. the

ran .

dog ♦ **6**|sat **9**|ran

hat ♦ **4**|sat **9**|ran

Speaker notes

and so on...

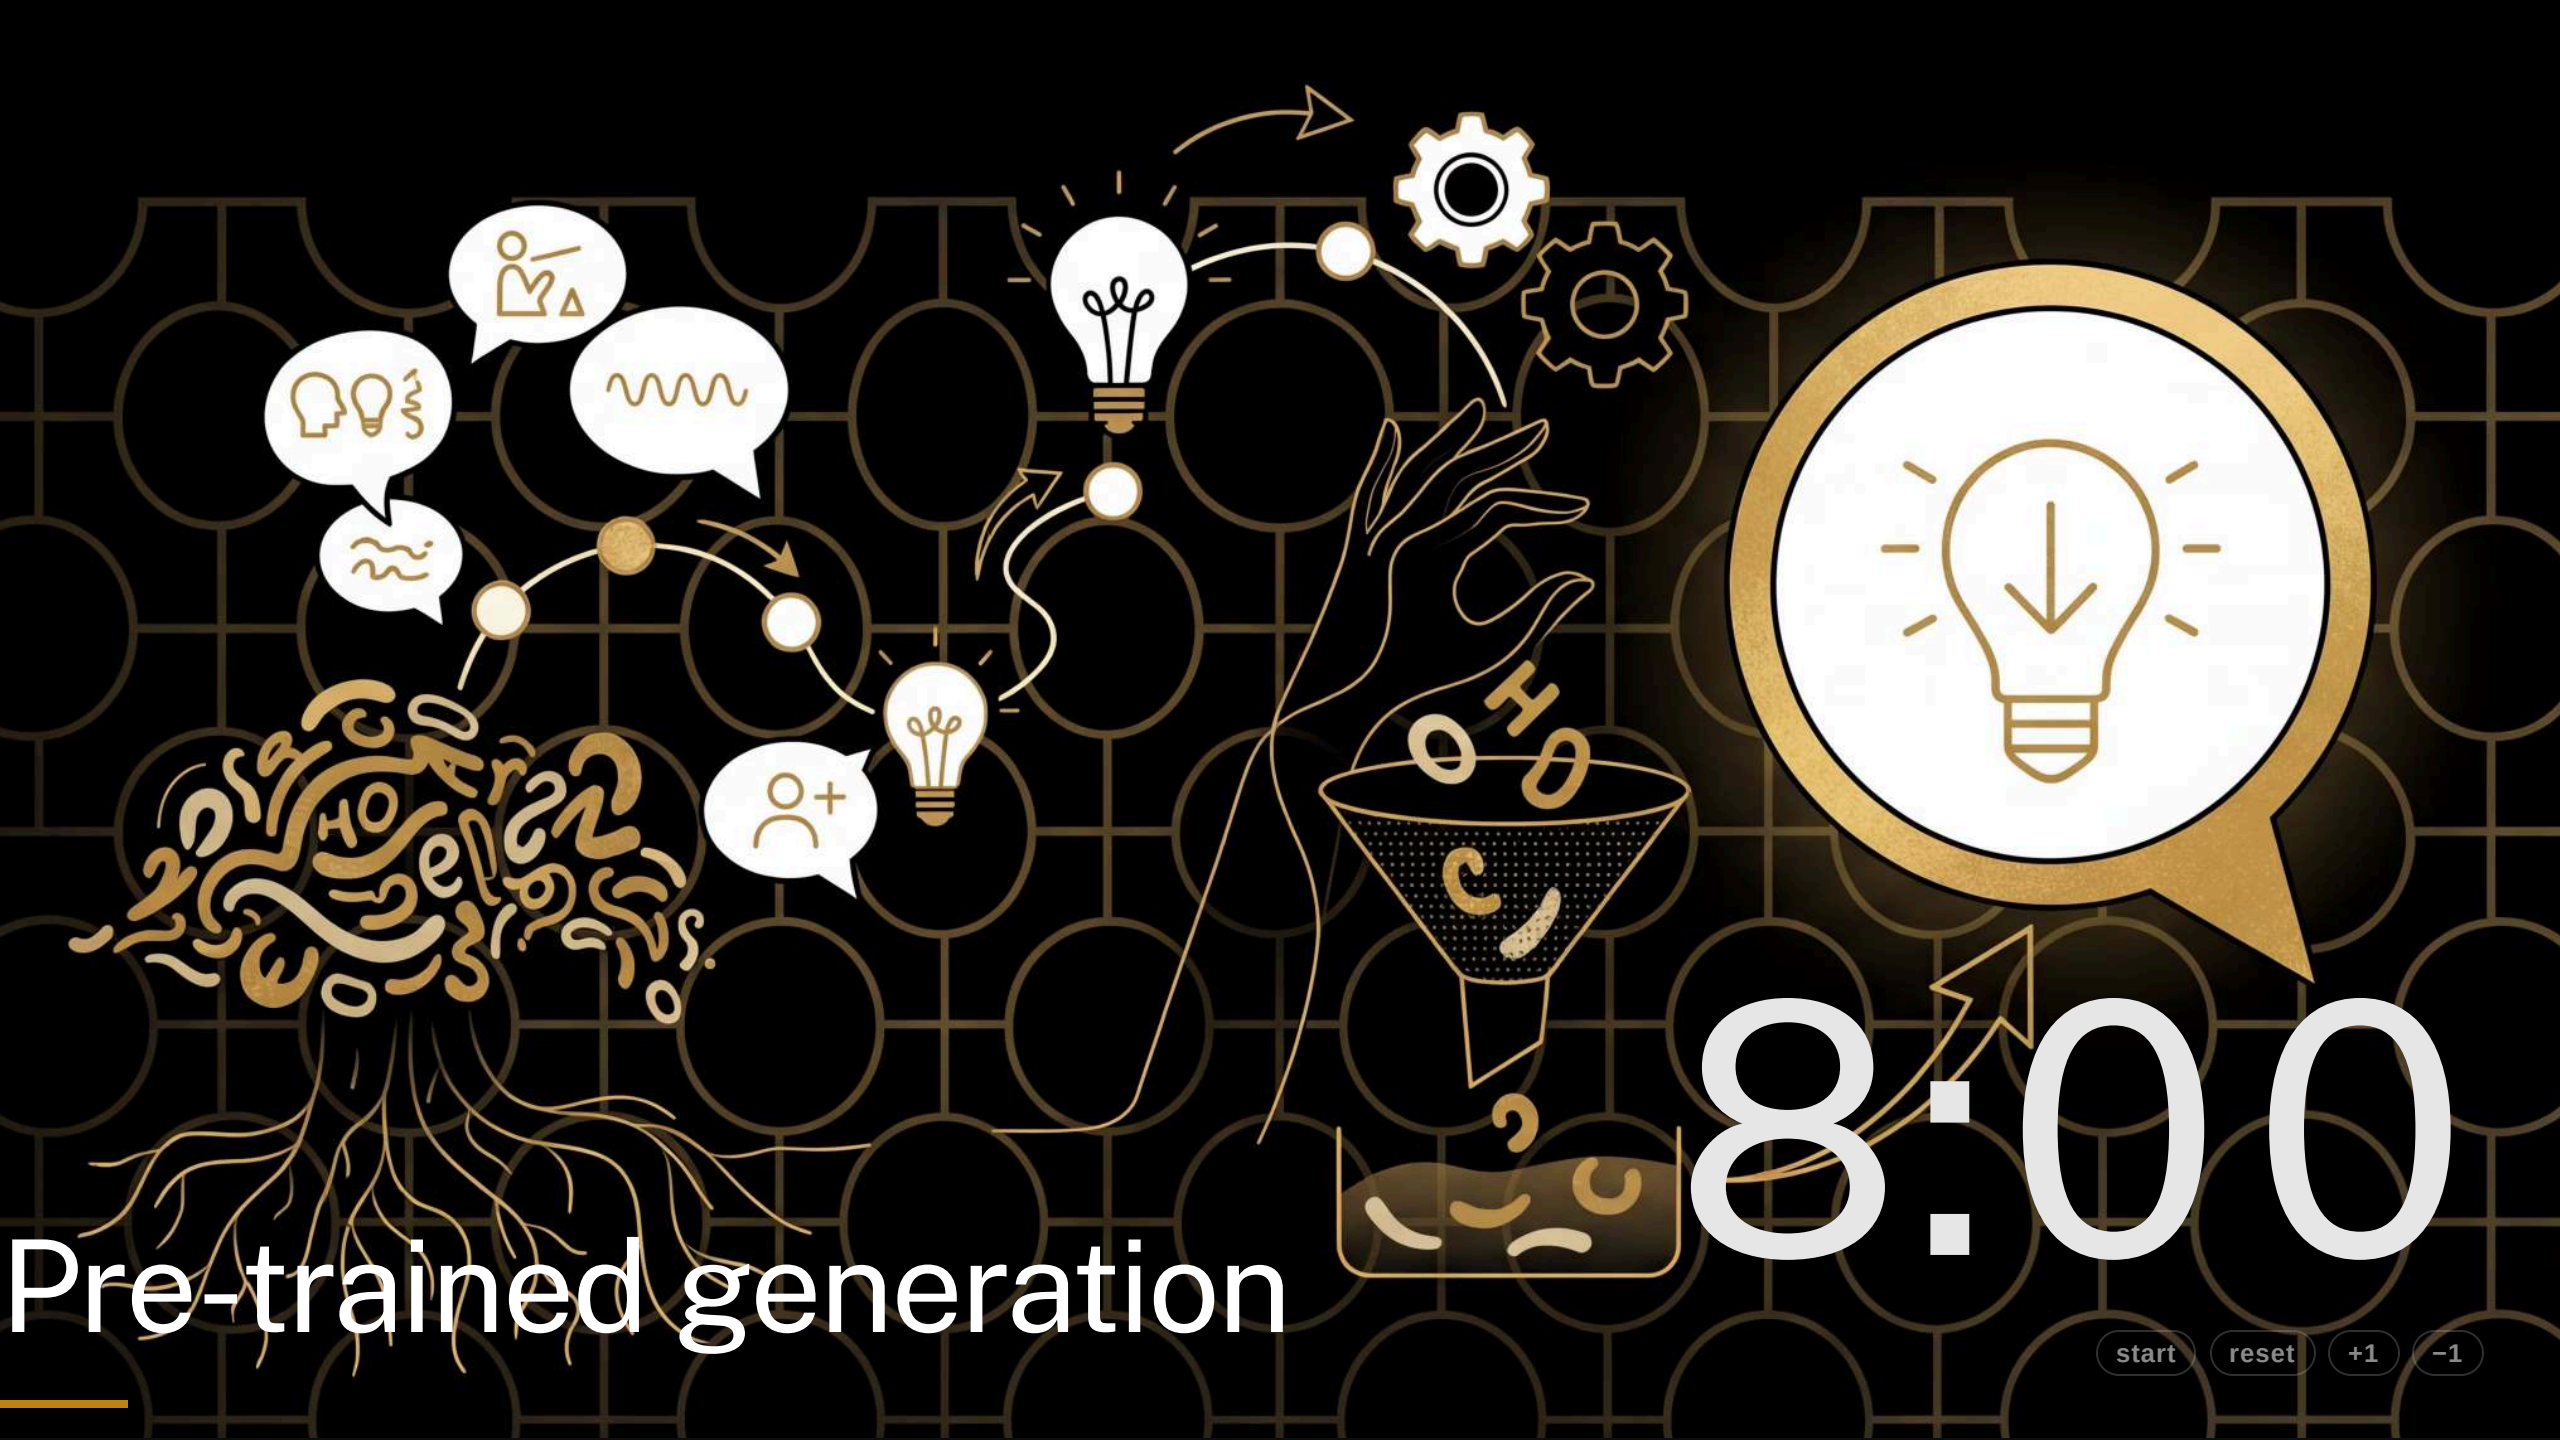

Pre-trained generation

8:00

start reset +1 −1

Speaker notes

Same as before but from the booklet — a bigger model, one or more d10s. 8 min, circulate.

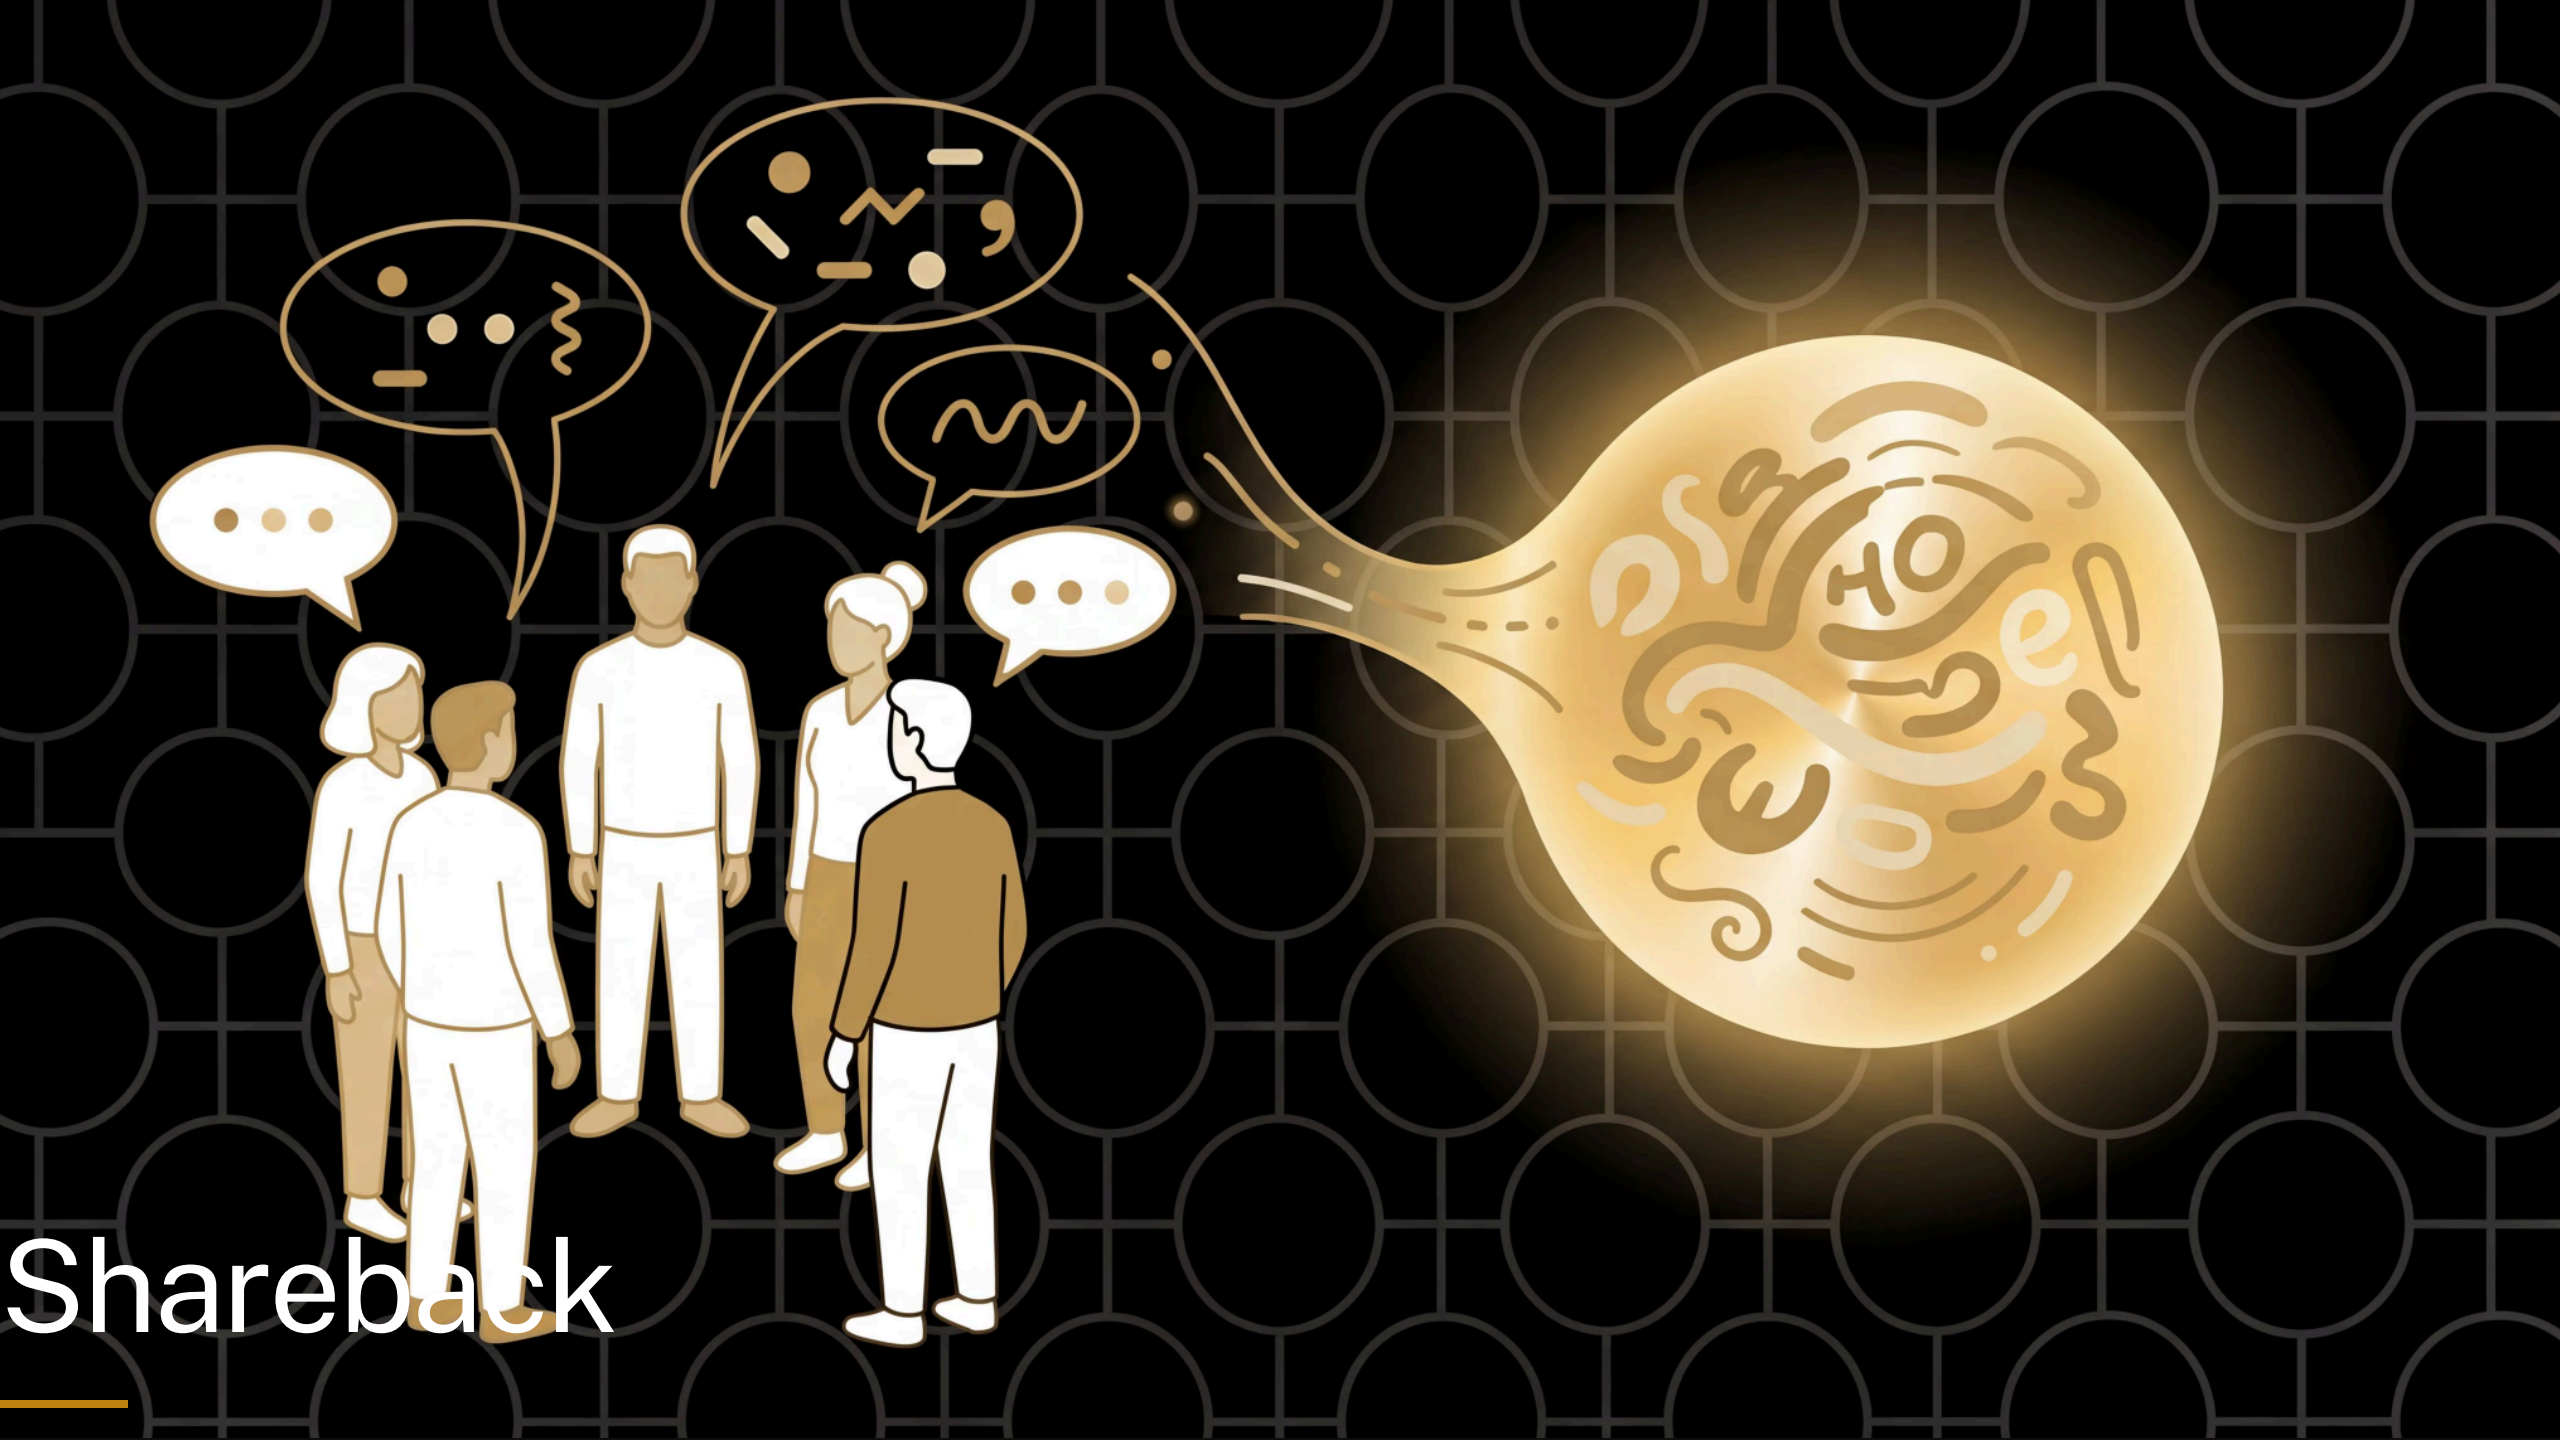

Shareback

Speaker notes

Compare with the hand-built model: a bigger training text gives more varied, more natural output. Invite ~5 groups to share back rather than the whole room. ~5 min.

The *language* of language models

- foundation model

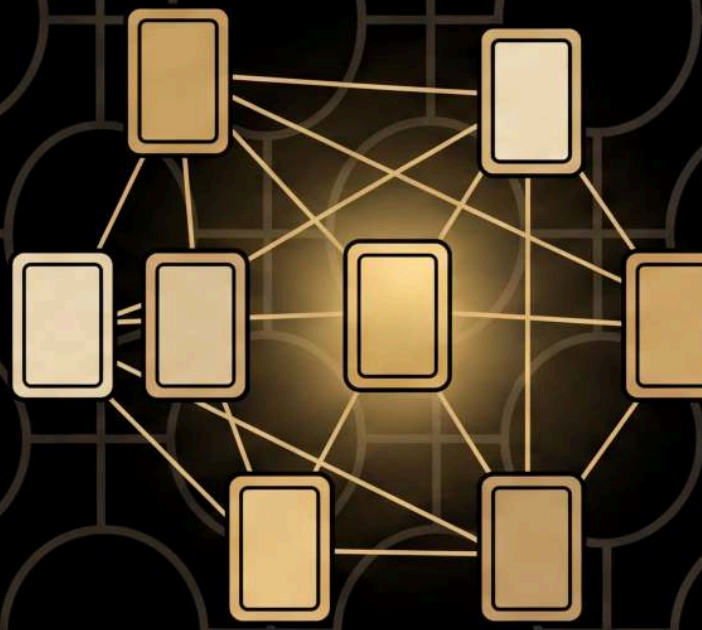

Speaker notes

Everyday sense first, then ours: a foundation is what a house sits on, or a charity --- here a foundation model is the general-purpose thing you get from one enormous training run, which everything else is built on. Worth saying that "pre-training" sounds like a warm-up but is the main event; the booklet they've just used is the output of one.

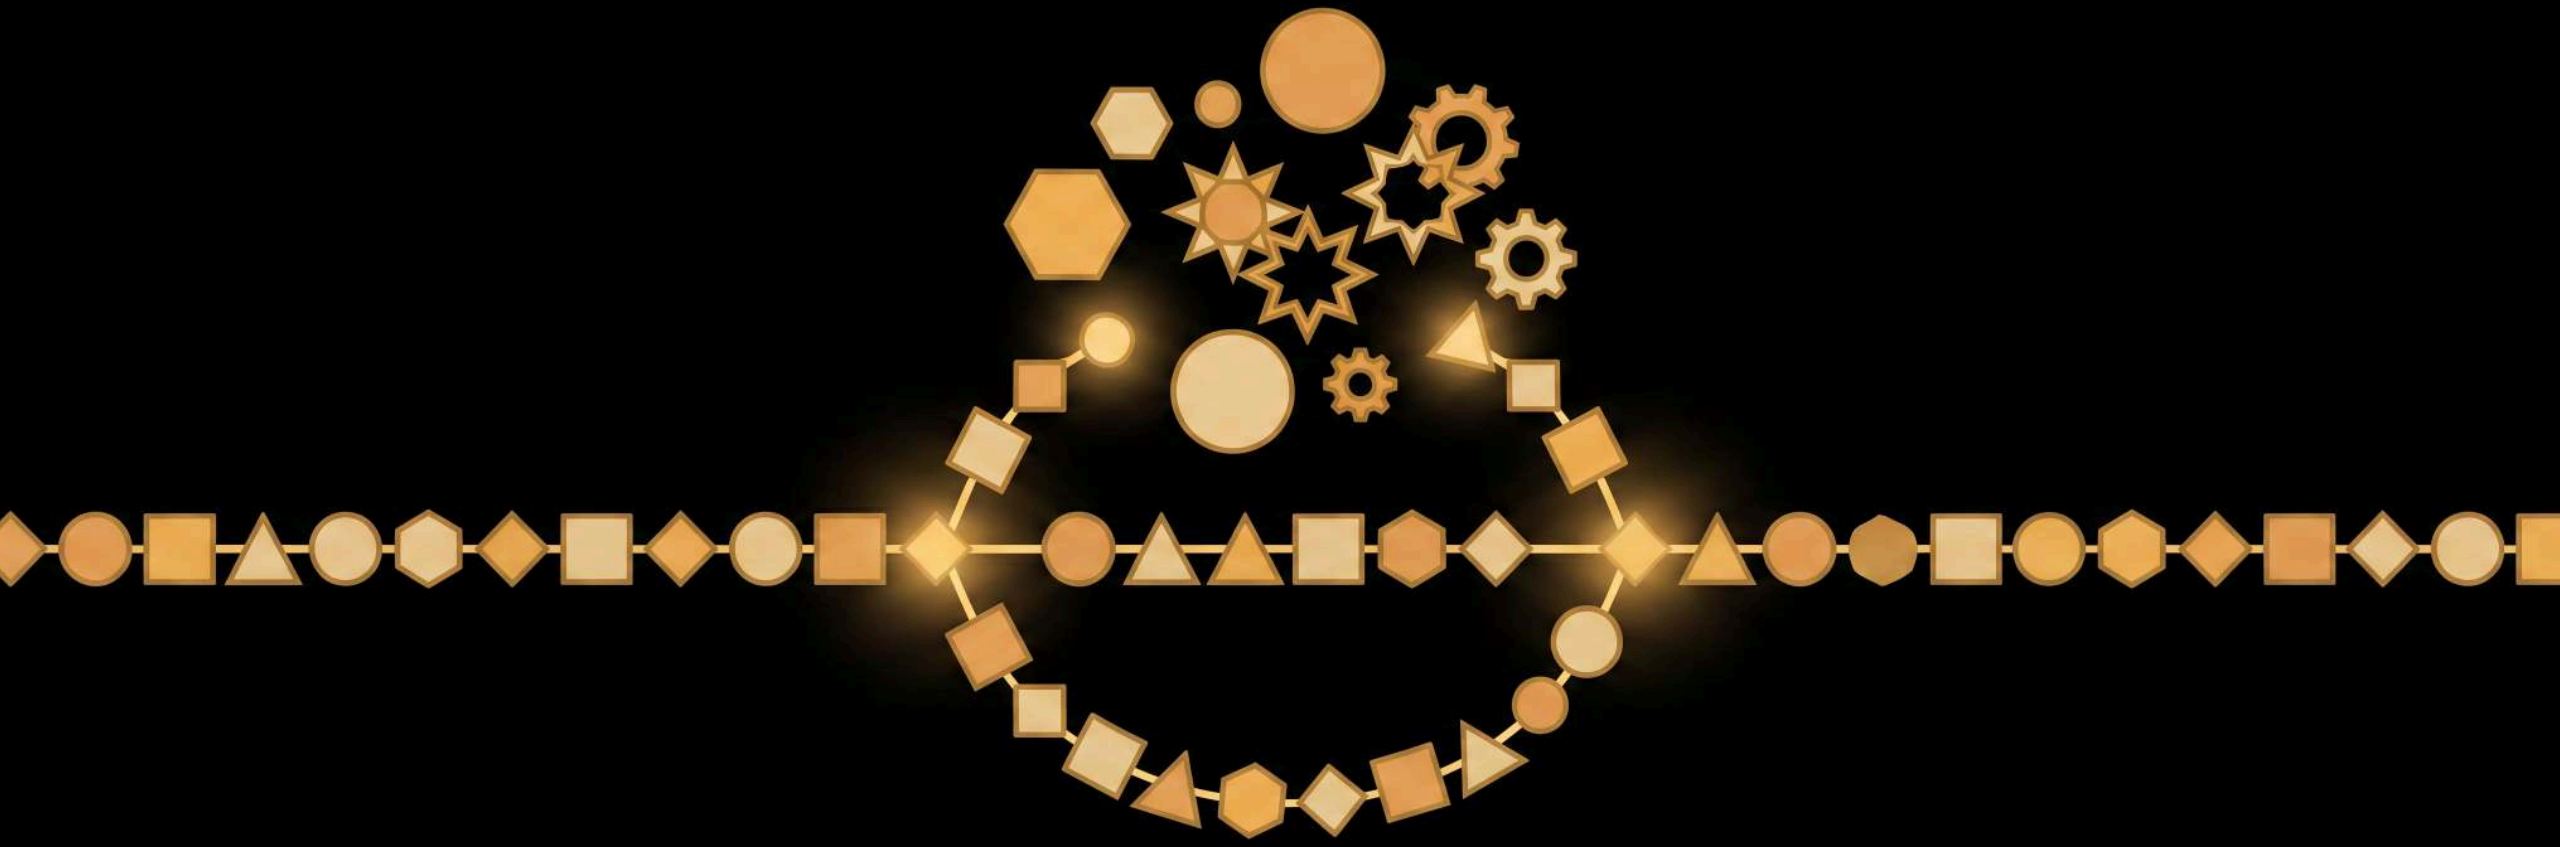

Agentic AI

Speaker notes

Section beat: an agent is a model that can pause, call a tool, and use the result. We bolt that onto your model.

What makes a model an “agent”?

an agent is a model that can **call tools**: pause generation, get information from “outside” the model, and continue

the loop: generate → trigger token → call tool → write the result → keep going

Speaker notes

The beat: the model pauses, gets information from outside itself, then continues. The loop is the thing.

The recipe

generate from your model as before, but every punctuation token triggers a **tool call**: text `What comes next? "[sentence so far]..."` to 3 friends, and the whole first reply goes into your text — followed by the punctuation you rolled

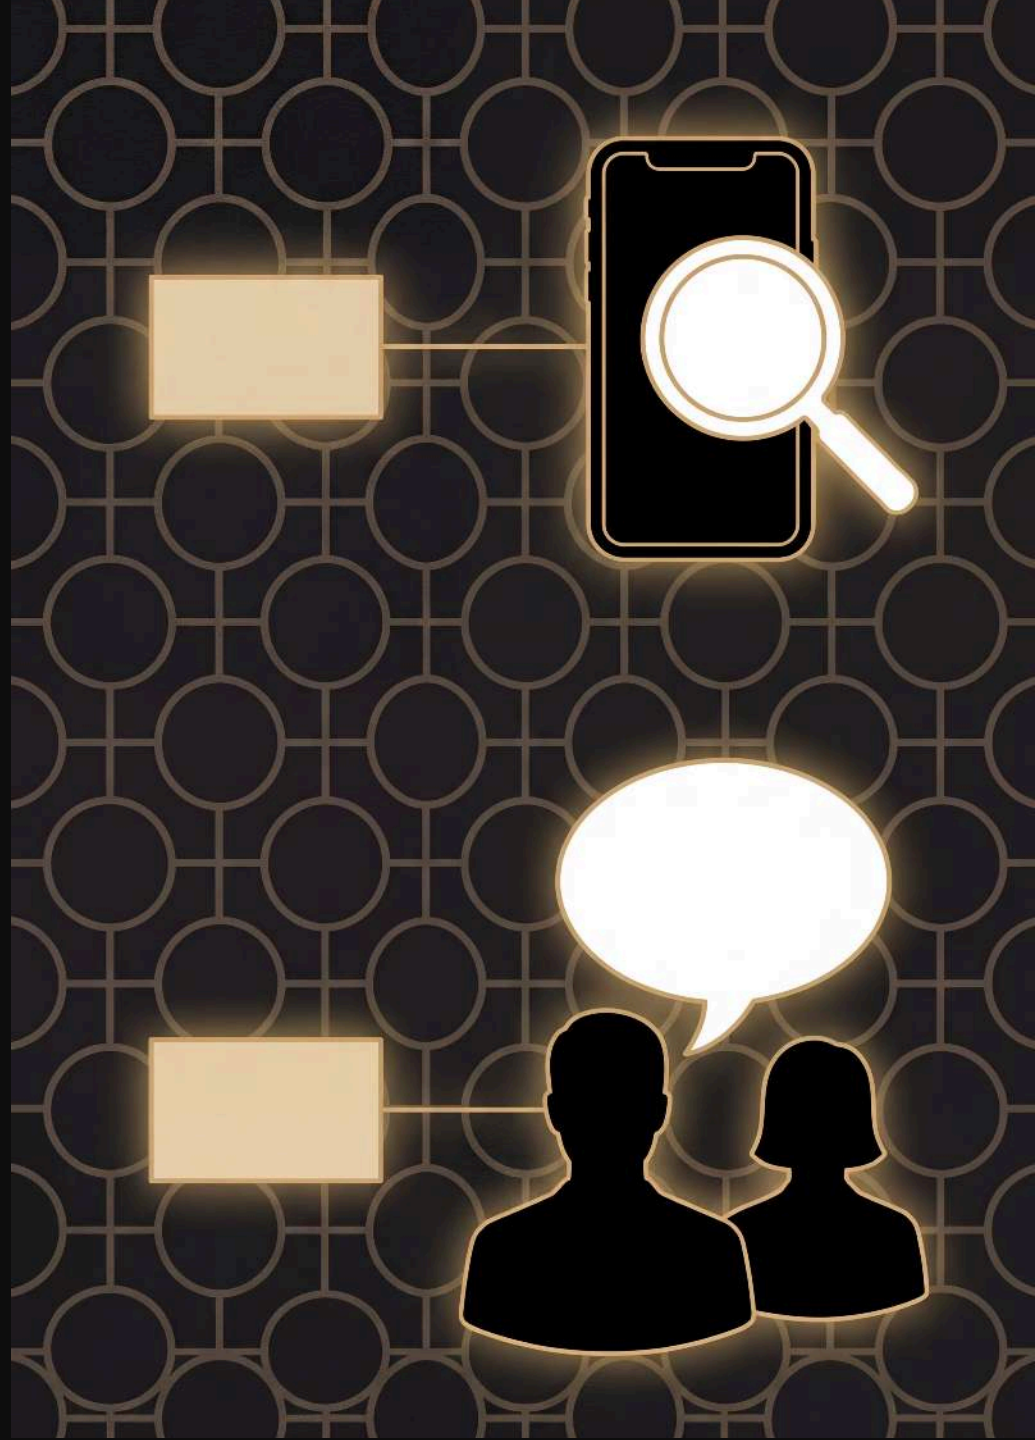

Speaker notes

1. **generate** from your model as before, rolling for each next word
2. when you land on a **punctuation token** (like `.` or `,`), pause --- that's a tool call
3. text **What comes next? "[sentence so far]..."** (everything since the last full stop) to 3 friends
4. **write down** the whole first reply, then the punctuation token you rolled
5. **continue** generating from the punctuation token

Worked example

your text so far is “the cat sat”, and the next dice roll gives you `.` as the next token — pause, that’s a tool call

text `What comes next? "the cat sat..."` to 3 friends; the first reply back might be “down by the river”

write `down by the river`, then the `.` you rolled anyway, and continue generating from `.`

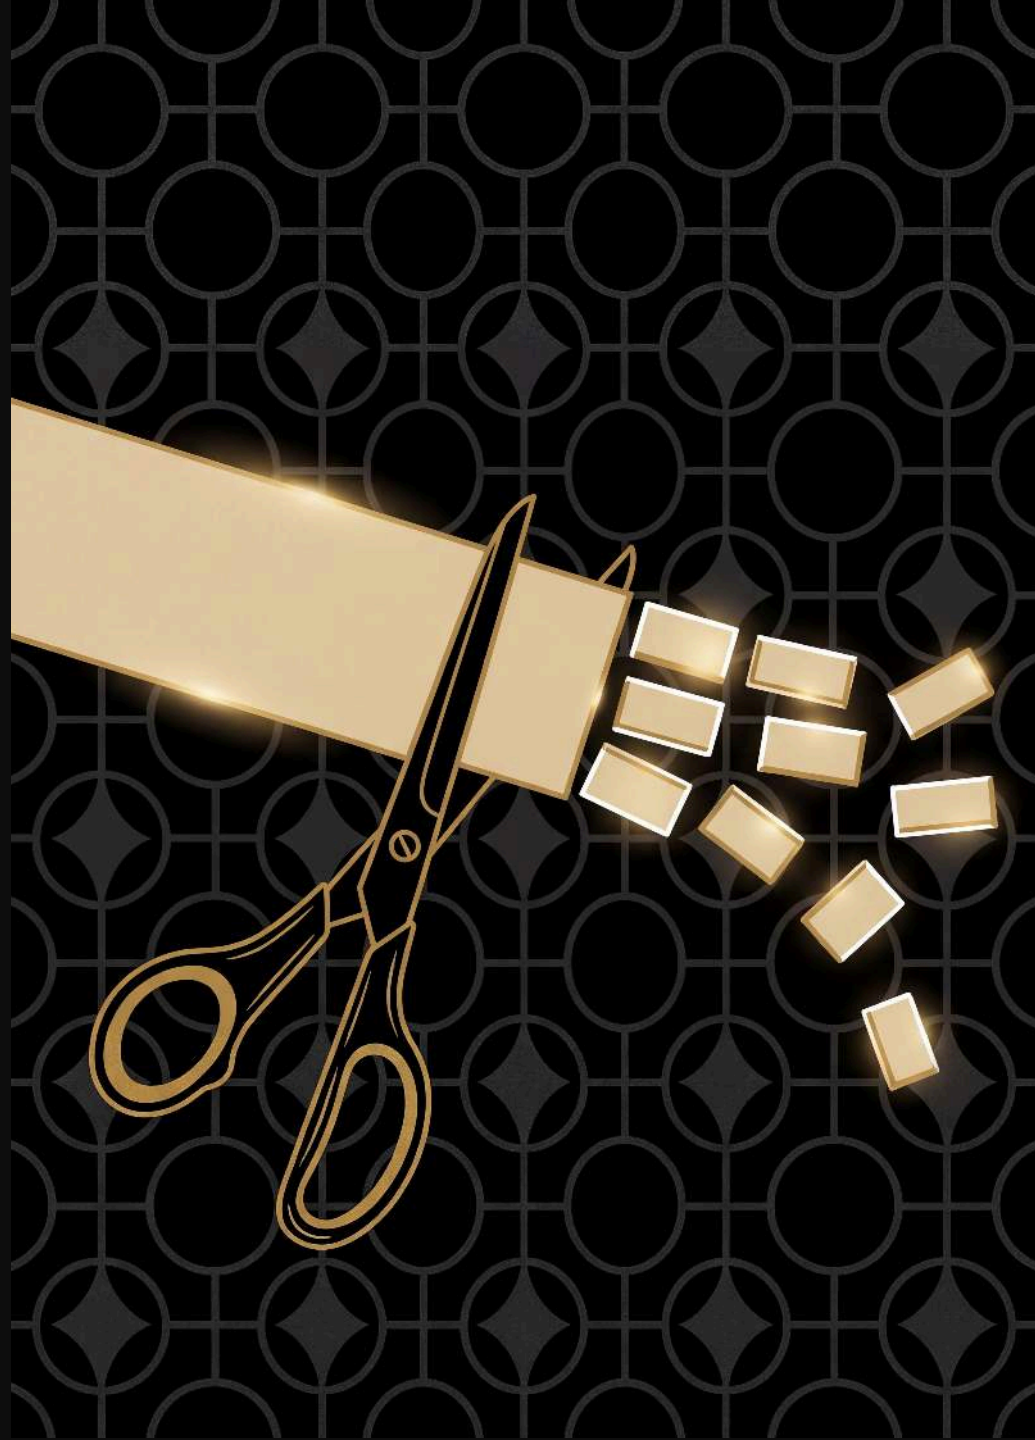

Speaker notes

Walk it slowly: the punctuation is the trigger, the completion prompt plus the sentence so far is the tool input, and the whole first reply is the tool result. The prompt gives friends enough context to produce a useful continuation rather than just react to a mysterious fragment. The punctuation still gets written --- the tool result slots in just before it. If asked why we don't continue from the friend's reply: those words probably aren't in your model's vocabulary, so resuming from the punctuation token is the cleanest way to keep generating.

You will need

your model, dice, pen and paper (as before)

a phone

3 friends who text back quickly

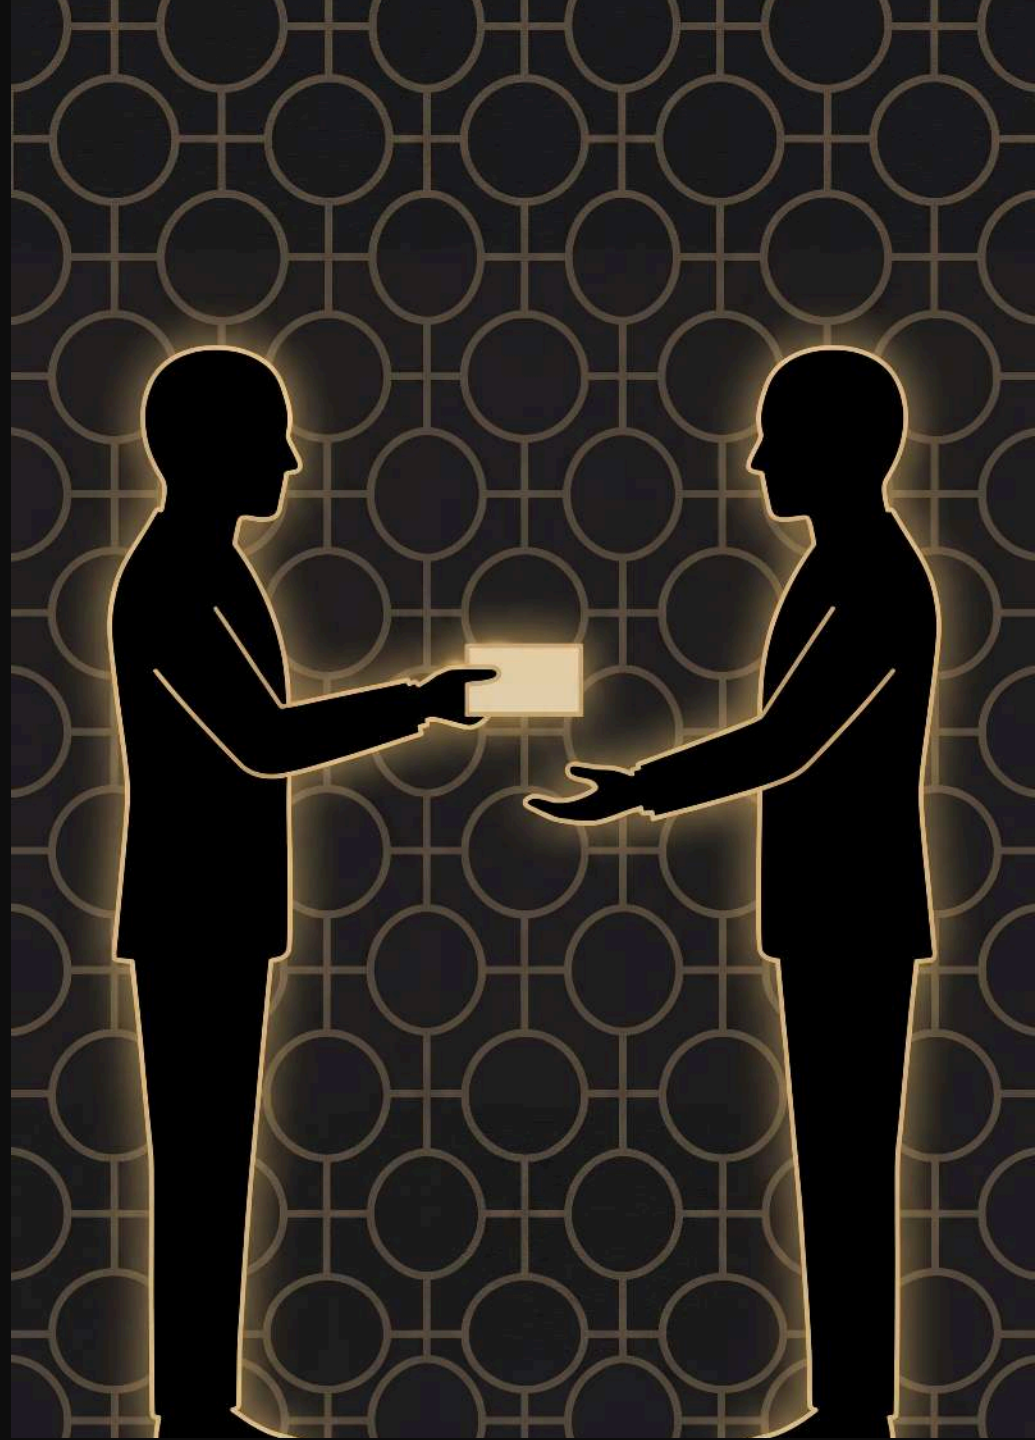

Speaker notes

Texting three friends is the retry logic --- first reply wins. If no-one replies before the next punctuation token, queue it up and back-fill. ~2 min set-up.

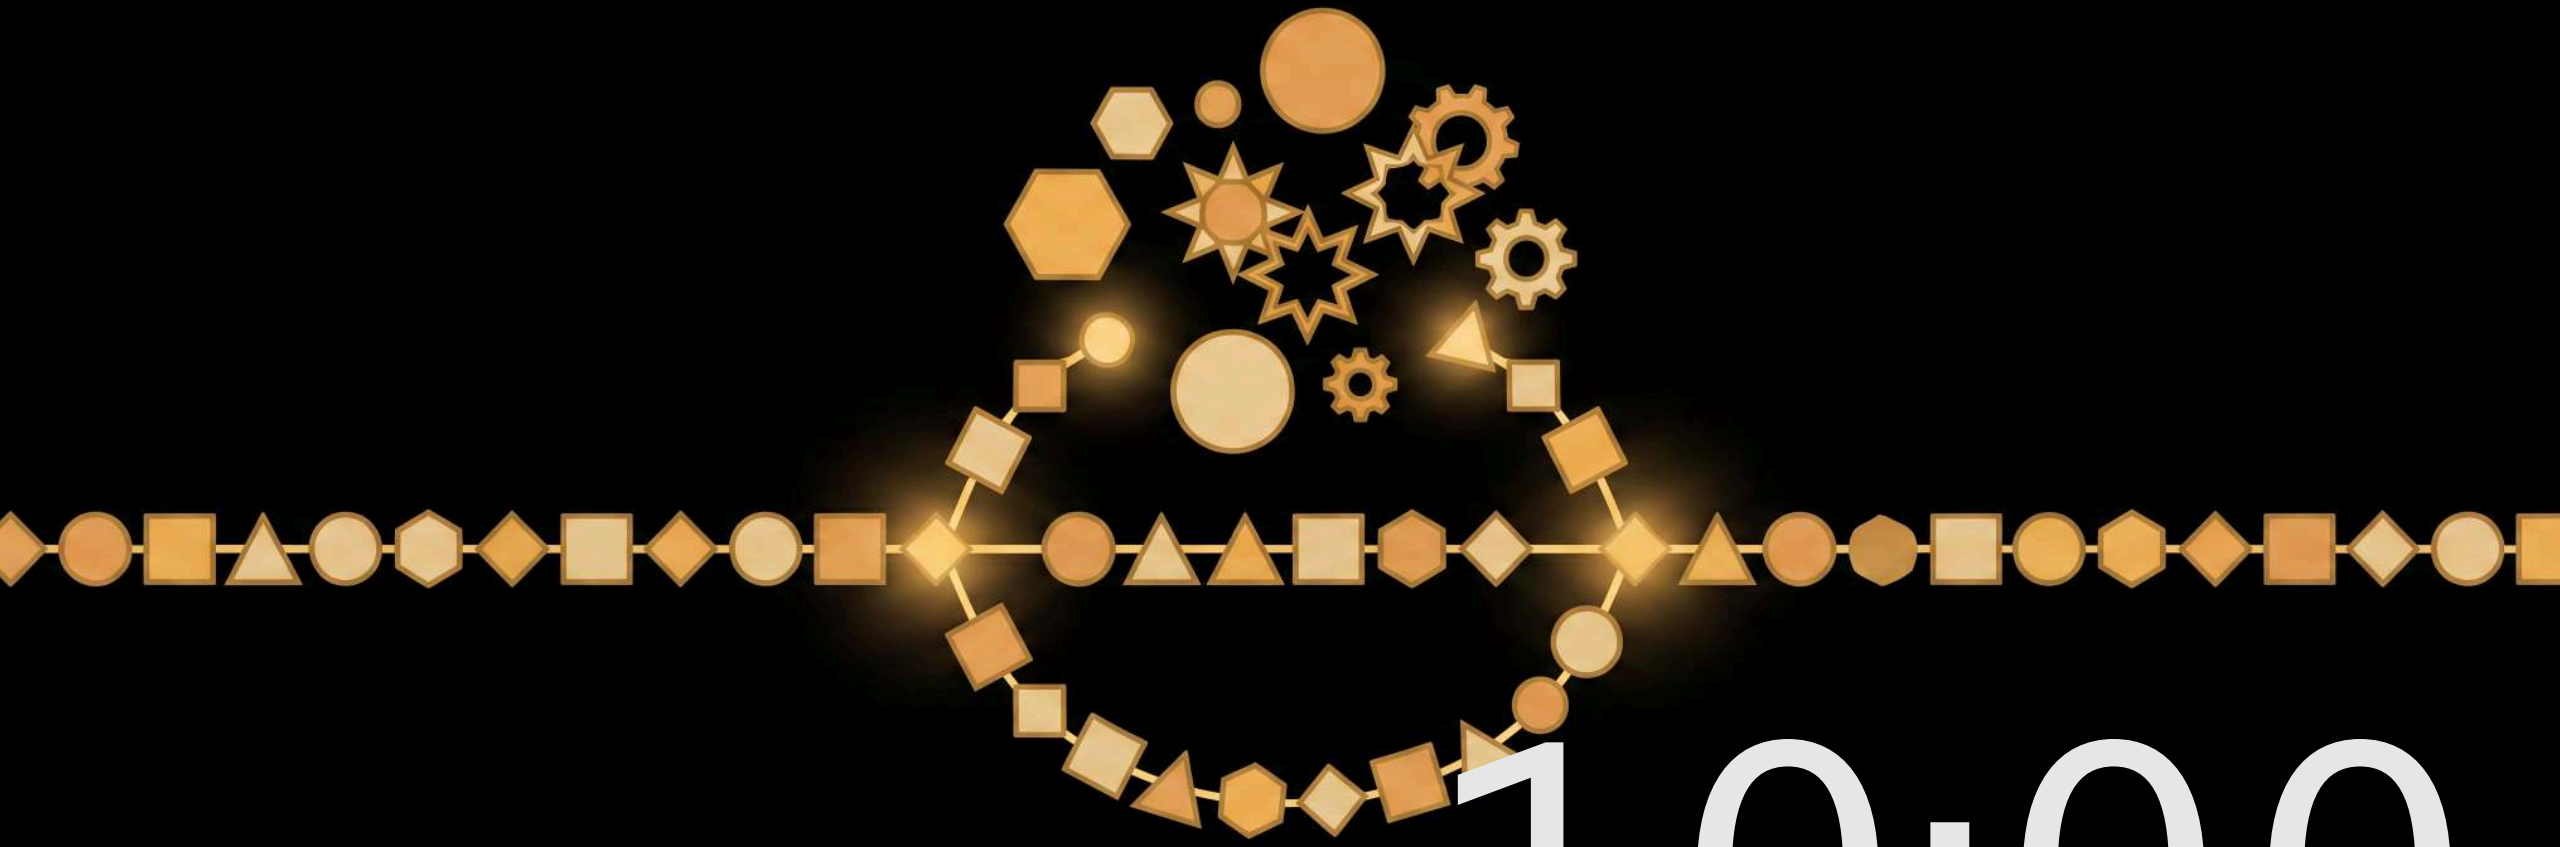

10:00

start reset +1 −1

Agentic AI

Speaker notes

Generate from your model --- on punctuation, text `What comes next? "[sentence so far]..."` to 3 friends and splice in the whole first reply. 10 min --- texting latency is real, so keep groups moving; if a reply hasn't landed by the next punctuation token, back-fill it later. Circulate.

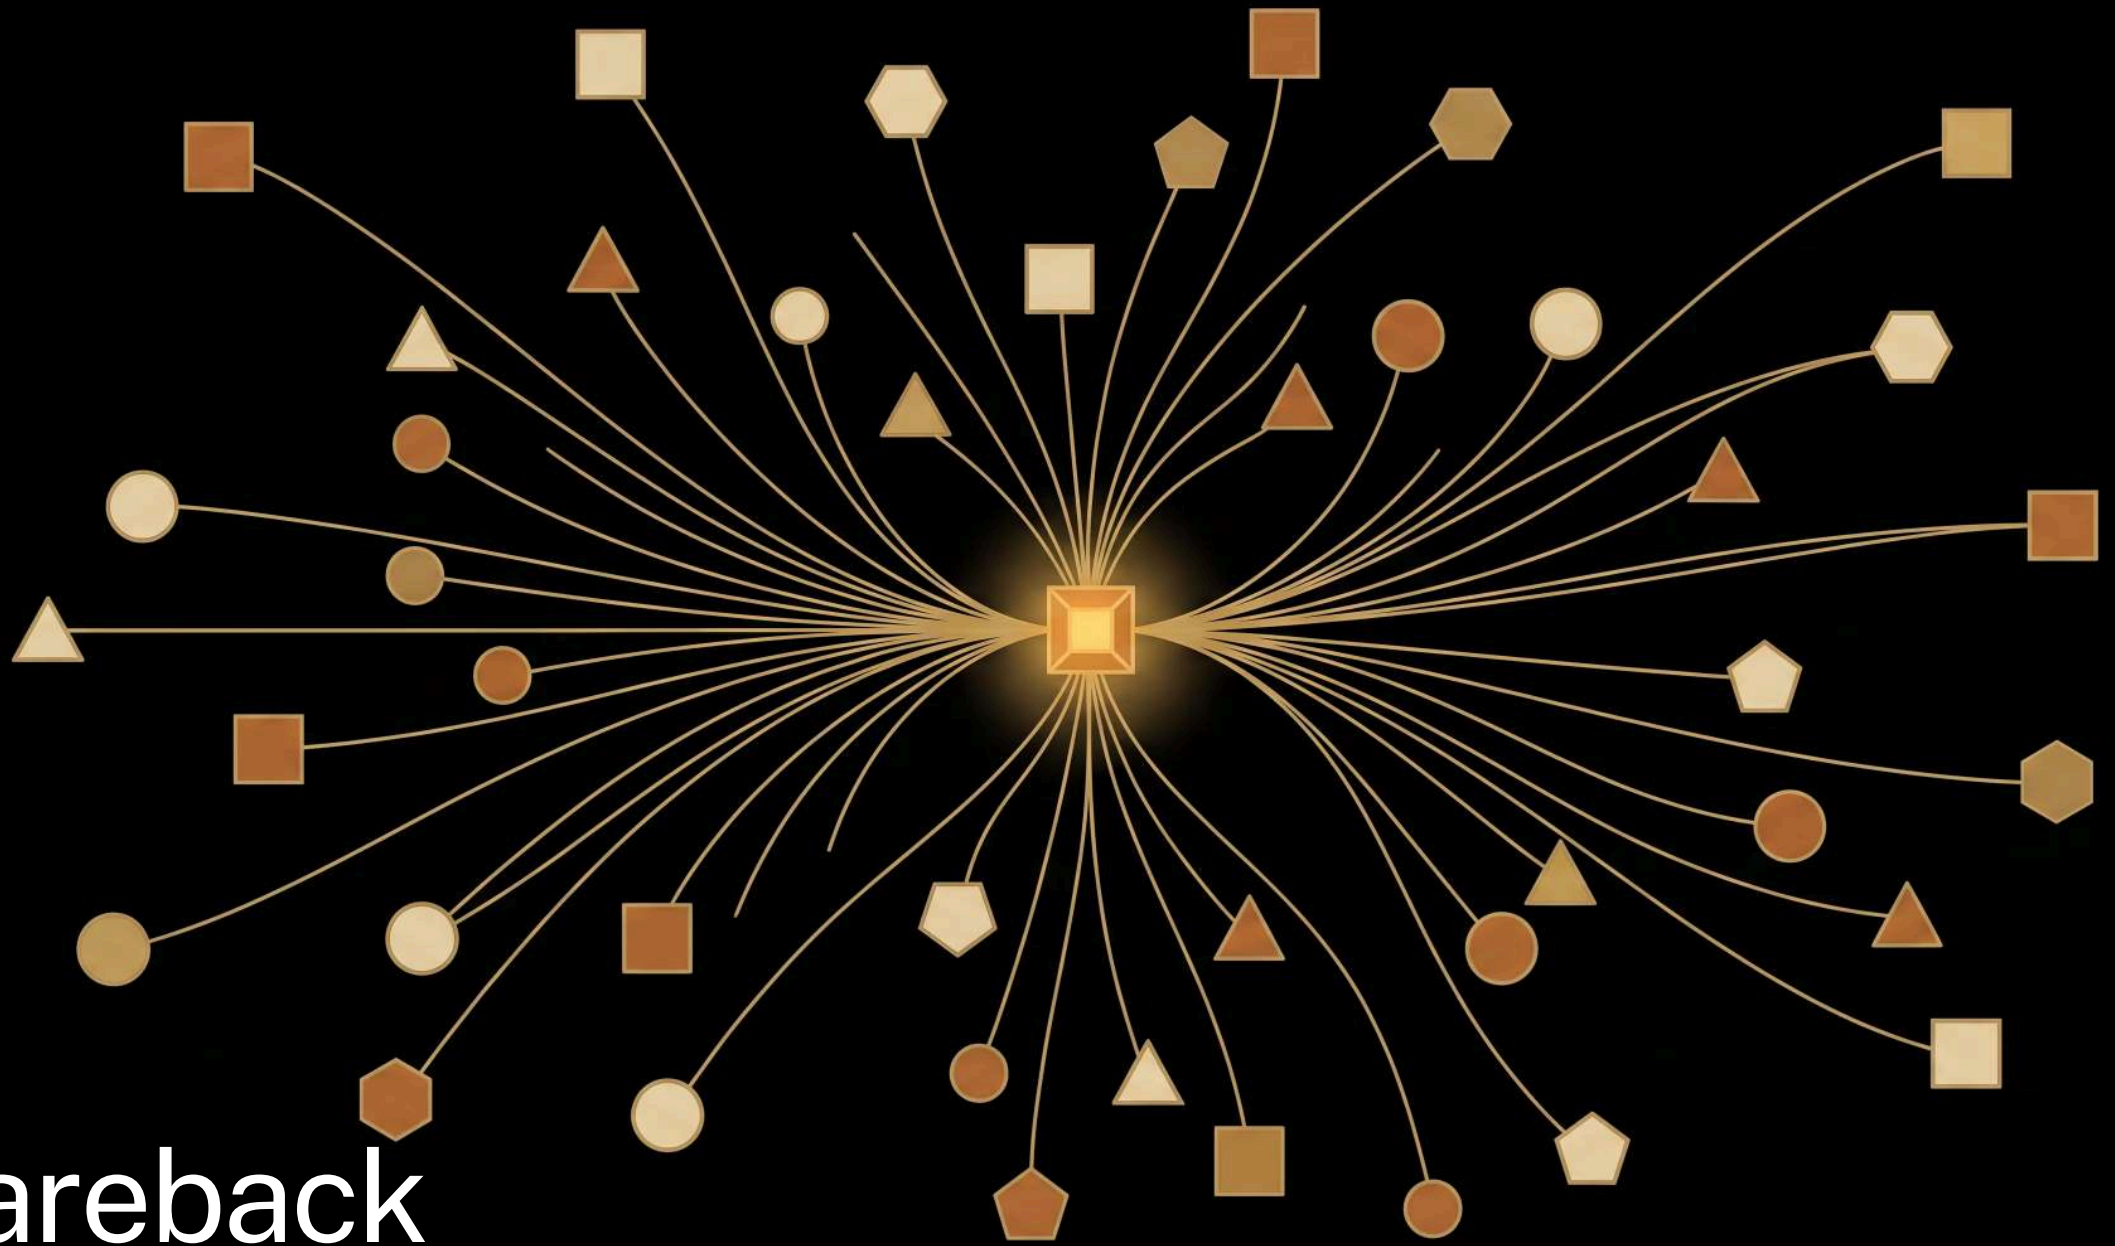

Shareback

Speaker notes

What did the tool calls add that the bare model couldn't --- fresh, outside information? Did the friends' words fit the text, or derail it? Invite ~5 groups to share back rather than the whole room. ~5 min.

The *language* of language models

- agent
- tool call

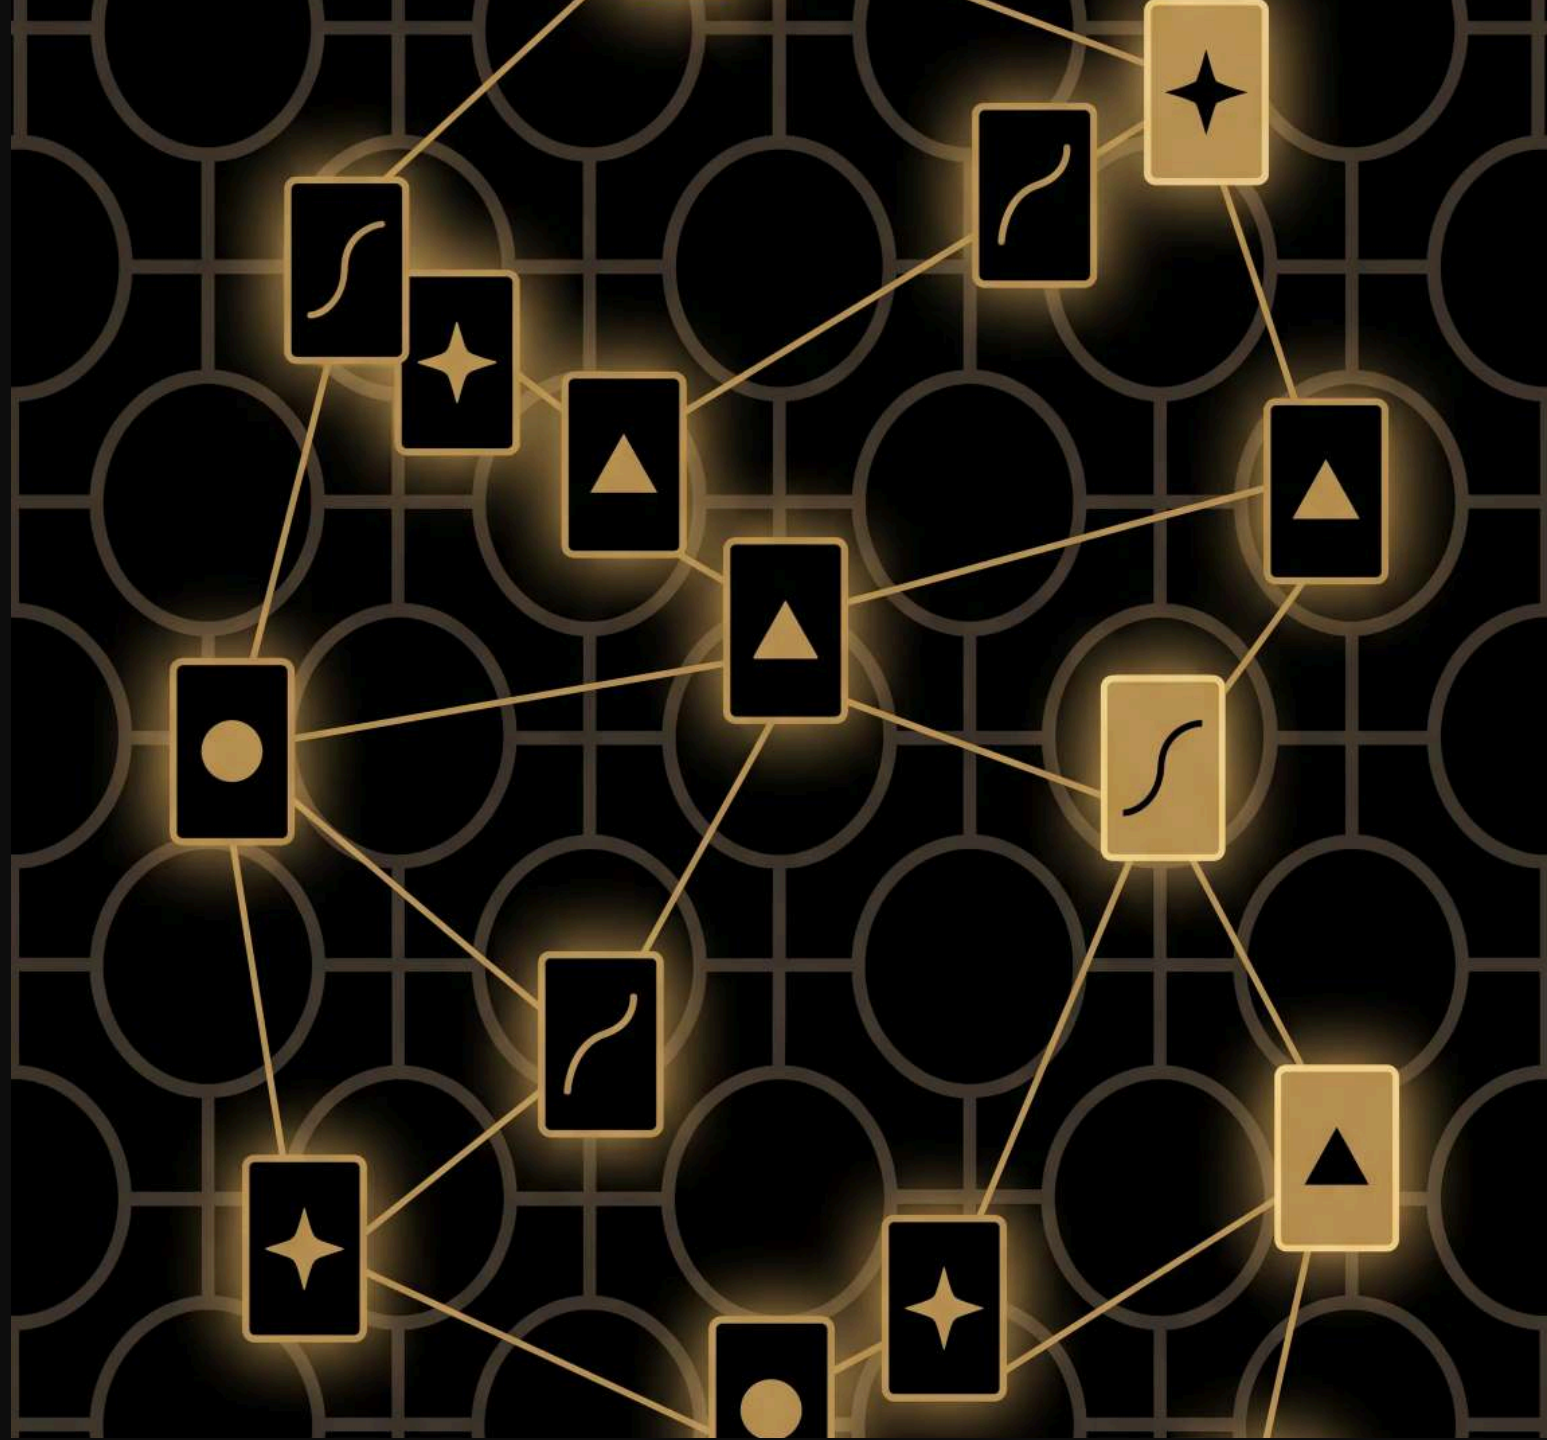

Speaker notes

Everyday sense first, then ours: an agent is a travel agent or a secret agent, someone who acts for you---here it's a model that can act, not just talk. A tool call sounds like reaching for a hammer---here the model asks for something outside itself and waits for the answer. If anyone asks what runs the loop: *you* did. The harness is the scaffolding around the model that pauses it, makes the call, and splices the result back---here, that was the students themselves.

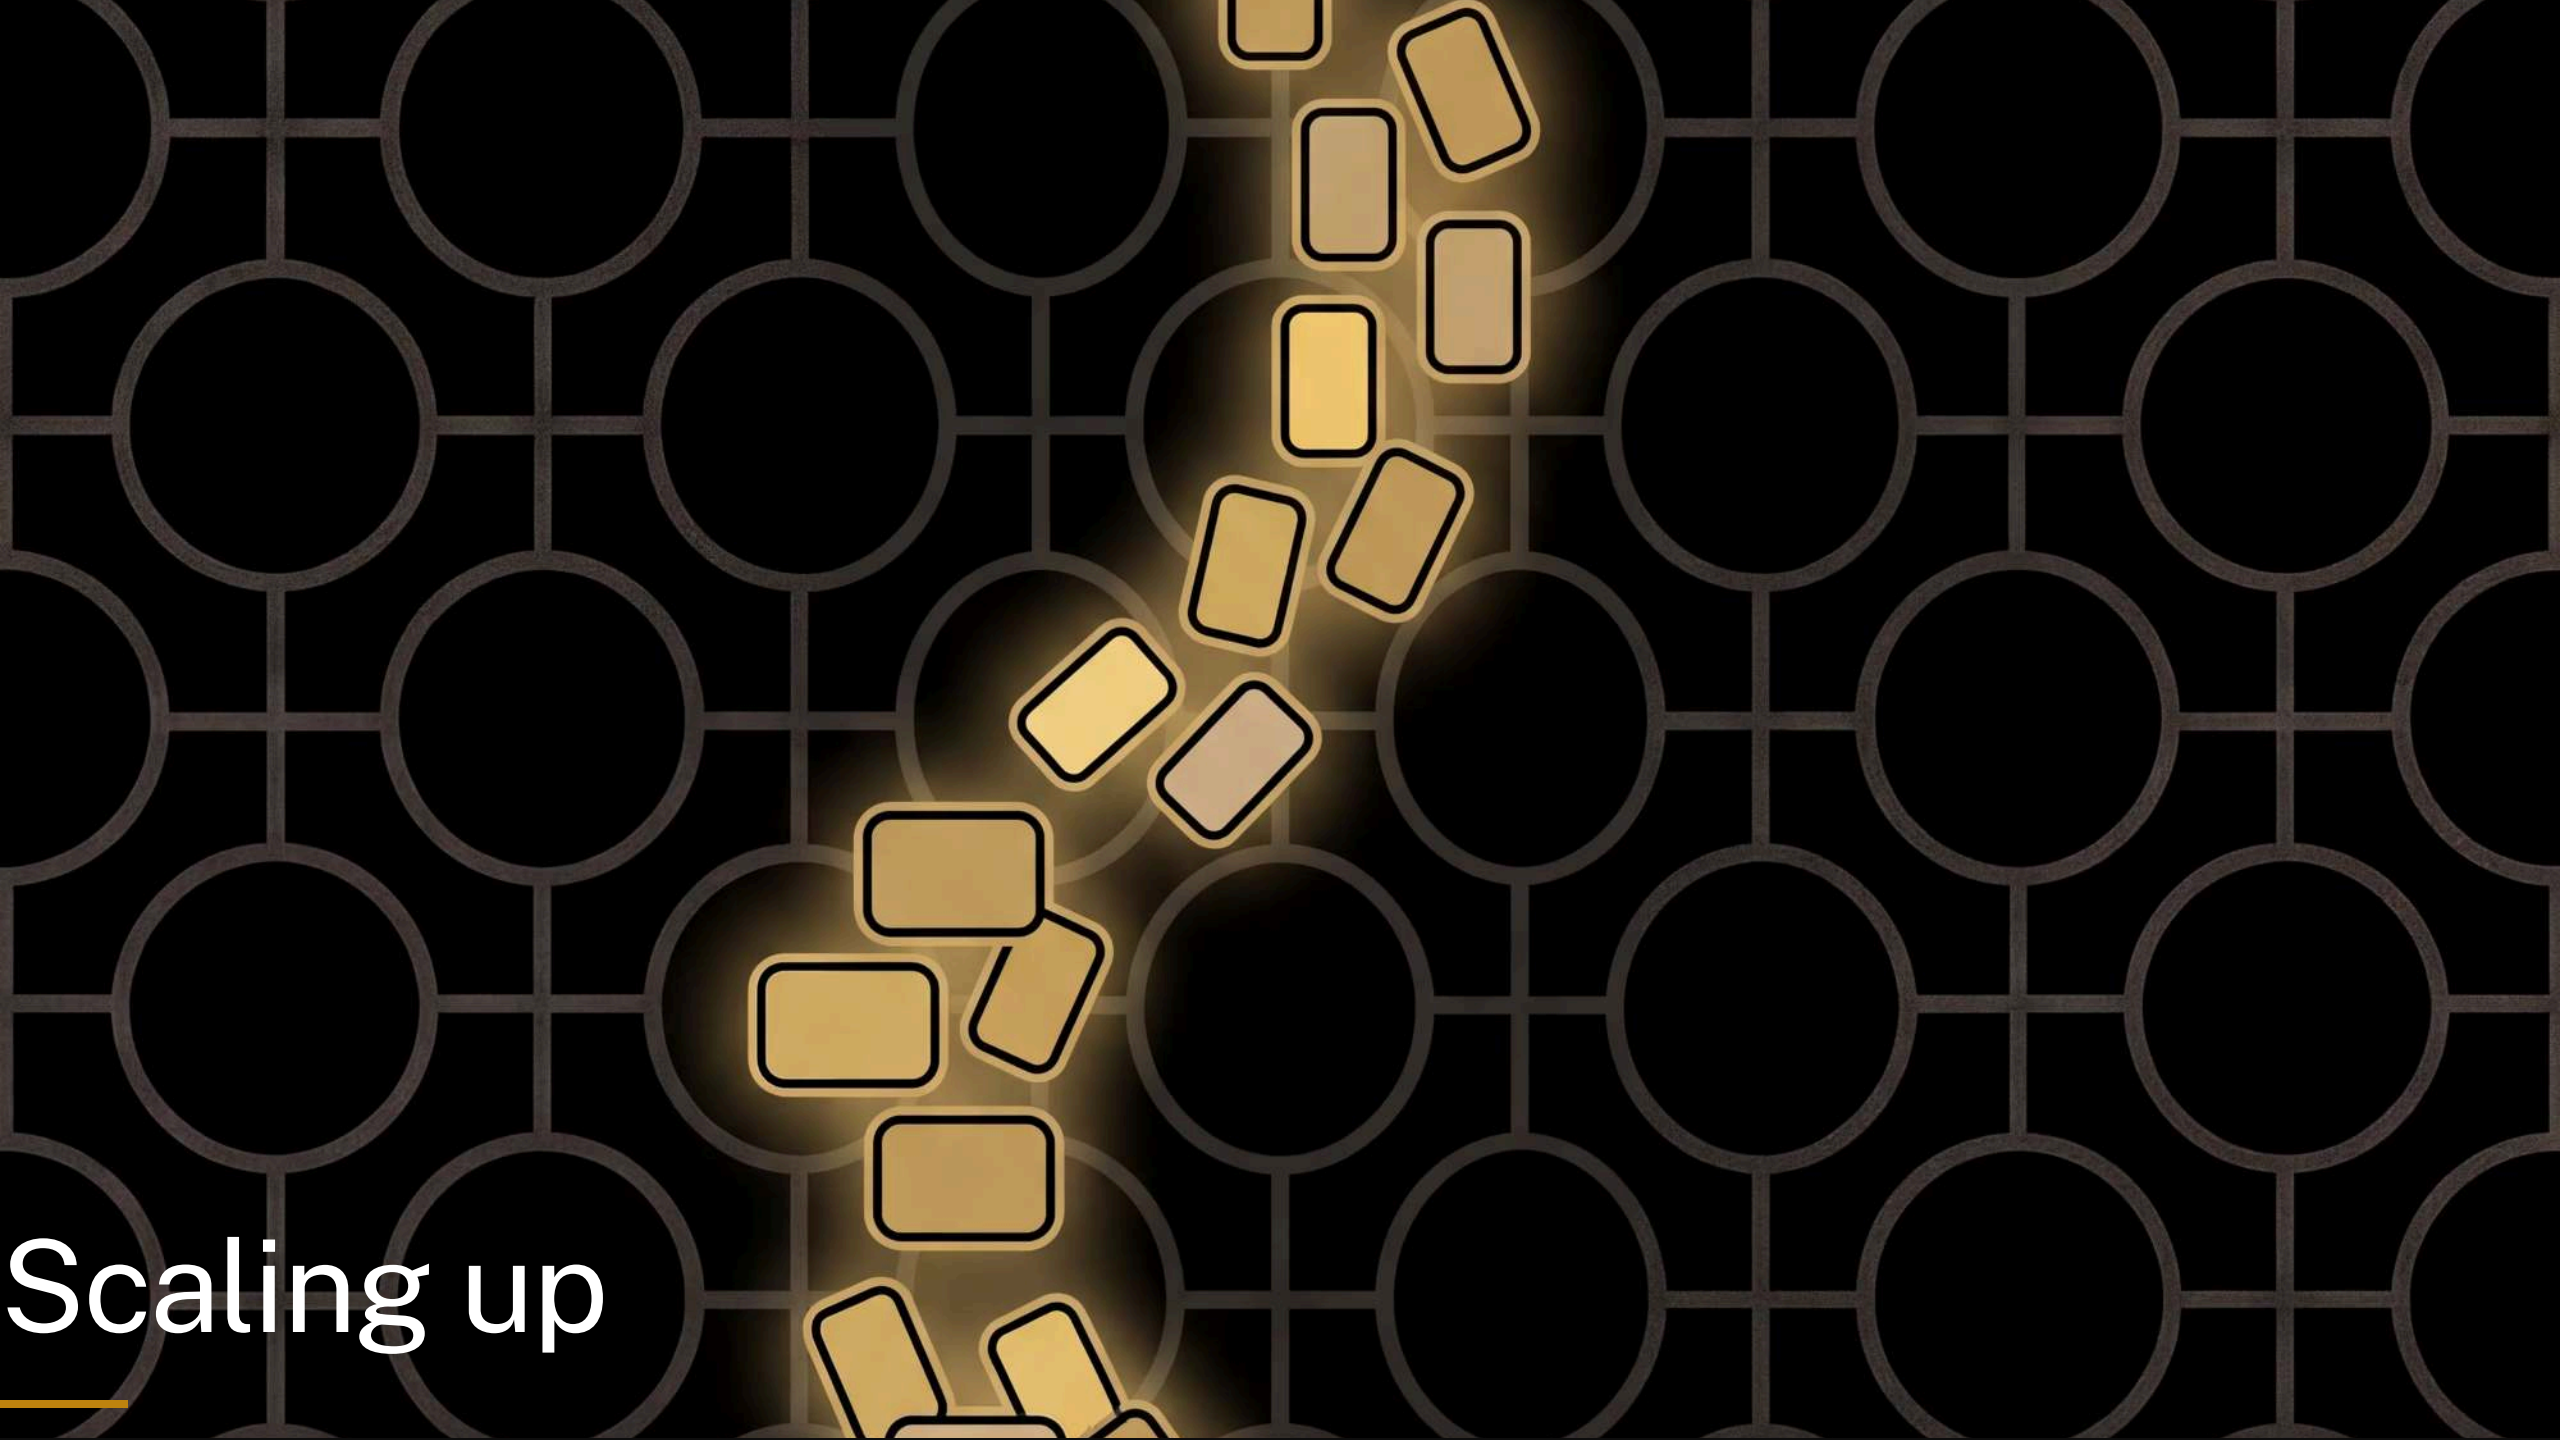

Scaling up

Speaker notes

The reassurance beat: what they built isn't a cartoon of an LLM, it's the real mechanism at human scale---look at the context, weigh every possible next word, roll the dice, write it down, repeat. So what's actually different? The whole section is run against their own grid rather than as a set of separate diagrams. Four times over, the grid is made to fail at something on stage, and the thing that fixes it turns out to be a real part of a transformer: more context (and why the grid can't have it), attention, parameters, post-training. Then scale. ~12 min---it's a lot of slides but most are one beat each, so keep moving.

One word of context

see spot ,

2 options — roll the die!

	run	,	spot	.	see
run		|		||	
,	|		|		
spot	|	|			
.					|
see			|		

Speaker notes

Back to the worked example. We're sitting on a comma with two options and no way to choose between them --- so we roll. But look at why we're stuck: the model is looking at a comma, and every comma in the text looks identical to it. It cannot tell whether it just wrote `run` or `spot`. That's not bad luck, it's the whole design: one word of context.

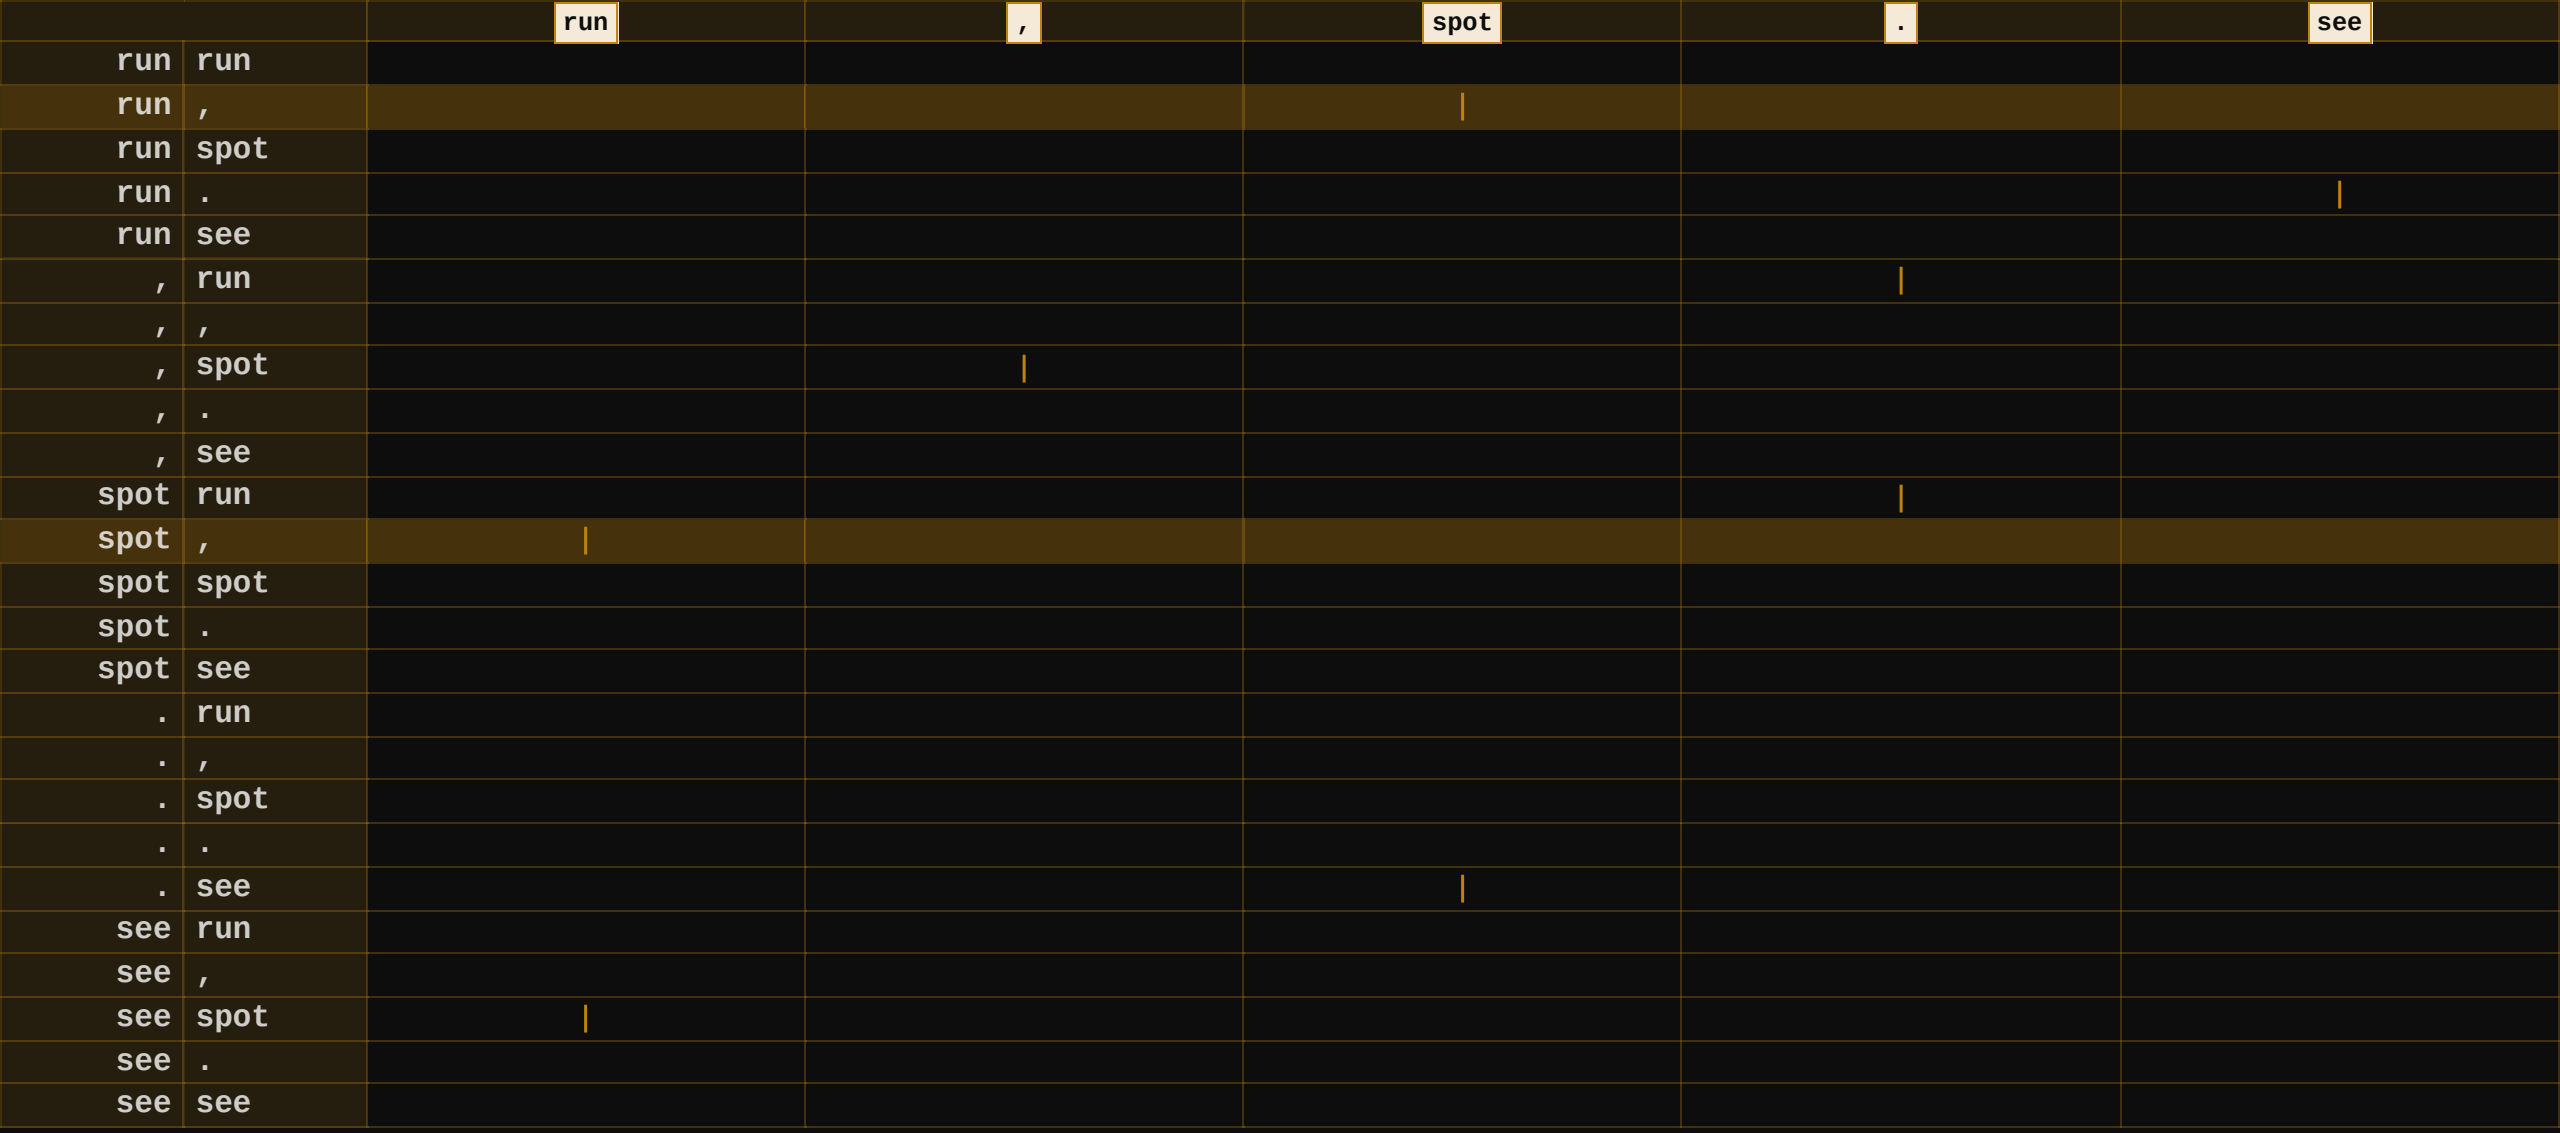

so providing two words of context instead of one helps — what does that cost?

Speaker notes

No heading on this one, so say it: two words of context. Give it two words and the guess stops being a guess. `run ,` is always followed by `spot`; `spot ,` is always followed by `run`. Same text, same tallies, no dice needed --- the ambiguity was never in the language, it was in how little the model could see. This is a trigram model, and there's a lesson on the website that runs it with pen and paper.

one word: 5 rows

two words of context

25 rows, one per possible pair

and almost every one of them is empty

Speaker notes

This is the bill. The grid needs a row for every context it might find itself in, so going from one word to two takes 5 rows to 25 --- and most of them stay empty, because our text never contains `see see` or `. .`. Three words would be 125. Each extra word of memory multiplies the grid by the size of the vocabulary.

The grid nobody can build

context	rows in the grid
1 word	5
2 words	25
3 words	125
10 words	9,765,625

Claude keeps **200,000 words** in view, over a vocabulary of **100,000**

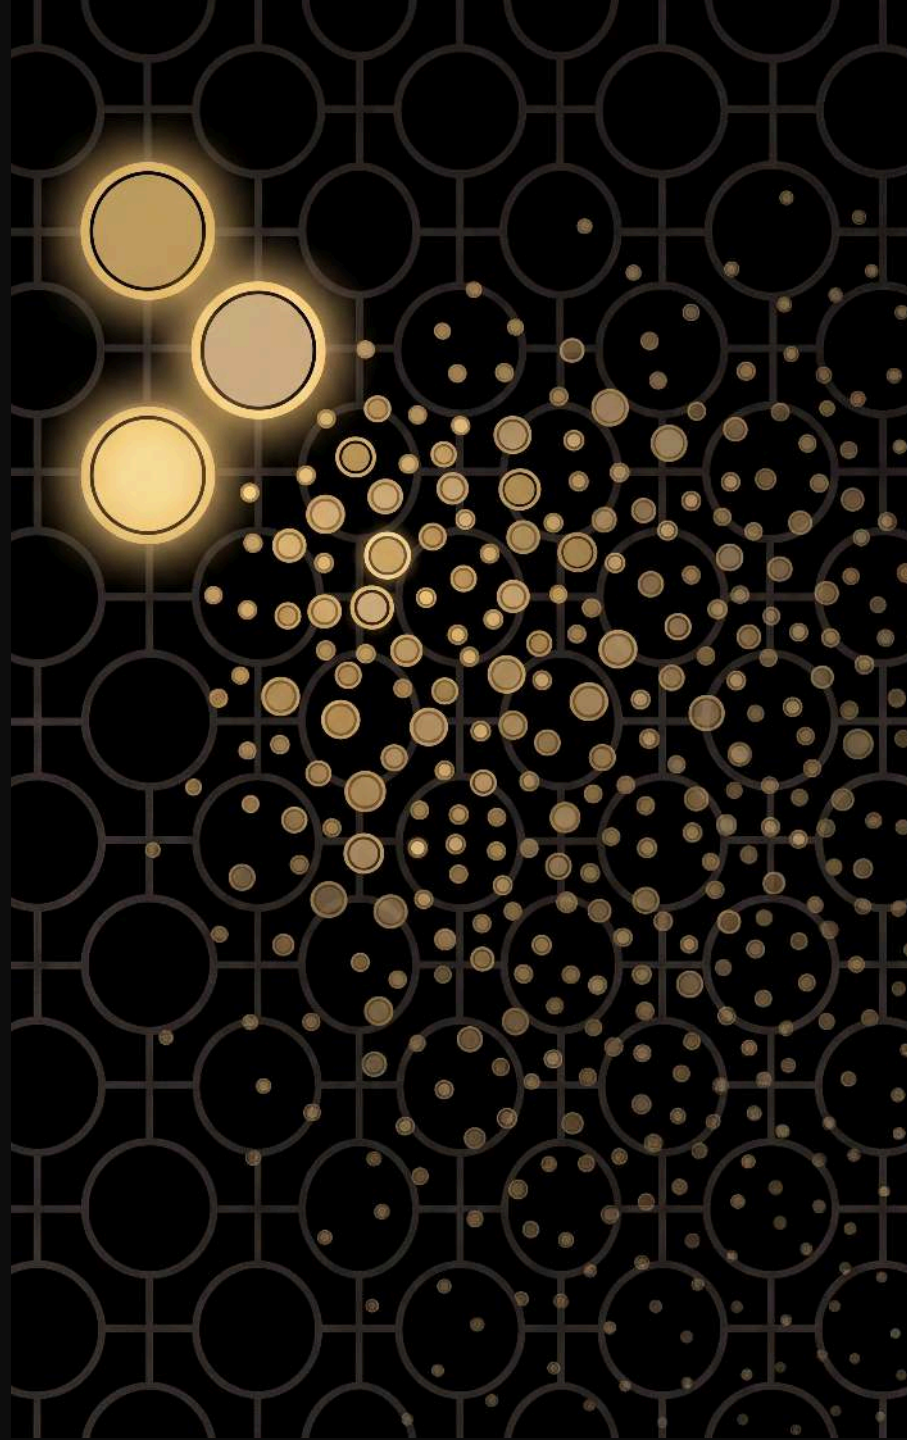

Speaker notes

Read the table, then drop the last line. Their vocabulary is 5, so it's 5-to-the-n. A real model's vocabulary is around 100,000, which is 10-to-the-5 --- so seventeen words of context is already 10-to-the-85 rows, and the observable universe only has about 10-to-the-80 atoms. Seventeen words. Claude holds two hundred thousand. The grid is not big, it is impossible, and that's not rhetoric --- you run out of universe long before you run out of sentence. So something else has to be going on.

Attention: a row that changes shape

	run	,	spot	.	see
run		|		||	
,	|		|		
spot	|	|			
.					|
see			|		

no new rows — this is the grid you built

Speaker notes

Back to their own grid, whole, one last time. Point at the comma row: one mark under `run`, one under `spot`, and that is the 50/50 we rolled on. Attention is not a bigger grid---watch what happens to that one row while everything else stays exactly where it is.

Attention: a row that changes shape

	`run`	`,`	`spot`	`.`	`see`
`run`					
`run` `,`					
`spot`					
`.`					
`see`					

with `run` behind it, the row leans to `spot`

Speaker notes

Same grid, same row, but the model now looks back over what it can see and weighs which earlier words matter here. With `run` before the comma the row is redrawn on the spot: nearly all the weight on `spot`. Nothing was stored --- this row didn't exist a moment ago.

Attention: a row that changes shape

	run	,	spot	.	see
run					
see spot ,					
spot					
.					
see					

with `see spot` behind it, the same row leans the other way

Speaker notes

And the payoff: `see spot` before the comma pulls the same row the other way. One row, re-shaped per context, instead of a row per context---which is exactly the bill the last slide said nobody can pay. That's attention, and it's why transformers can have a context window at all.

Attention: learning where to look

the wifi **password** is **swordfish** … the **password** is **?**

the next word is *swordfish* — attention finds where `password` last appeared and copies what came after

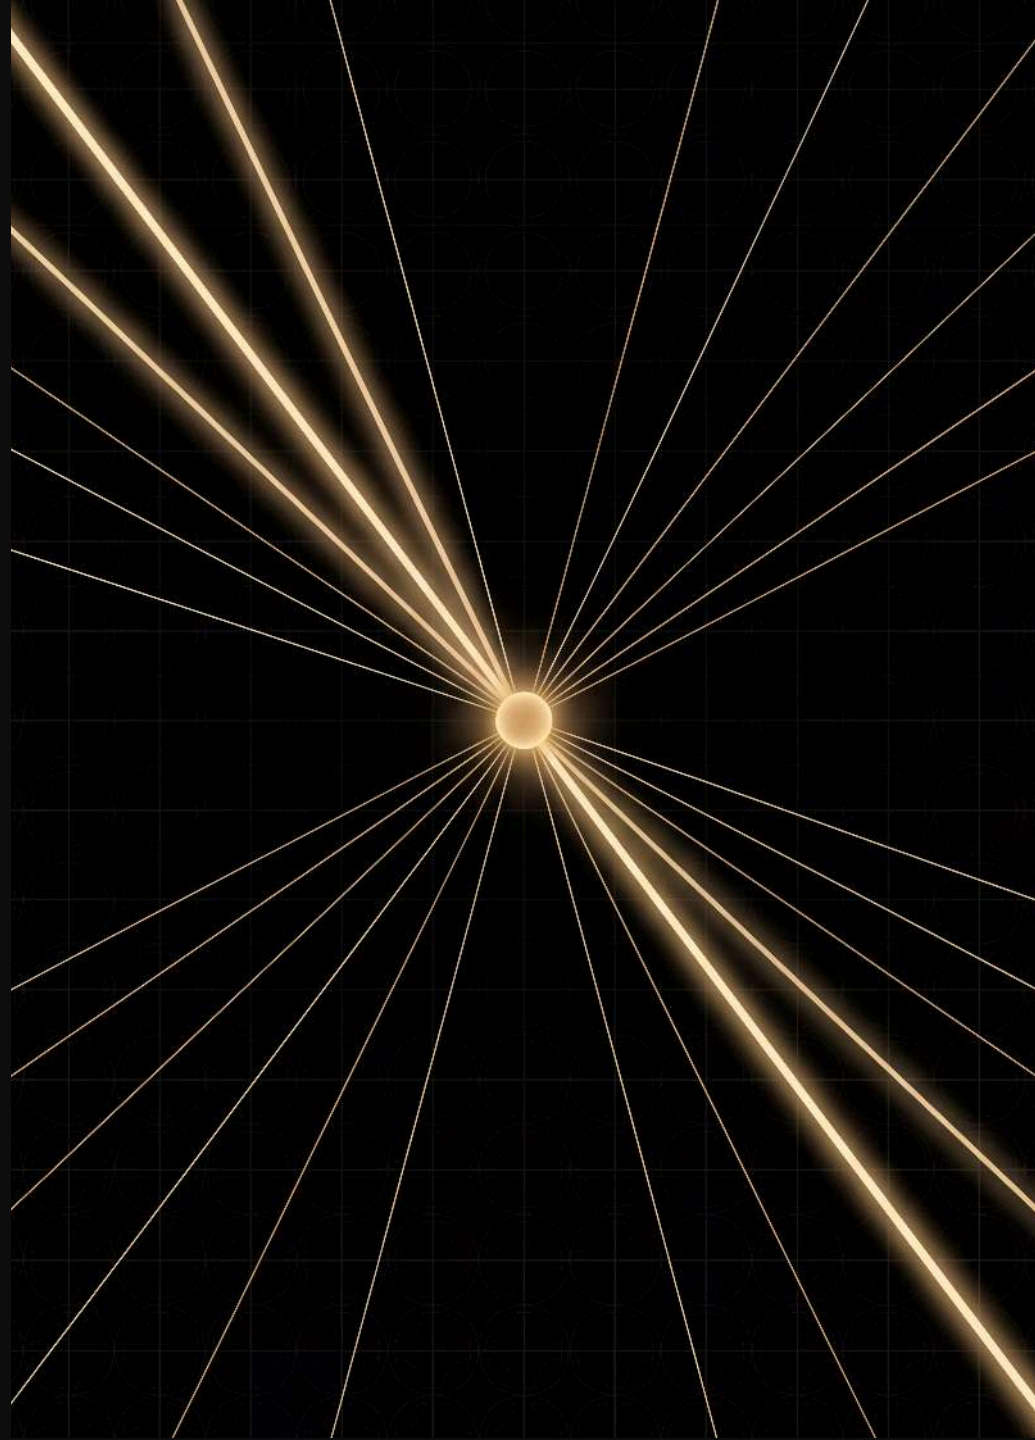

Speaker notes

A concrete one to make it stick. The model spots the earlier "the password is swordfish" and repeats what followed. That look-back-and-copy move is an induction head, and the Induction Heads lesson on the website is exactly this pattern --- "find the last time I saw this word, copy what came next" --- run with pen and paper. The weights aren't hand-set either: the model learns where to look.

You can point at the pattern

see spot , run .

rolled 3 → .

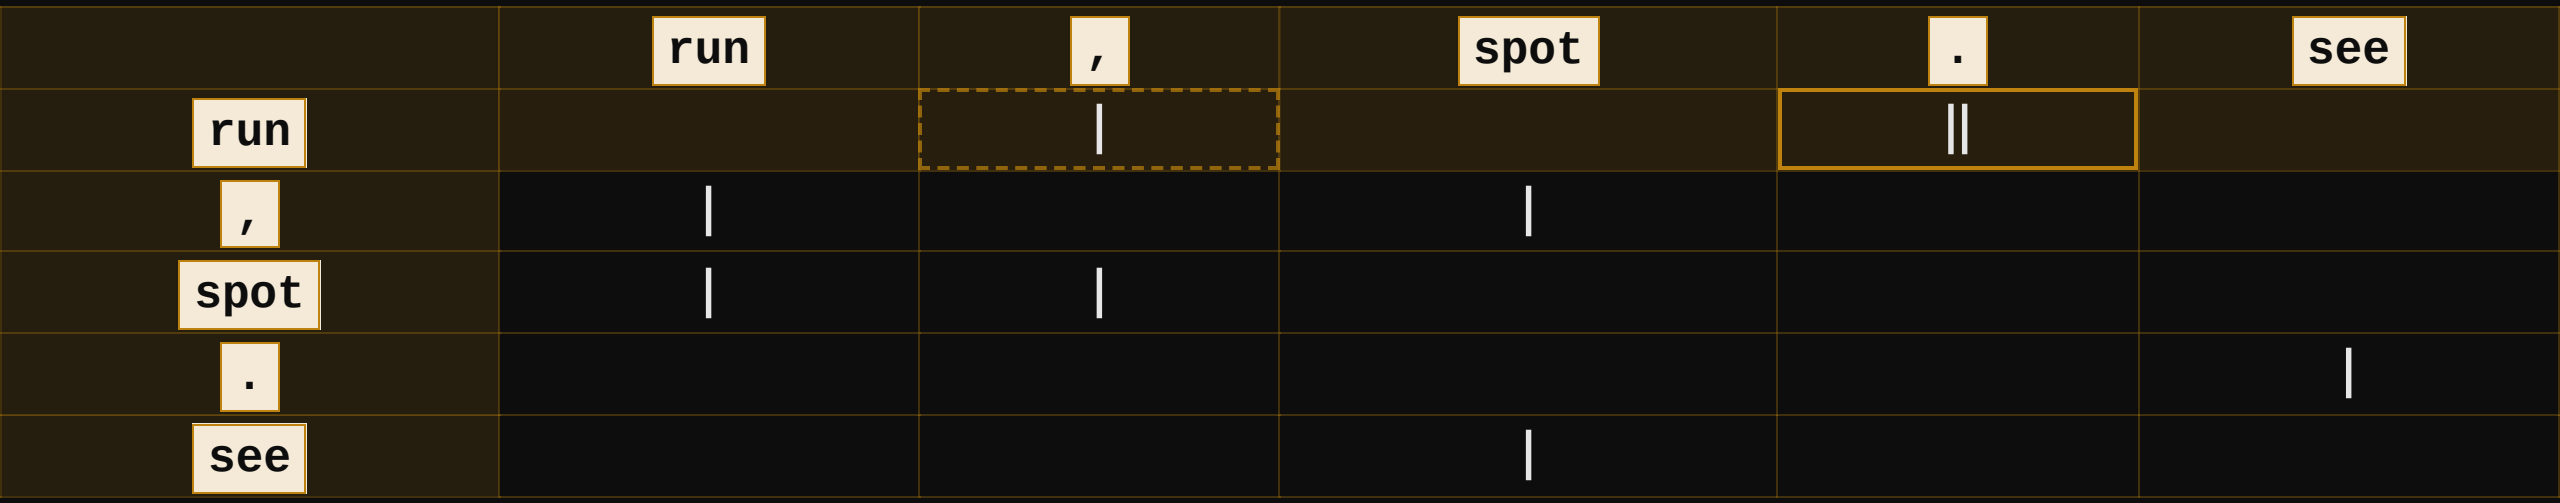

Speaker notes

Second failure, and it starts with something the grid is good at. `run` is followed by a full stop twice, and there are the two marks, in that cell, and you can put your finger on them. Everything the model knows is somewhere you can point at. Hold that thought.

cat and dog are strangers

the cat sat . the cat ran . the dog sat .

	the	cat	sat	.	ran	dog
the		‖				|
cat			|		|	
sat				‖		
.	‖					
ran				|		
dog			|			

Speaker notes

A different text, same idea. `cat` and `dog` do exactly the same job here---both come after `the`, both get followed by verbs. But look at the two rows: `cat` has seen `sat` and `ran`; `dog` has only seen `sat`. Ask this grid to continue "the dog" and it will never, ever say `ran`, because that cell is empty and an empty cell means impossible. You know "the dog ran" is fine. The grid has no way to know, because nothing it learned about `cat` is connected to `dog` in any way. Every cell is its own little island.

Tallies become knobs

	the	cat	sat	.	ran	dog
the	0.03	0.46	0.04	0.02	0.04	0.41
cat	0.04	0.01	0.48	0.03	0.42	0.02
sat	0.12	0.02	0.01	0.78	0.04	0.03
.	0.82	0.05	0.03	0.02	0.04	0.04
ran	0.11	0.03	0.03	0.77	0.02	0.04
dog	0.05	0.02	0.46	0.03	0.40	0.04

a number in every cell — even the one your grid left empty — because it is **billions of knobs** doing the remembering, not cells

Speaker notes

The fix for the island problem, shown on the same grid---same text, same rows, the tally marks ghosted in place with a number under each one. Two things to point at. First, the ringed cell: `dog` followed by `ran`, which the tallies left empty and which now has 0.40 in it. Second, the `cat` and `dog` rows, which have come out nearly identical, because in this text the two words do the same job. Where those numbers come from is the slide. A transformer doesn't give each pattern its own cell; it spreads every pattern across billions of numbers, each nudged a tiny amount after every wrong guess in training. Because `cat` and `dog` turn up in the same kinds of places, training pushes them onto overlapping numbers---so what the model learns from `cat` is already, for free, something it knows about `dog`. That's generalisation, and it's the thing the grid structurally cannot do. (Careful with the picture: the model isn't storing this grid either. Nothing here is looked up; every number is computed on demand, which is why it can fill a cell it never saw.) The cost is the flip side of the slide before: there is no longer a cell to point at. Ask where a transformer stores "the dog ran" and the honest answer is "smeared across a few billion numbers, and we're still working on reading it back". That's most of what interpretability research is.

prompt: “what is the capital of France?”

your grid `“see spot, run. see”`

it isn’t refusing to answer — **answering isn’t a thing it does**

Speaker notes

No heading --- the line is "let's ask your grid a question", and the slide just shows what comes back. Third failure, and the funniest one. Feed their model a question and it does what it always does: continues. It isn't being unhelpful and it isn't broken; nothing in "tally which word follows which" has any notion of a question and an answer. This is the honest state of a model that has only ever been trained to continue text.

From continuing to answering

the `?` row, tallied over text from the web

? →	what	how	why
	卌	|||	||

the same row, tallied over conversations

? →	it	sure	yes
	卌	|||	||

post-training: the same tallying again, over text made of *conversations*

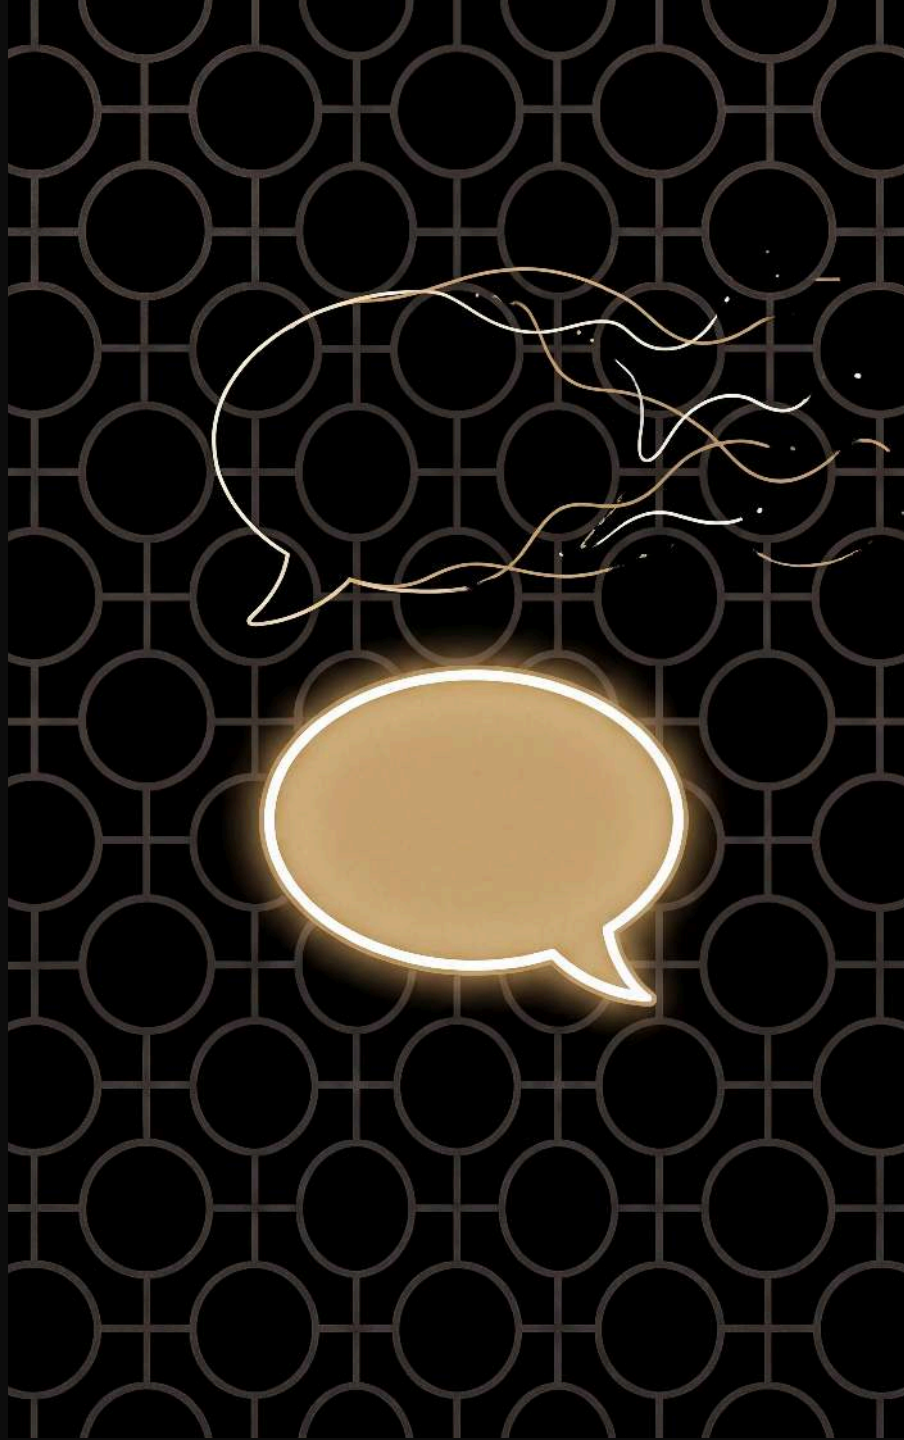

Speaker notes

Same machinery one more time, so read it off the rows. On the web, what follows a question mark is usually another question --- so a model trained on it continues with questions, which is exactly what a frontier model straight out of pre-training does: their problem, with better grammar. Tally the same row over transcripts of conversations instead and what follows a `?` is the start of an answer. That's post-training: a second, much smaller round of training on example conversations, plus feedback on which answers people preferred. So the thing that comes back with "Paris" a slide ago isn't running a different mechanism --- it learned from text where answering is what happens next. Two honesty notes if anyone asks. A real model uses everything before the question mark, not just the mark itself --- these rows are the one-word-of-context version of the point. And the mechanism hasn't changed: same predict-the-next-word loop, same dice, different text to learn from, so different habits on top.

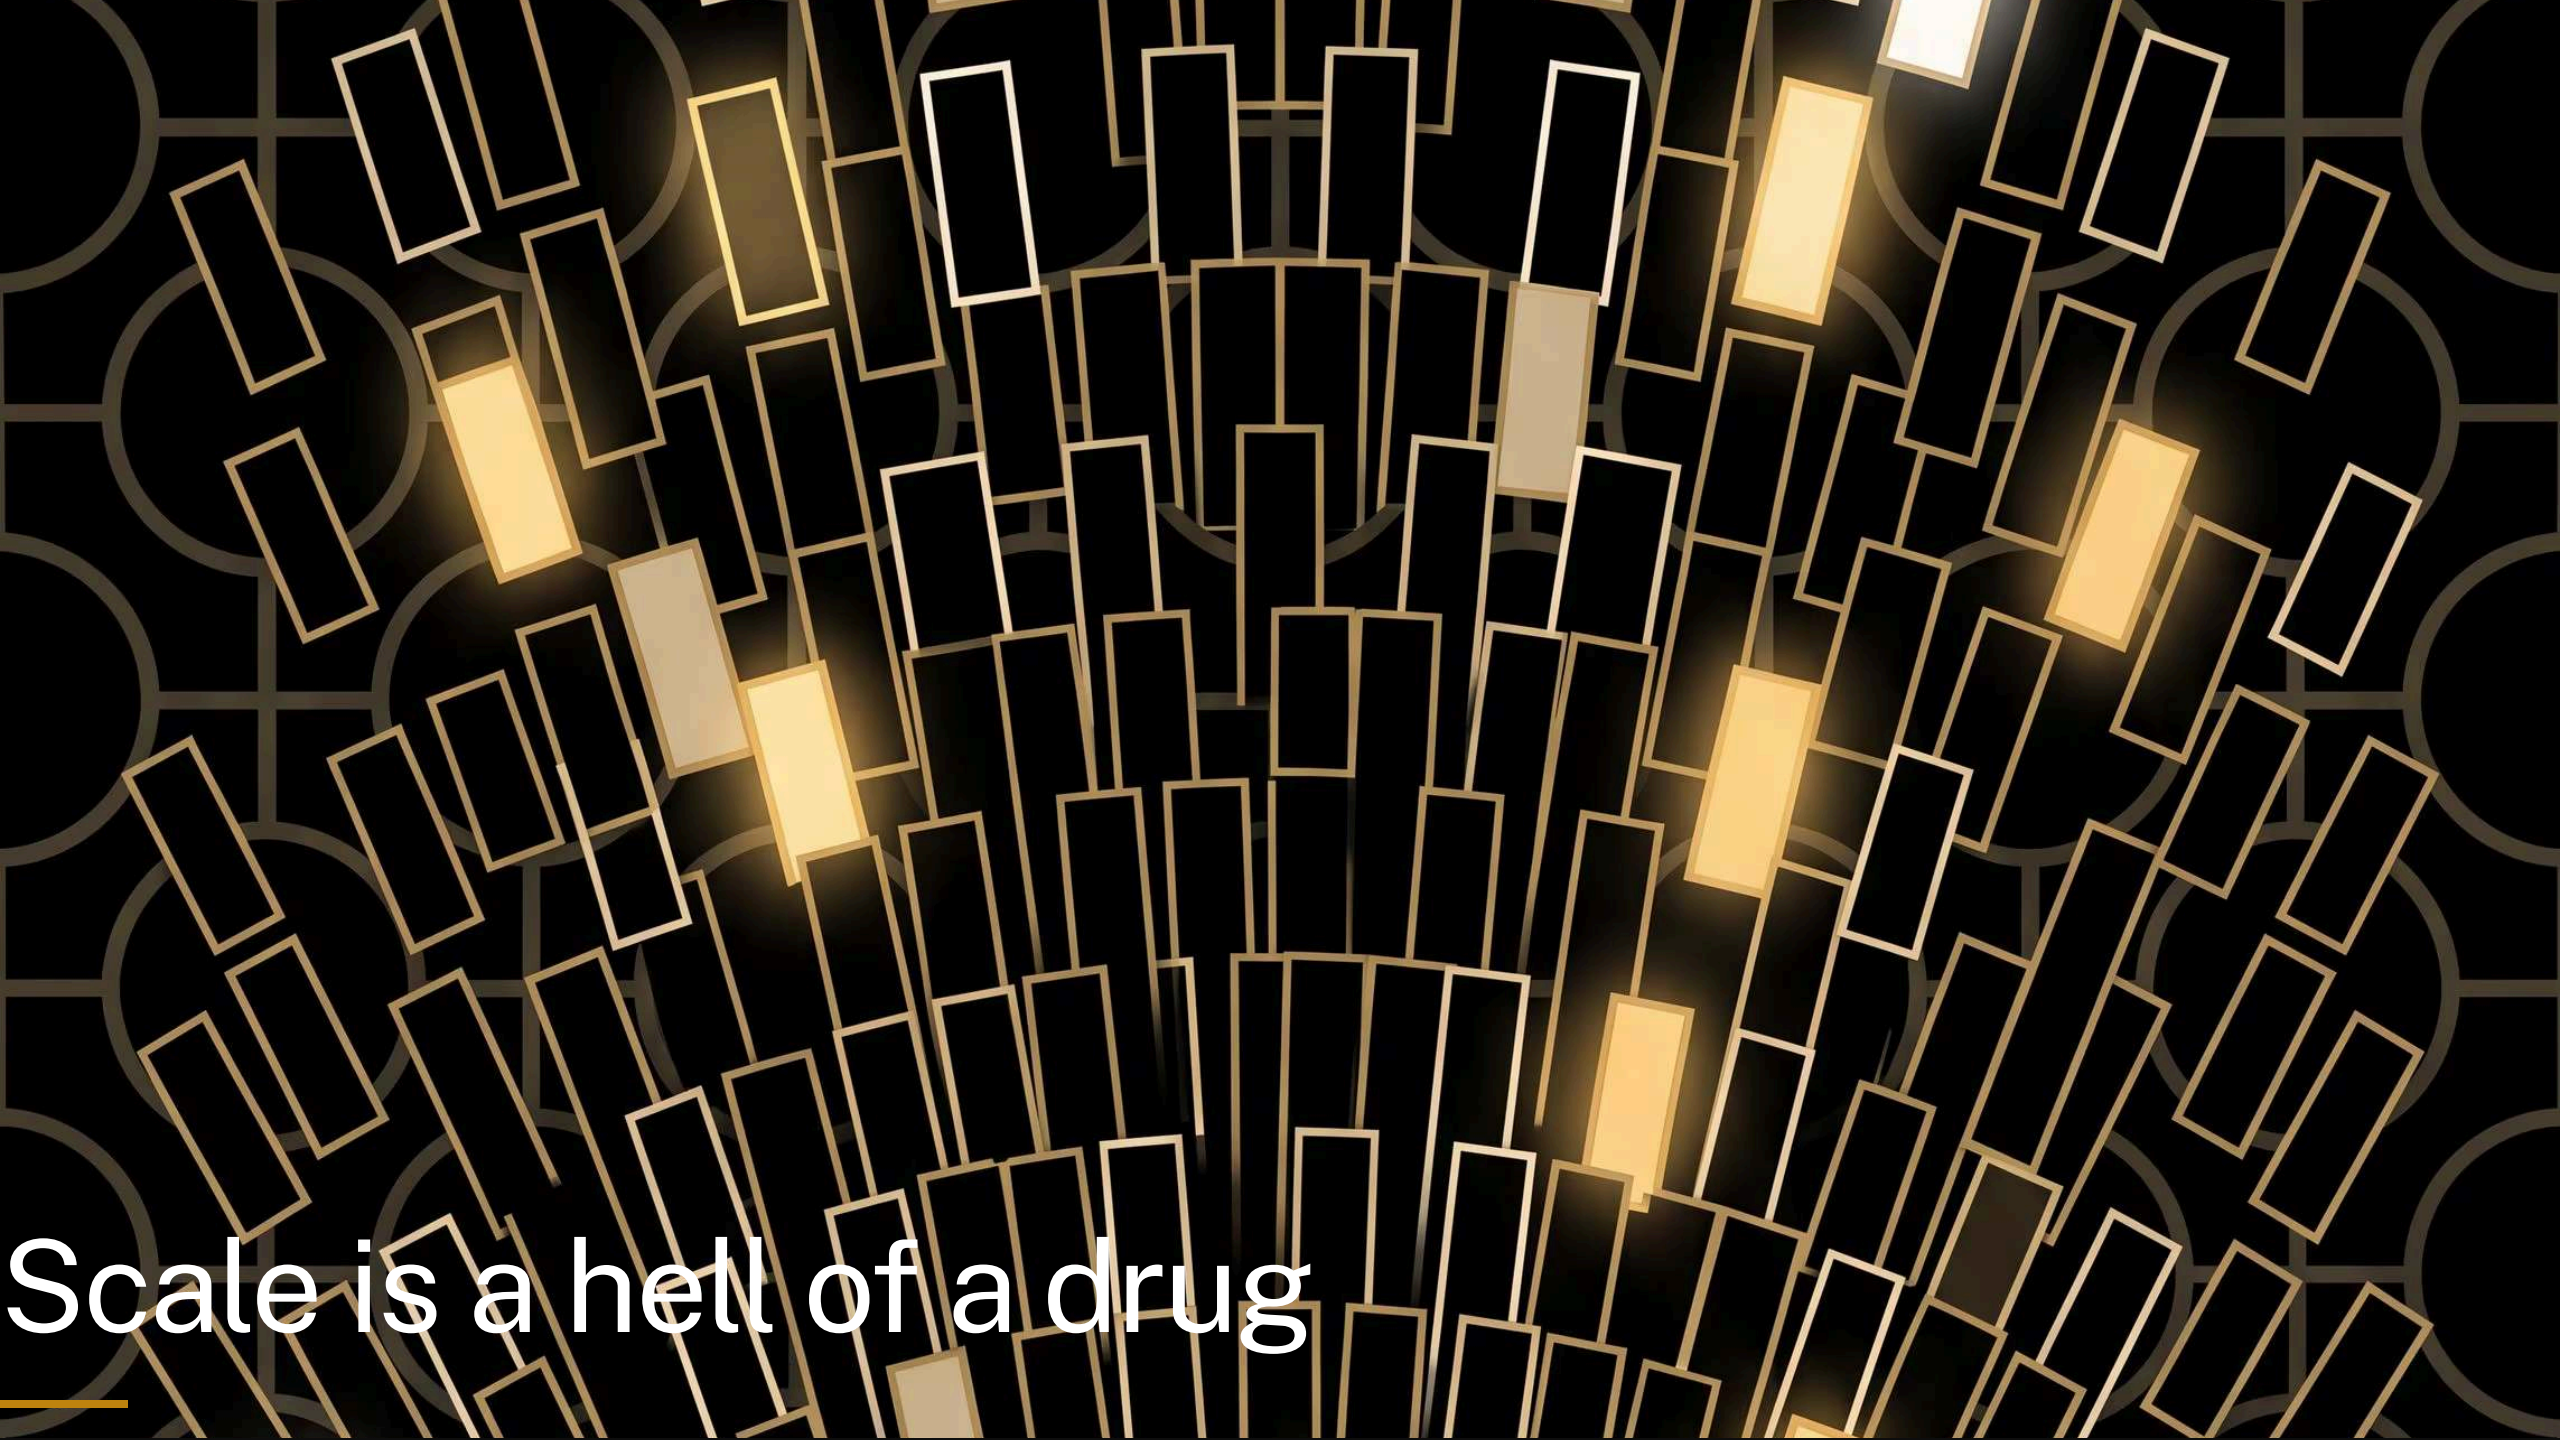

Scale is a hell of a drug

Speaker notes

Four things fixed. None of them changed the loop. The last difference is just dosage --- and it's worth taking slowly, because this is the one people's intuition has no grip on. Keep your eye on the gold square: that's their grid, drawn at true scale, on every slide.

your grid

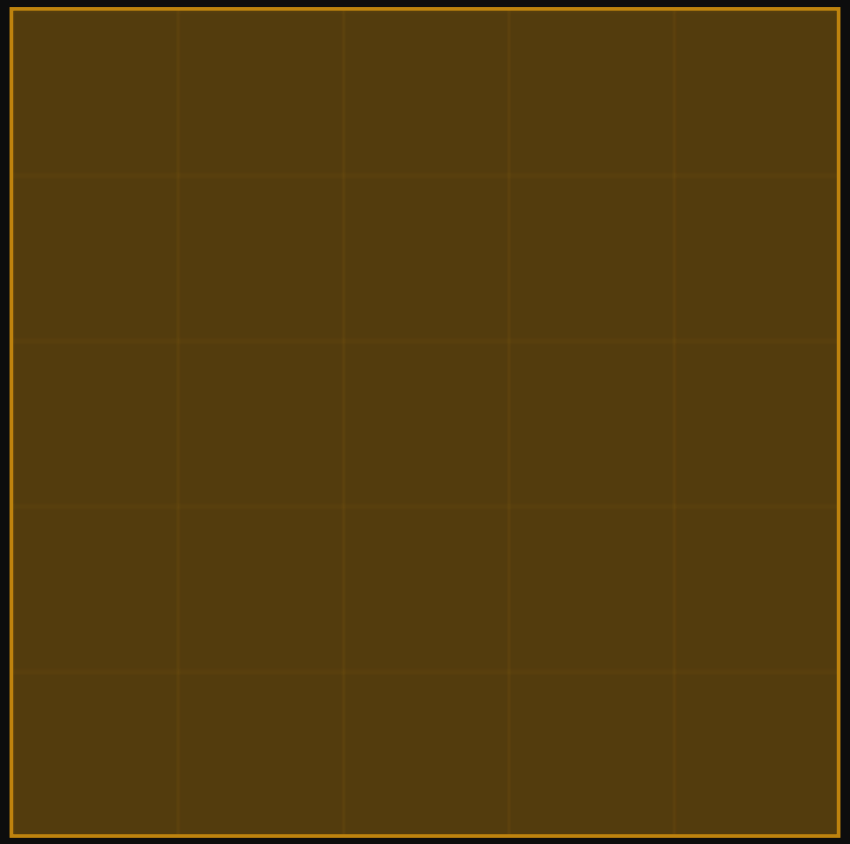

“Run, Spot, run. See Spot run.”

5 tokens · 25 cells

Speaker notes

The next three slides are figure-only, so you carry them: this one is "your grid". Five tokens, twenty-five cells. At the booklet's 1cm per cell that's a five-centimetre square --- smaller than a beer coaster.

your grid

Green Eggs and Ham

70 tokens · 4,900 cells

Speaker notes

Same again, no heading --- read the caption out. A whole Dr Seuss book is 70 distinct tokens, so ~4,900 cells --- a sheet about 70cm square, call it a poster. And their grid is the little gold square in the corner. Note the vocabulary is what matters here, not the length: the book is a couple of thousand words, but Seuss reuses them relentlessly.

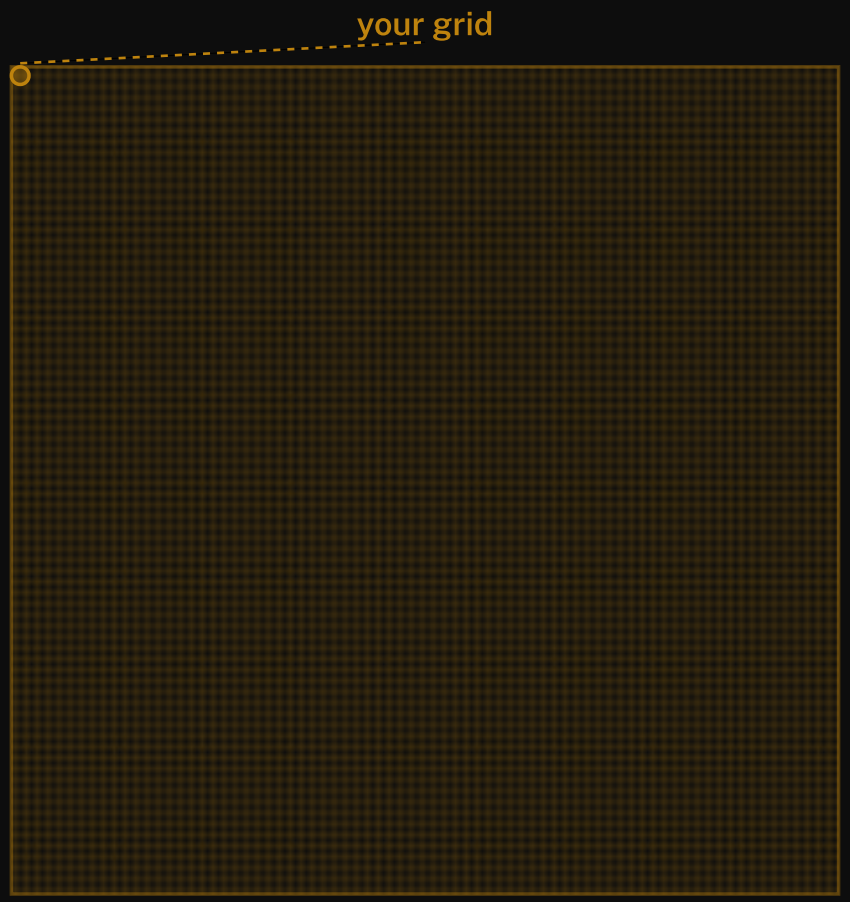

Frankenstein

7,441 tokens · 55 million cells

the cells are now too small to draw

Speaker notes

And the third: Frankenstein, the novel in the booklet. 7,441 distinct tokens, 55 million cells. The grid lines have stopped being drawable --- that grey is the grid. On paper it's about 5,500 square metres, most of a football field, and their model is the gold speck. This is also why the booklet doesn't print a grid: at this size you list only the cells that aren't empty.

a real vocabulary: **100,000** tokens, **10 billion** cells, **1 km^2** of paper

Speaker notes

Stop drawing and start measuring, because the picture gives out here. A commercial model's vocabulary is around 100,000, so its bigram grid alone is 10 billion cells --- a square kilometre of paper, near enough the whole Acton campus. And that is still the toy version: one word of context, tally marks, none of the four fixes we just went through. Two more ballparks if you want them. The real thing: a frontier model's ~1.5 trillion parameters at 1cm each is about 150 km^2, a square 12km on a side, a hundred Acton campuses, twenty-odd Lake Burley Griffins. And time: if hand-training 100 tokens took you 10 minutes, the tens of trillions of tokens these are trained on would take about 5.7 million years at your rate --- you'd have had to start around when our ancestors split from the chimpanzees to be finishing now. The punchline isn't any one figure. It's that the mechanism didn't change, only the dosage.

still just **tokens in → tokens out**

Speaker notes

Every reply from Claude or ChatGPT is sampled one word at a time, exactly like their dice rolls---that's why "regenerate" gives a different answer. More context, learned where to look, billions of numbers instead of tallies, a round of manners---but the loop is the one they ran by hand.

Questions

what new questions do you have about large language models?

Speaker notes

Open the floor on the first question and let it run. Bring up the second only once the questions thin out---it turns the session outward, from what the grid does to what they'll do differently on Monday.

Questions

what new questions do you have about large language models?

how will this change the way you think about and use LLMs in the future?

Australian National University

School of Cybernetics

www.llmsunplugged.org